



Red Hat

Training and Certification

Student Workbook

ACM 2.4 DO480

Multicluster Management with Red Hat OpenShift Platform Plus

Edition 0





Join a community dedicated to learning open source

The Red Hat® Learning Community is a collaborative platform for users to accelerate open source skill adoption while working with Red Hat products and experts.



Network with tens of thousands of community members



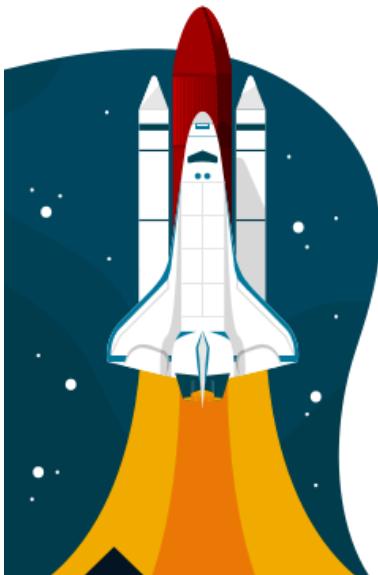
Engage in thousands of active conversations and posts



Join and interact with hundreds of certified training instructors



Unlock badges as you participate and accomplish new goals



This knowledge-sharing platform creates a space where learners can connect, ask questions, and collaborate with other open source practitioners.

Access free Red Hat training videos

Discover the latest Red Hat Training and Certification news

Connect with your instructor - and your classmates - before, after, and during your training course.

Join peers as you explore Red Hat products

Join the conversation learn.redhat.com



Copyright © 2020 Red Hat, Inc. Red Hat, Red Hat Enterprise Linux, the Red Hat logo, and Ansible are trademarks or registered trademarks of Red Hat, Inc. or its subsidiaries in the United States and other countries.

Multicluster Management with Red Hat OpenShift Platform Plus



ACM 2.4 DO480

**Multicluster Management with Red Hat OpenShift Platform
Plus**

Edition 0 f54b122

Publication date 20220101

Author: John Smith
Course Architect: Jane Smith
DevOps Engineer: John Smith
Editor: Jane Smith

Copyright © 2021 Red Hat, Inc.

The contents of this course and all its modules and related materials, including handouts to audience members, are
Copyright © 2021 Red Hat, Inc.

No part of this publication may be stored in a retrieval system, transmitted or reproduced in any way, including, but not limited to, photocopy, photograph, magnetic, electronic or other record, without the prior written permission of Red Hat, Inc.

This instructional program, including all material provided herein, is supplied without any guarantees from Red Hat, Inc. Red Hat, Inc. assumes no liability for damages or legal action arising from the use or misuse of contents or details contained herein.

If you believe Red Hat training materials are being used, copied, or otherwise improperly distributed, please send email to training@redhat.com or phone toll-free (USA) +1 (866) 626-2994 or +1 (919) 754-3700.

Red Hat, Red Hat Enterprise Linux, the Red Hat logo, JBoss, OpenShift, Fedora, Hibernate, Ansible, CloudForms, RHCA, RHCE, RHCSA, Ceph, and Gluster are trademarks or registered trademarks of Red Hat, Inc. or its subsidiaries in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a registered trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack Logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation or the OpenStack community.

All other trademarks are the property of their respective owners.

Contributors: John Doe, Jane Doe

Document Conventions	ix
	ix
Introduction	xi
DO480: Multicluster Management with Red Hat OpenShift Platform Plus	xi
Orientation to the Classroom Environment	xiii
Performing Lab Exercises	xiv
1. Managing Multicluster Kubernetes Architectures	1
Introducing Kubernetes Multicluster and Hybrid Cloud Architectures Landscape	2
Quiz: Introducing Kubernetes Multicluster and Hybrid Cloud Architectures Landscape	8
Describing Red Hat OpenShift Platform Plus	12
Quiz: Describing Red Hat OpenShift Platform Plus	16
Deploying Red Hat Advanced Cluster Management for Kubernetes	18
Guided Exercise: Deploying Red Hat Advanced Cluster Management for Kubernetes	23
Lab: Installing the RHACM Operator	31
Summary	40
2. Inspecting Resources from Multiple Clusters Using the RHACM Web Console	41
Navigating the Red Hat Advanced Cluster Management Web Console	42
Guided Exercise: Navigating the Red Hat Advanced Cluster Management Web Console	47
Configuring Access Control for Multicluster Management	53
Guided Exercise: Configuring Access Control for Multicluster Management	57
Lab: Inspecting Resources from Multiple Clusters Using the RHACM Web Console	66
Summary	74
3. Deploying and Managing Policies for Multiple Clusters with RHACM	75
Deploying and Managing Policies with RHACM	76
Guided Exercise: Deploying and Managing Policies with RHACM	79
Deploying and Configuring the Compliance Operator for Multiple clusters using RHACM	85
Guided Exercise: Deploying and Configuring the Compliance Operator for Multiple clusters using RHACM	87
Integrating RHACM with Other Policy Engines	96
Guided Exercise: Integrating RHACM with Other Policy Engines	97
Lab: Deploying and Managing Policies with RHACM	110
Summary	116
4. Configuring the Observability Stack in RHACM	117
Introducing the RHACM Observability Stack	118
Guided Exercise: Introducing the RHACM Observability Stack	120
Customizing the RHACM Observability Stack	128
Guided Exercise: Customizing the RHACM Observability Stack	130
Lab: Installing and Customizing the RHACM Observability Stack	136
Summary	145
5. Managing Applications Across Multiple Clusters with ACM	147
Introducing RHACM GitOps and the Application Model	148
Quiz: Introducing RHACM GitOps and the Application Model	153
Managing Multicluster Application Resources with RHACM	157
Guided Exercise: Managing Multicluster Application Resources with RHACM	166
Customizing Resources with Kustomize for RHACM	175
Guided Exercise: Customizing Resources with Kustomize for RHACM	180
Lab: Managing Applications Across Multiple Clusters with RHACM	192
Summary	199

Document Conventions

This section describes various conventions and practices used throughout all Red Hat Training courses.

Admonitions

Red Hat Training courses use the following admonitions:



References

These describe where to find external documentation relevant to a subject.



Note

These are tips, shortcuts, or alternative approaches to the task at hand. Ignoring a note should have no negative consequences, but you might miss out on something that makes your life easier.



Important

These provide details of information that is easily missed: configuration changes that only apply to the current session, or services that need restarting before an update will apply. Ignoring these admonitions will not cause data loss, but may cause irritation and frustration.



Warning

These should not be ignored. Ignoring these admonitions will most likely cause data loss.

Inclusive Language

Red Hat Training is currently reviewing its use of language in various areas to help remove any potentially offensive terms. This is an ongoing process and requires alignment with the products and services covered in Red Hat Training courses. Red Hat appreciates your patience during this process.

Introduction

DO480: Multicluster Management with Red Hat OpenShift Platform Plus

Multicluster Management with Red Hat OpenShift Platform Plus teaches the skills required to maintain a diverse portfolio of applications, running across a fleet of OpenShift clusters. Applications follow placement rules determined by capacity and criticality; cluster configurations comply with governance and security policies; all automated according to DevOps principles.

Course Objectives

- Deploy Red Hat Advanced Cluster Management (ACM) in a hub cluster.
- Add a managed cluster to ACM (configure a cluster to be managed by ACM).
- Define and apply cluster configuration policies.
- Detect and correct non-conformance to cluster configuration policies.
- Define and apply application placement policies.
- Deploy and update applications from CI/CD and GitOps workflows.
- Deploy Red Hat Quay in the hub cluster.
- Deploy Red Hat Advanced Cluster Security (ACS) in the hub cluster.
- Integrate Quay and ACS security with ACM management.

Audience

- System Administrators, Developers, SREs, and IT Architects interested in managing and automating the management of a fleet of OpenShift clusters, possibly in different data centers and cloud providers.

Prerequisites

- Red Hat Certified Specialist in OpenShift Administration certification (EX280) or equivalent knowledge for the roles of Red Hat OpenShift cluster administrator or SRE.
- Red Hat Certified Systems Administrator certification (EX200) or equivalent knowledge of Linux system administration is recommended for all roles.
- Basic knowledge of Kubernetes and OpenShift administration skills are recommended.

Orientation to the Classroom Environment

Orientation to the Classroom Environment

In this course, the main computer system used for hands-on learning activities is **workstation**.

The **bastion** system must always be running.

Two other machines are also used by students for these activities:

servera, **serverb**, **serverc**, and **serverd**.

All six of these systems are in the `lab.example.com` DNS domain.

Classroom Machines

Machine name	IP addresses
ROLE	<code>bastion.lab.example.com</code>
172.25.250.254	Router that links VMs to central servers
<code>workstation.lab.example.com</code>	172.25.250.9
Graphical workstation for system administration	<code>servera.lab.example.com</code>
172.25.250.10	Managed server "A"
<code>servera.lab.example.com</code>	172.25.250.11
Managed server "B"	<code>servera.lab.example.com</code>
172.25.250.12	Managed server "C"
<code>servera.lab.example.com</code>	172.25.250.13

Performing Lab Exercises

Run the `lab` command from `workstation` to prepare your environment before each hands-on exercise, and again to clean up after an exercise. Each hands-on exercise has a unique name within a course. The exercise is prepended with `lab-` as its file name in `/usr/local/lib`. For example, the `instances-cli` exercise has the file name `/usr/local/lib/lab-instances-cli`. To list the available exercises, use tab completion in the `lab` command:

```
[student@workstation ~]$ lab Tab Tab
administer-users  deploy-overcloud-lab  prep-deploy-ips      stacks-autoscale
analyze-metrics   instances-cli        prep-deploy-router  stacks-deploy
assign-roles       manage-interfaces  public-instance-deploy verify-overcloud
```

Exercises are in two types. The first type, a *guided exercise*, is a practice exercise that follows a course narrative. If a narrative is followed by a quiz, then usually the topic did not have an achievable practice exercise. The second type, an *end-of-chapter lab*, is a gradable exercise to help verify your learning. When a course includes a comprehensive review, the review exercises are structured as gradable labs. The syntax for running an exercise script is:

```
[student@workstation ~]$ lab action exercise
```

The `action` is a choice of `start`, `grade`, or `finish`. All exercises support `start` and `finish`. Only end-of-chapter labs and comprehensive review labs support `grade`. Older courses might still use `setup` and `cleanup` instead of the current `start` and `finish` actions.

start

Formerly `setup`. A script's start logic verifies the required resources to begin an exercise. It might include configuring settings, creating resources, checking prerequisite services, and verifying necessary outcomes from previous exercises. You can take an exercise at any time, even without taking prerequisite exercises.

grade

End-of-chapter labs help verify what you have learned, after practicing with earlier guided exercises. The `grade` action directs the `lab` command to display a list of grading criteria, with a `PASS` or `FAIL` status for each. To achieve a `PASS` status for all criteria, fix the failures and rerun the `grade` action.

finish

Formerly `cleanup`. A script's finish logic deletes exercise resources that are no longer necessary. You can take an exercise as many times as you want.

Troubleshooting Lab Scripts

Exercise scripts do not exist on `workstation` until each is first run. When you run the `lab` command with a valid exercise and action, the script named `lab-exercise` is downloaded from the `classroom` server content share to `/usr/local/lib` on `workstation`. The `lab` command creates two log files in `/var/tmp/labs`, plus the directory if it does not exist. One file, named `exercise`, captures standard output messages that are normally displayed on your terminal. The other file, named `exercise.err`, captures error messages.

```
[student@workstation ~]$ ls -l /usr/local/lib
-rwxr-xr-x. 1 root root 4131 May  9 23:38 lab-instances-cli
-rwxr-xr-x. 1 root root 93461 May  9 23:38 labtool.cl110.shlib
-rwxr-xr-x. 1 root root 10372 May  9 23:38 labtool.shlib

[student@workstation ~]$ ls -l /var/tmp/labs
-rw-r--r--. 1 root root 113 May  9 23:38 instances-cli
-rw-r--r--. 1 root root 113 May  9 23:38 instances-cli.err
```

**Note**

Scripts download from the `http://content.example.com/courses/course/release/grading-scripts` share, but only if the script does not yet exist on `workstation`. When you need to download a script again, such as when a script on the share is modified, manually delete the current exercise script from `/usr/local/lib` on `workstation`, and then run the `lab` command for the exercise again. The newer exercise script then downloads from the `grading-scripts` share.

To delete all current exercise scripts on `workstation`, use the `lab` command's `--refresh` option. A refresh deletes all scripts in `/usr/local/lib` but does not delete the log files.

```
[student@workstation ~]$ lab --refresh
[student@workstation ~]$ ls -l /usr/local/lib

[student@workstation ~]$ ls -l /var/tmp/labs
-rw-r--r--. 1 root root 113 May  9 23:38 instances-cli
-rw-r--r--. 1 root root 113 May  9 23:38 instances-cli.err
```

Interpreting the Exercise Log Files

Exercise scripts send output to log files, even when the scripts are successful. Step header text is added between steps, and additional date and time headers are added at the start of each script run. The exercise log normally contains messages that indicate successful completion of command steps. Therefore, the exercise output log is useful for observing messages that are expected if no problems occur, but offers no additional help when failures occur.

Instead, the exercise error log is more useful for troubleshooting. Even when the scripts succeed, messages are still sent to the exercise error log. For example, a script that verifies that an object already exists before attempting to create it should cause an *object not found* message when the object does not exist yet. In this scenario, that message is expected and does not indicate a failure. Actual failure messages are typically more verbose, and experienced system administrators should recognize common log message entries.

Although exercise scripts are always run from `workstation`, they perform tasks on other systems in the course environment. Many course environments, including OpenStack and OpenShift, use a command-line interface (CLI) that is invoked from `workstation` to communicate with server systems by using API calls. Because script actions typically distribute tasks to multiple systems, additional troubleshooting is necessary to determine where a failed task occurred. Log in to those other systems and use Linux diagnostic skills to read local system log files and determine the root cause of the lab script failure.

Chapter 1

Managing Multicloud Kubernetes Architectures

Goal

Describe multicloud architectures and solve its challenges with Red Hat OpenShift Platform Plus.

Objectives

- Describe common use cases for multiple Kubernetes clusters and examine the challenges of multicloud architectures.
- Identify the tools provided by OpenShift Platform Plus to facilitate end-to-end management, security, and compliance of Kubernetes fleets on hybrid clouds.
- Describe and deploy the RHACM operator from the OperatorHub in an OpenShift hub cluster.

Sections

- Introducing Kubernetes Multicloud and Hybrid Cloud Architectures Landscape (and Quiz)
- Describing Red Hat OpenShift Platform Plus (and Quiz)
- Deploying Red Hat Advanced Cluster Management for Kubernetes (and Guided Exercise)

Lab

Installing the RHACM Operator

Introducing Kubernetes Multicluster and Hybrid Cloud Architectures Landscape

Objectives

After completing this section, you should be able to describe common use cases for multiple Kubernetes clusters and examine the challenges of multicluster architectures.

Interpreting the Background of Kubernetes Multicluster Architectures

Digital transformation is challenging traditional Information Technology (IT) departments and teams.

The wide adoption of the Internet brought new business models to different industries, with new technologies and ways of working. One example is the music and video streaming services that completely changed the entertainment industry. Another example is the growth of online stores that pushed traditional retailers to change how they operated their traditional stores. These disruptions affect other sectors also, such as telecommunications, the automotive business, and financial companies.

Furthermore, these new business models must frequently support operations that serve a huge number of users. An online business with one thousand customers could succeed and grow to hundreds of thousands or millions of users in a few days or weeks.

The following list enumerates some of the challenges that face IT departments:

- Modernizing systems and processes to deliver new software services quickly
- Building new services for customers to adapt to new industry trends and legislation
- Adapting new and existing services to be capable of running on a global scale
- Establishing new methodologies that support the work needed to modernize, build, and adapt their services

Three modern technologies and ways of working are helping IT departments to adapt to digital transformation:

- Cloud computing
- Containers technology
- Automation of processes

Describing Cloud Computing: The Way to the Hybrid Cloud

In IT, a cloud is an environment that abstracts, pools, and shares resources across a network. Cloud computing is the act of running workloads in a cloud environment. The following are some of the main features of a cloud computing IT system:

- You can access all resources in the cloud computing environment through a network.
- The cloud computing environment contains a repository of IT resources, such as applications, software services, and hardware abstractions.
- You can provision and scale the cloud computing environment quickly.

There are different types of clouds, such as public clouds, private clouds, and hybrid clouds. A public cloud provider is a company in charge of maintaining the base hardware, the network, and

the software used to virtualize the servers. Cloud providers also offer service agreements and pay-per-use options for using their services.

The arrival of public cloud providers was a disruptive change in the IT market. Now, an enterprise can run all its software workloads without maintaining on-premises data centers or spending on new hardware.

However, not all companies and workloads are suitable for migration to a public cloud. The following list suggests some reasons to limit the use of public clouds:

- The company already has many critical workloads running in its own data centers.
- The cost of migrating to cloud-native developments is very high.
- Some legacy systems are not adapted to run on cloud environments.
- The operational cost of public cloud services can be high, especially if the enterprise grows rapidly.
- Some verticals, such as healthcare or finance, must comply with important regulatory restrictions on where their data is stored. The General Data Protection Regulation law of the European Union is an excellent example of a data location restriction.
- Concerns about security and losing control of the IT infrastructure are prioritized.

There are many other models of interacting with one or multiple clouds, such as between a fully externalized public provider and a private data center. One model is a fully private cloud.

A private cloud is an IT infrastructure hosted in private data centers whose services are requested and served with a model similar to public clouds:

- Self-provisioning on demand
- Auto-scaling resources based on infrastructure usage
- Better resource allocation and utilization over the same physical resources

Private clouds can be isolated from the external world. Ideally, a private cloud cannot and should not communicate with a public cloud. However, in the real world, IT departments use mixed models. Private data centers, hosted data centers, data centers on public clouds, or data centers on private clouds need to share resources between them. A Hybrid Cloud is an IT infrastructure that can communicate and share services between the following resource types:

- One or more data centers with bare metal or virtualized environments owned by the company
- One or more data centers with bare metal or virtualized environments rented to other companies
- One or more private clouds
- One or more public clouds

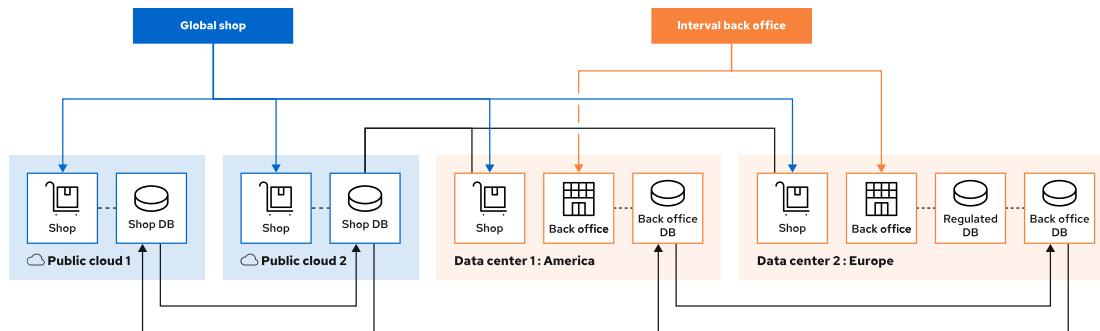
The boundaries between some of these kinds of infrastructures are not always clear, particularly as the technologies evolve. For example, many public cloud providers install some services, and physical resources in some cases, in their customer's private data centers.

Using multiple clouds has the following benefits:

- High availability and failover capabilities: If the primary cloud is down, then other clouds can assume the workload.
- Using specific services from specific clouds: Administrators can use a preferred cloud service from each cloud.
- Geographical proximity: Some services can be faster if they run in close proximity to the user, although this depends on the customer's location.
- Regulatory restrictions: In some cases, data regulatory restrictions can apply.

- Network latency: For some industries like Telco, network traffic latency is a critical concern.

As an example of a hybrid cloud, the following diagram displays the distribution of two software services, an online shop for global access to customers and a back-office application to manage the online shop. The shop service has the benefits of global distribution. Customers always connect to the nearest shop location. The back-office service only runs in the private data centers, whereas the shop service runs in the private data centers and in two public clouds. However, the regulated back-office data is only present in the Europe data center.



Explaining the Need for Containers and Orchestration

As discussed elsewhere in this chapter, software development techniques are evolving quickly. Scale economies and new disruptive players in different industries push to bring new features to their applications as soon as possible. Nevertheless, user applications must still need to be developed with quality, stability, performance, and high availability.

Linux containers are one of the best approaches to providing the infrastructure and development framework to run both cloud-native applications and legacy workloads.

Container technologies use software packaging techniques to make applications fully portable to any hybrid cloud. The packaging and portability of containers make them suitable for every cloud-computing model and in traditional data centers.

The following list describes some of the benefits of using Linux containers:

- Security, storage, and network isolation, similar to virtual machines (VMs)
- Require fewer hardware resources and are quick to start and terminate compared to VMs
- Quick deployment because you do not need to install a complete operating system
- Portable and reusable applications by using environments

In contrast, the use of Linux containers can introduce some of the following new challenges:

- Managing communications between services provided by containers
- Managing versions of running containers and controlling their deployment
- Scaling containers up and down to respond to the real demands of the service
- Determining the location of the containers in the existing infrastructure
- Limiting and balancing available compute resources for containers

You can address these challenges by using a container orchestration system. Kubernetes is an orchestration system that simplifies the deployment, management, and scaling of containerized applications.

Automation Processes

Companies must reduce the time to market for their new services and offerings but keep the ability to scale their infrastructure up or down quickly if demands change.

By using hybrid clouds and Linux containers orchestration systems, companies can cope with changing demands. Moreover, companies must also increase agility and automate their processes at both the technical and organizational levels.

Modern IT teams are adapting their work to adopt agile models that automate a good part of the software development lifecycle. Methods like CI/CD, DevSecOps, and GitOps are difficult to implement, but are the key to an organization's ability to rapidly bring new business value through software development.



Note

To learn more about CI/CD and GitOps see the GLS course D0380 - Red Hat OpenShift Administration III : Scaling Kubernetes Deployments in the Enterprise, chapter 4.

Challenges of Multicloud Architectures

The adoption of automation, containers, and Kubernetes in Hybrid Cloud environments means that companies must find a way to easily create and manage Kubernetes clusters across the hybrid environment.

The following sections discuss the main challenges to adopting a multicloud architecture.

Managing Fleets of Clusters

As software development processes evolve, it might be necessary to use ephemeral Kubernetes clusters. For example, if many DevOps pipelines create a whole cluster to do the quality assurance (QA) end-to-end test, then you need ephemeral Kubernetes clusters. Sometimes, you destroy the cluster after completing the test. Other times, you must keep the cluster temporarily to perform research on it after the tests are done.

When an organization has many development and QA teams belonging to different areas or regions, or in different business units, the number of Kubernetes clusters can grow quickly. The higher the number of clusters, the more difficult it is to locate information about the performance and the status of the fleet.

Managing a high number of clusters is time-consuming and error prone. Administrators must work with multiple consoles and distributed business applications, and sometimes inconsistent security controls across the diverse clusters deployed on premises or in public clouds.

Organizing Compliance and Security of Clusters

IT departments are often under different regulatory compliance requirements. Suppose a company must ensure that one cluster is compliant with standard or custom policies. If it is easy for a team to manually implement the policies on one cluster and ensure that they are correctly applied. On the other hand, manually applying custom policies is impossible if the company has hundreds or thousands of clusters. In such cases, companies need a mechanism to set consistent security policies across diverse environments and to ensure enforcement.

Another concern of IT departments from a security viewpoint is the origin of the software that runs in the Kubernetes clusters. The packaging techniques of containers make it very easy to package software inside a container with security problems or obsolete versions of some dependencies. The developer's freedom to choose the software used in a container has the associated risk of adding layers of defective software to the container image.

Organizing Applications Between Clusters

The placement of the workloads must be automatic and based on the capacity of the clusters, availability zones, or other policies previously defined.

Furthermore, if the applications have dependencies with services in different clusters, having visibility of the network topology is a difficult challenge. Multicloud architectures need repositories of IT artifacts available from different clouds. This need represents another challenge that affects not only the applications but also the deployment of infrastructure in the multicloud environment. In the containers world, many infrastructure components are packaged as container images.

To summarize, the following table relates the new technologies and processes to the challenges that those technologies and new processes present in a multicloud architecture:

Technology or Process	Challenges Introduced by Multicloud Architectures
Cloud computing	<ul style="list-style-type: none"> - Easily provisioning clusters on all types of clouds - Locating information about objects present in the fleet of clusters - Managing configuration compliance and the fleet of cluster's security - Monitoring the behavior of all clusters and their workloads from performance and scalability points of view
Containers	<ul style="list-style-type: none"> - Requires a containers orchestration system - Managing communications between services provided by containers, even when they are in different clusters or clouds - Requires a controlled place to store the artifacts and images that containers use - Security scanning of the artifacts and images that containers use
Automation processes	<ul style="list-style-type: none"> - An easy way for developers to implement CI/CD across a fleet of clusters - An easy way for DevOps teams to implement GitOps across a fleet of clusters



References

Understanding cloud computing

<https://www.redhat.com/en/topics/cloud>

Types of cloud computing

<https://www.redhat.com/en/topics/cloud-computing/public-cloud-vs-private-cloud-and-hybrid-cloud>

What is Hybrid Cloud?

<https://www.redhat.com/en/topics/cloud-computing/what-is-hybrid-cloud>

For more information, refer to the official Documentation of Red Hat Advanced Cluster Management for Kubernetes at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4

► Quiz

Introducing Kubernetes Multicluster and Hybrid Cloud Architectures Landscape

Choose the correct answers to the following questions:

- ▶ 1. **Which two of the following modern technologies are helping IT departments to adapt to digital transformation? (Choose two.)**
 - a. Cloud computing
 - b. Distributed computing
 - c. Linux containers
 - d. Databases
 - e. Private data centers

- ▶ 2. **Which three of the following advantages are a result of using multiple clouds? (Choose three.)**
 - a. Geographical proximity to the end users
 - b. Centralized data sources
 - c. Managing regulatory restrictions per country
 - d. Ability to use specific cloud services from different providers
 - e. High availability and service failover between clouds

- ▶ 3. **Which two of the following challenges arise from using Linux containers? (Choose two.)**
 - a. Managing communication between services running in containers
 - b. Easily updating the server RPM packages
 - c. Allocating resources effectively for the running containers
 - d. Reverting failed operating system upgrades

- ▶ 4. **Which three of the following advantages result from using Linux containers compared to legacy deployments? (Choose three.)**
 - a. Portability and reusability of the applications
 - b. Better resource utilization
 - c. Simpler deployment architecture
 - d. Reduced time to deployment
 - e. Reduced need to keep software versions updated

► 5. **Which four of the following challenges arise from using a Kubernetes multicluster architecture? (Choose four.)**

- a. Managing multiple clusters effectively
- b. Keeping high availability of the running applications
- c. Distributing the applications efficiently among multiple clusters
- d. Ensuring the security and compliance of the fleet of clusters
- e. Monitoring the fleet of clusters from a single pane of glass
- f. Failover when a cloud environment is unavailable

► Solution

Introducing Kubernetes Multicluster and Hybrid Cloud Architectures Landscape

Choose the correct answers to the following questions:

- ▶ 1. **Which two of the following modern technologies are helping IT departments to adapt to digital transformation? (Choose two.)**
 - a. Cloud computing
 - b. Distributed computing
 - c. Linux containers
 - d. Databases
 - e. Private data centers

- ▶ 2. **Which three of the following advantages are a result of using multiple clouds? (Choose three.)**
 - a. Geographical proximity to the end users
 - b. Centralized data sources
 - c. Managing regulatory restrictions per country
 - d. Ability to use specific cloud services from different providers
 - e. High availability and service failover between clouds

- ▶ 3. **Which two of the following challenges arise from using Linux containers? (Choose two.)**
 - a. Managing communication between services running in containers
 - b. Easily updating the server RPM packages
 - c. Allocating resources effectively for the running containers
 - d. Reverting failed operating system upgrades

- ▶ 4. **Which three of the following advantages result from using Linux containers compared to legacy deployments? (Choose three.)**
 - a. Portability and reusability of the applications
 - b. Better resource utilization
 - c. Simpler deployment architecture
 - d. Reduced time to deployment
 - e. Reduced need to keep software versions updated

► 5. **Which four of the following challenges arise from using a Kubernetes multicluster architecture? (Choose four.)**

- a. Managing multiple clusters effectively
- b. Keeping high availability of the running applications
- c. Distributing the applications efficiently among multiple clusters
- d. Ensuring the security and compliance of the fleet of clusters
- e. Monitoring the fleet of clusters from a single pane of glass
- f. Failover when a cloud environment is unavailable

Describing Red Hat OpenShift Platform Plus

Objectives

After completing this section, you should be able to identify the tools provided by OpenShift Platform Plus to facilitate end-to-end management, security, and compliance of Kubernetes fleets on hybrid clouds.

Describing Red Hat OpenShift Platform Plus

Red Hat OpenShift Container Platform (RHOC) is a certified enterprise Kubernetes distribution. It runs on a wide variety of infrastructure providers: public and private clouds, virtualization software, and bare metal servers. OpenShift Container Platform allows administrators to build a portable, production ready, and hybrid application development container platform.

Red Hat OpenShift Platform Plus is a collection of software tools running on Red Hat OpenShift Container Platform. Red Hat OpenShift Platform Plus allows administrators to manage the software life cycle across Kubernetes and RHOC clusters. It provides end-to-end visibility and control over multiple Kubernetes clusters from a single pane of glass. Furthermore, OpenShift Platform Plus brings Kubernetes-native security capabilities to protect the software supply chain, infrastructure, and workloads. Finally, OpenShift Platform Plus also provides a globally-distributed and scalable registry.

The capabilities mentioned previously are provided by the following Red Hat products, included in the Red Hat OpenShift Platform Plus subscription, and distributed and installed as Kubernetes operators in OpenShift Container Platform.

Red Hat Advanced Cluster Management for Kubernetes

Red Hat Advanced Cluster Management for Kubernetes provides multi cluster and application control, along with built-in security policies.

By using Red Hat Advanced Cluster Management for Kubernetes, administrators can perform the following tasks:

- Create, update, or delete RHOC or Kubernetes clusters across multiple private and public clouds.
- Find and modify any Kubernetes resource across the entire domain using the built-in search engine.
- Troubleshoot and resolve issues using the built-in search engine.
- Automate and apply tasks in the clusters through the integration with Red Hat Ansible Automation Platform.
- Enforce compliance policies across the fleet of managed clusters.
- Manage lifecycle of applications across datacenters and hybrid clouds.
- Visualize cluster metrics from a centralized monitoring stack.
- Create and manage alerts generated across all the managed clusters.

Red Hat Advanced Cluster Security for Kubernetes

Powered by StackRox technology, Red Hat Advanced Cluster Security for Kubernetes helps administrators to enforce DevOps and security best practices.

By using Red Hat Advanced Cluster Security for Kubernetes, administrators get the following advantages:

- Visibility – Centralized view of your deployments, traffic in all clusters, and critical system level events in each running container.
- Vulnerability management – Image vulnerability scanning, vulnerability correlation to running deployments, and policy enforcing at build, deploy, and run time.
- Compliance – Assessment for CIS Benchmarks, payment card industry (PCI), Health Insurance Portability and Accountability Act (HIPAA), and NIST SP 800-190. Centralized compliance dashboard with evidence export for auditors, and detailed view compliance to pinpoint specific clusters, nodes, or namespaces.
- Network segmentation – Visualization of allowed and active traffic, simulation of network policy changes, network policy recommendations, and network enforcement capabilities.
- Risk profiling – Deployment ranking based on security risk calculation, and security tracking to validate changes in configuration.
- Configuration management – Pre-built DevOps and security policies, Kubernetes role-based access control (RBAC) analysis, audit control for access to Kubernetes Secrets, and configuration policies at build time for CI/CD integration and deploy time.
- Runtime detection and response – System level event monitoring, automatic allowing of process activity, pre-built policies to detect crypto mining, privilege escalation, and various exploits, and system level data collection using external Berkeley Packet Filter (eBPF) or other Linux kernel modules.
- Integration – API and prebuilt plugins to integrate with DevOps systems, CI/CD tools, image scanners, registries, container runtimes, security integration event management (SIEM) solutions, and notification tools.

Red Hat Quay

Red Hat Quay is a scalable container registry to serve as a single source of truth for container images. Administrators can use Red Hat Quay to automate container builds with integration with GitHub, Bitbucket, and more. Red Hat Quay also provides vulnerability scanning for the container images.

By using Red Hat Quay, administrators get the following advantages:

- Time machine, with the ability to rollback image changes to a previous state.
- Geographic replication, for improved performance and availability.
- Continuous garbage collection, removing unused images automatically.
- Advanced access control management, including support for mapping teams and organizations integrating with an existing identity infrastructure.
- TLS security encryption between Red Hat Quay and your servers.
- Security scanning integration with vulnerability detectors like Clair [<https://www.redhat.com/en/topics/containers/what-is-clair>].

- Notifications for alerting about detected vulnerabilities.
- Integrating with CI/CD pipelines using build triggers, git hooks, and robot accounts.
- Auditing of CI and API actions.
- Integrated Prometheus metrics export.

Challenges Addressed by Red Hat OpenShift Platform Plus

The following table relates the challenges of multicloud architectures, as discussed in the *Challenges of Multicloud Architectures* section, to the Red Hat OpenShift Platform Plus software collection.

Challenges Introduced by Multicloud Architectures	Red Hat OpenShift Platform Plus software addressing it
Easily provisioning clusters on all types of clouds	Red Hat Advanced Cluster Management for Kubernetes
Locating information about objects present in a fleet of clusters	Red Hat Advanced Cluster Management for Kubernetes search engine
Managing configuration compliance and the fleet of cluster's security	Red Hat Advanced Cluster Management for Kubernetes governance engines
Monitoring the behavior of all clusters and their workloads from performance and scalability points of view	Red Hat Advanced Cluster Management for Kubernetes observability engine
Requires a containers orchestration system	Red Hat OpenShift Container Platform or Kubernetes
Managing communications between services provided by containers, even when they are in different clusters or clouds	Red Hat Advanced Cluster Management for Kubernetes Submariner services
Requires a controlled place to store the artifacts and images that containers use	Red Hat Quay
Security scanning of the artifacts and images that containers use	Red Hat Quay for static scanning and Red Hat Advanced Cluster Security for Kubernetes for dynamic scanning
Automation processes	Red Hat Advanced Cluster Management for Kubernetes and its integration with the Ansible Automation Platform Resource Operator from Red Hat OpenShift Container Platform
An easy way for developers to implement CI/CD across a fleet of clusters	Red Hat Advanced Cluster Management for Kubernetes applications engine
An easy way for DevOps teams to implement GitOps across a fleet of clusters	Red Hat Advanced Cluster Management for Kubernetes applications engine



References

Red Hat OpenShift Container Platform

<https://www.redhat.com/en/technologies/cloud-computing/openshift/container-platform>

Red Hat OpenShift Platform Plus

<https://www.redhat.com/en/technologies/cloud-computing/openshift/platform-plus>

Red Hat Advanced Cluster Management for Kubernetes

<https://www.redhat.com/en/technologies/management/advanced-cluster-management>

Red Hat Advanced Cluster Security for Kubernetes

<https://www.redhat.com/en/resources/advanced-cluster-security-for-kubernetes-datasheet>

Red Hat Quay

<https://www.redhat.com/en/technologies/cloud-computing/quay>

Red Hat Quay datasheet: Private container registry

<https://www.redhat.com/en/resources/quay-datasheet>

► Quiz

Describing Red Hat OpenShift Platform Plus

Choose the correct answers to the following questions:

- ▶ 1. **Which three of the following products include Red Hat OpenShift Platform Plus? (Choose three.)**
 - a. Red Hat Advanced Cluster Management for Kubernetes
 - b. Red Hat Satellite
 - c. Red Hat Quay
 - d. Red Hat Advanced Cluster Security for Kubernetes
 - e. Red Hat Identity Management

- ▶ 2. **Which three of the following tasks can administrators perform using Red Hat Advanced Cluster Management for Kubernetes? (Choose three.)**
 - a. Create, update, or delete RHOCP clusters across multiple private and public clouds.
 - b. Create virtual networks to connect clusters across multiple private and public clouds.
 - c. Enforce compliance policies across a fleet of managed clusters.
 - d. Create custom ISO files to deploy Red Hat Enterprise Linux CoreOS (RHCOS).
 - e. Visualize cluster metrics from a centralized monitoring stack.

- ▶ 3. **Which three of the following advantages arise from using Red Hat Quay? (Choose three.)**
 - a. Integrate with CI/CD pipelines using build triggers, git hooks, and robot accounts.
 - b. Preserve a private source code control system for infrastructure as a service (IaaS) repositories.
 - c. Rollback container image changes to a previous state.
 - d. Receive alerts about detected vulnerabilities in the hosted container images.
 - e. Deploy RHOCP clusters from the Red Hat Quay dashboard.

- ▶ 4. **Which four of the following advantages arise from using Red Hat Advanced Cluster Security for Kubernetes? (Choose four.)**
 - a. Assessment for CIS benchmarks
 - b. Simulation of changes in the network policies
 - c. Configuration policies at build time for CI/CD integration
 - d. Automatic application updates to avoid recent vulnerabilities
 - e. Centralized vulnerability management with correlation to running deployments
 - f. Automatic cluster updates to avoid recent vulnerabilities

► Solution

Describing Red Hat OpenShift Platform Plus

Choose the correct answers to the following questions:

- ▶ 1. **Which three of the following products include Red Hat OpenShift Platform Plus? (Choose three.)**
 - a. Red Hat Advanced Cluster Management for Kubernetes
 - b. Red Hat Satellite
 - c. Red Hat Quay
 - d. Red Hat Advanced Cluster Security for Kubernetes
 - e. Red Hat Identity Management

- ▶ 2. **Which three of the following tasks can administrators perform using Red Hat Advanced Cluster Management for Kubernetes? (Choose three.)**
 - a. Create, update, or delete RHOPC clusters across multiple private and public clouds.
 - b. Create virtual networks to connect clusters across multiple private and public clouds.
 - c. Enforce compliance policies across a fleet of managed clusters.
 - d. Create custom ISO files to deploy Red Hat Enterprise Linux CoreOS (RHCOS).
 - e. Visualize cluster metrics from a centralized monitoring stack.

- ▶ 3. **Which three of the following advantages arise from using Red Hat Quay? (Choose three.)**
 - a. Integrate with CI/CD pipelines using build triggers, git hooks, and robot accounts.
 - b. Preserve a private source code control system for infrastructure as a service (IaaS) repositories.
 - c. Rollback container image changes to a previous state.
 - d. Receive alerts about detected vulnerabilities in the hosted container images.
 - e. Deploy RHOPC clusters from the Red Hat Quay dashboard.

- ▶ 4. **Which four of the following advantages arise from using Red Hat Advanced Cluster Security for Kubernetes? (Choose four.)**
 - a. Assessment for CIS benchmarks
 - b. Simulation of changes in the network policies
 - c. Configuration policies at build time for CI/CD integration
 - d. Automatic application updates to avoid recent vulnerabilities
 - e. Centralized vulnerability management with correlation to running deployments
 - f. Automatic cluster updates to avoid recent vulnerabilities

Deploying Red Hat Advanced Cluster Management for Kubernetes

Objectives

After completing this section, you should be able to describe and deploy the RHACM operator from the OperatorHub in an OpenShift hub cluster.

Introducing Red Hat Advanced Cluster Management for Kubernetes

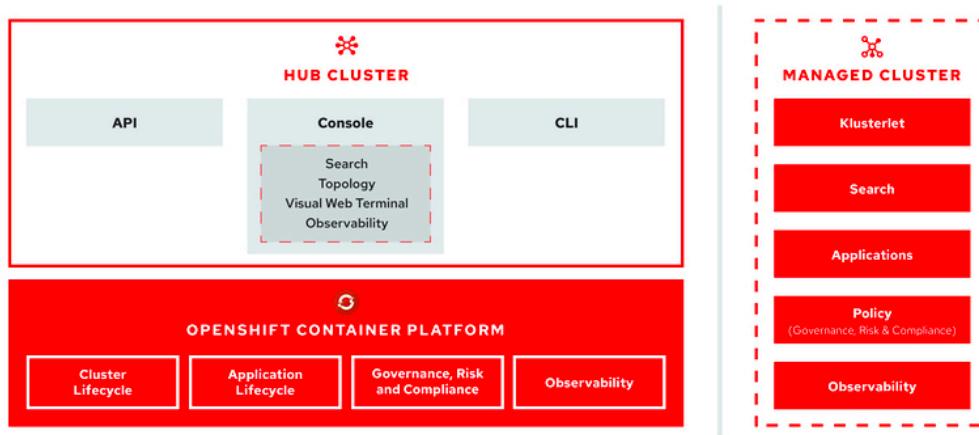
Red Hat Advanced Cluster Management for Kubernetes (RHACM) provides visibility of your OpenShift Container Platform and Kubernetes clusters from a view. RHACM provides built-in governance, cluster lifecycle management, and application lifecycle management, and includes observability features. RHACM is cloud-native by design. It runs on containers, does not use relational persistence, and manages Kubernetes objects directly through their APIs.

The core components of RHACM run in the hub cluster, the central controller that runs the RHACM management APIs, the RHACM web console, and an integrated command-line tool. The `MultiClusterHub` operator, part of the RHACM operator installation, is responsible of managing, upgrading, and installing hub cluster components.

In RHACM, the local cluster is the RHOCP cluster hosting the hub cluster. The hub cluster is also a managed cluster itself unless specified otherwise during the creation of the `MultiClusterHub` object. A managed cluster is an additional cluster managed by the hub cluster.

The agent controlling the connection between clusters is the `klusterlet`. The hub cluster can control managed clusters communicating with the `klusterlet` running in them. The operator creates all the necessary custom resource definitions (CRDs) for the RHACM components.

The following diagram provides a high level overview of an RHACM deployment on RHOCP.



As displayed in the previous diagram, there are three ways to interact with RHACM:

- The RHACM API

- The RHACM web console
- CLI tools such as `oc` and `kubectl`

The web console provides visual, simplified access to the search engine, topology view, visual web terminal, and the observability console among other features.

Apart from the core components running in the hub cluster, every managed cluster runs agents to communicate with the hub cluster. These agents, also called addons are:

- the `klusterlet`
- the search engine addon
- the addon for managing application deployments from RHACM
- the addon for applying policies
- the addon for the observability components

RHACM leverages capabilities from the RHOCP local cluster for cluster lifecycle management, application lifecycle management, applying governance, risk, and compliance policies, and the prometheus and grafana instances running on it for observability.

Installing RHACM

Red Hat recommends administrators to install the Advanced Cluster Management for Kubernetes operator from the **OperatorHub** menu in the RHOCP web console.



Note

A list of supported hub cluster providers is available in the RHACM documentation [https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters-supported-clouds#supported-hub-cluster-providers].

The installation requires selecting an **Update channel**. An update channel is used to deliver streams of updates to the operators. Typically, update channels correspond to minor software versions. The installation also requires to choose an **Update approval** strategy, between automatic or manual. Choosing the manual update approval approach requires human intervention to update when a new version of the operator is available in the update stream. Alternatively, the automatic update approval will trigger an update when a new operator version is available.

After the operator is installed, administrators must create the **MultiClusterHub** custom resource to deploy all the RHACM components. You will be required to perform this steps elsewhere in this course.

Administrators can also install the RHACM operator using the CLI. Both installation methods can be performed in connected and disconnected environments.



Note

For more information about all the available installation methods, check the *Red Hat Advanced Cluster Management for Kubernetes Install Guide* at https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/install/index

When the **MultiClusterHub** finishes deploying, it creates the **multicloud-console** route to access the RHACM web console. Administrators can find that route using the RHOCP web console, in the **Networking → Routes** menu. Alternatively, administrators can use the following command:

```
[user@demo ~]$ oc get routes -n open-cluster-management
NAME           HOST/PORT ...
multicloud-console multicloud-console.apps.ocp4.example.com
```

Importing Existing Clusters to RHACM

When the **MultiClusterHub** custom resource installation completes, the next step is to aggregate or import more Kubernetes or RHOCP clusters into RHACM.



Note

A list of supported managed cluster providers is available in the RHACM documentation [https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/supported-clouds#supported-managed-cluster-providers].

The local cluster (the RHOCP cluster where the hub cluster runs) is automatically imported into RHACM, unless specified during the creation of the **MultiClusterHub** object.

There are two ways for importing new clusters: from the web console, and using the CLI.

The following prerequisites are common for both import methods:

1. Internet connectivity between the hub and the managed cluster
2. The `oc` or `kubectl` CLI tools
3. The `base64` utility
4. In case of importing a Kubernetes cluster, a `multicloudhub.spec.imagePullSecret` defined containing the `cloud.redhat.com` pull secret
5. In case of importing a Red Hat OpenShift Dedicated cluster, the hub cluster running in it, and `cluster-admin` permissions to use RHACM.

Importing Existing Clusters Using the RHACM Web Console

To import an existing cluster from the RHACM web console, administrators must follow these steps:

1. From the RHACM web console, navigate to the **Infrastructure → Clusters** menu.
2. Click **Add a cluster**.
3. Click **Import an existing cluster**. Provide a name for it.
4. Add any desired **Additional labels** (optional).
5. Click **Save import and generate code**.
6. Click **Copy command** to copy the generated command and token to the clipboard.
7. Use the `oc` CLI tool to log into the cluster you want to import.

8. Paste and run the copied command.
9. Return to the **Infrastructure → Clusters** menu.

The newly imported cluster appears immediately in the **Infrastructure → Clusters** menu. It takes a while to see all the cluster details. This is because RHACM is installing all the add-ons in the managed cluster.



Note

For detailed steps and expanded information, review the *Chapter 8. Importing a target managed cluster to the hub cluster* section of the *Red Hat Advanced Cluster Management for Kubernetes Clusters Guide* at https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/index#importing-a-target-managed-cluster-to-the-hub-cluster

Importing Existing Clusters Using the CLI

The steps for importing an existing cluster using the CLI are different from the ones using the web console.

The following steps are a summary of the actions to be performed in each cluster, for a better understanding of the process.



Note

For detailed instructions on how to import clusters into RHACM using the CLI, visit https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/index#importing-a-managed-cluster-with-the-cli

In the hub cluster:

1. Create a project with the name of the cluster to import, for instance: `imported-cluster`
2. Add the label `cluster.open-cluster-management.io/managedCluster=imported-cluster` to the `imported-cluster` namespace
3. Create a new `ManagedCluster` CR with the name `imported-cluster`
4. Create a new `KlusterletAddonConfig` CR with klusterlet configuration file.

In the hub cluster, administrators must extract existing secrets and generate YAML files containing:

1. The `klusterlet` CRD. For instance in a file named `klusterlet-crd.yaml`.
2. The secret for importing new clusters. For instance in a file named `import.yaml`.

Then, in the managed cluster to import, administrators must use the YAML files generated in the previous step to create the necessary objects.

1. `oc create -f klusterlet-crd.yaml`
2. `oc create -f import.yaml`

Finally, validate the status of the imported cluster.

In the hub cluster, validate the JOINED and AVAILABLE status of the `managedcluster` named `imported-cluster`.

In the target cluster, validate the status of the pods of all the addons running on the target cluster namespace `open-cluster-management-agent-addon`.

Removing Imported Clusters from RHACM

To remove an imported cluster from the RHACM console, navigate to the **Clusters** page and click **Actions** → **Detach cluster** from the options of the cluster to remove.

To remove an imported cluster using the CLI, log in to the hub cluster, locate the correspondent `managedcluster` object and use `oc` to delete it. Because RHACM is removing all the addons and other RHACM components from the managed cluster, the detach process takes a few minutes.



References

For more information, refer to the *About* guide in the *Red Hat Advanced Cluster Management for Kubernetes* documentation at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/about/index

For more information, refer to the *Install* guide in the *Red Hat Advanced Cluster Management for Kubernetes* documentation at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/install/index

For more information, refer to the *Clusters* guide in the *Red Hat Advanced Cluster Management for Kubernetes* documentation at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/index

► Guided Exercise

Deploying Red Hat Advanced Cluster Management for Kubernetes

In this exercise you will install the RHACM operator and import an existing cluster.

Then, you will remove the imported cluster and uninstall the RHACM operator and all its components.

Outcomes

You should be able to:

- Install the RHACM operator from OperatorHub
- Create the `MultiClusterHub` object
- Import an existing cluster into RHACM
- Remove an imported cluster
- Uninstall the RHACM operator and all its components

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the RHACM operator is not already installed.

```
[student@workstation ~]$ lab start multicloud-acm
```



Note

The hub cluster must be installed in the cluster `ocp4.example.com`.

Instructions

- 1. Using OperatorHub, install the Advanced Cluster Management for Kubernetes operator in the `ocp4.example.com` cluster. The web console URL is <https://console-openshift-console.apps.ocp4.example.com>.
- 1.1. From the `workstation` machine, use Firefox to navigate to the RHOCP 4 web console at <https://console-openshift-console.apps.ocp4.example.com>.
When prompted, select `htpasswd_provider`
Type the credentials:
- Username: `admin`
 - Password: `redhat`

Click Log in.

- 1.2. Install the Advanced Cluster Management for Kubernetes operator from OperatorHub.

Navigate to **Operators → OperatorHub** and type **Advanced Cluster Management** in the **Filter by keyword** text field.

The screenshot shows the Red Hat OpenShift Container Platform interface. On the left, there's a sidebar with 'Administrator' and 'Home' sections, followed by an 'Operators' section with 'OperatorHub' selected. The main content area is titled 'OperatorHub' and displays a search result for 'Advanced Cluster Management'. The result shows a card with a red icon, the text 'Advanced Cluster Management for Kubernetes provided by Red Hat', and a link to 'https://console-openshift-console.apps.ocp4.example.com/#operator/...'. There are also tabs for 'All Items' and 'Advanced Cluster Management'.

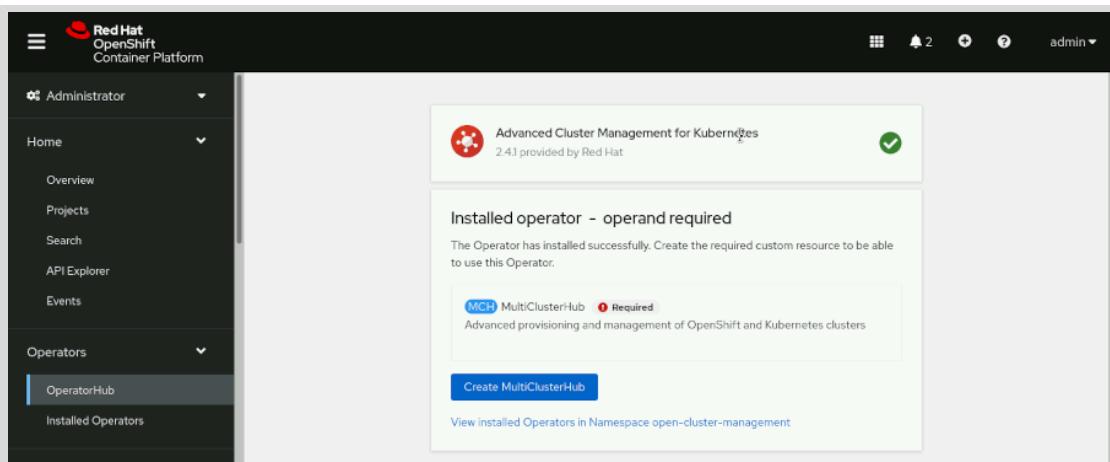
Click **Advanced Cluster Management for Kubernetes**, and then click **Install**.

In the **Update Channel**, ensure that the **release-2.4** radio button is selected. Select **Manual** as the **Approval strategy**. Then, click **Install**.

The screenshot shows the 'Install' dialog for the 'Advanced Cluster Management for Kubernetes' operator. It has fields for 'Select a Namespace' (set to 'default'), 'Update approval' (set to 'Manual'), and 'App Subscription' (disabled). A note explains that manual approval applies to all operators in the namespace. At the bottom are 'Install' and 'Cancel' buttons.

Next, you must approve the installation or updates to the RHACM operator manually. Click **Approve** in the next step. The installation can take a few minutes to complete.

When the operator is installed, you see the following message:



Click **Create MultiClusterHub**.

On the **Create MulticlusterHub** page, leave the default values and click **Create**. You will be redirected to the **MultiClusterHubs** tab.

Initially, the **Status** of the **multicloudhub** object will be **Phase: Installing**. After some minutes, the **Status** column will show **Phase: Running**.

► 2. Verify the deployment of the MultiClusterHub.

- 2.1. Open a terminal and log in to the `ocp4.example.com` cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 2.2. Check the status of all the objects in the `open-cluster-management` namespace.

```
[student@workstation ~]$ oc get all -n open-cluster-management
...output omitted...
[student@workstation ~]$
```

The output is quite long. Review the status of the pods, services, deployments, replica sets, and stateful sets.

- 2.3. Retrieve the route to the RHACM web console, named `multicloud-console`.

```
[student@workstation ~]$ oc get route multicloud-console -n
open-cluster-management
NAME          HOST/PORT      ...
multicloud-console  multicloud-console.apps.ocp4.example.com
```

Copy the URL `multicloud-console.apps.ocp4.example.com` for the next step.

► 3. Test the RHACM deployment by importing the `ocp4-mng.example.com` cluster from the RHACM web console.

- 3.1. In Firefox, open a new tab and log in to the RHACM web console with the URL `multicloud-console.apps.ocp4.example.com`.

When prompted, select `htpasswd_provider`.

Type the credentials:

- User name: `admin`
- Password: `redhat`

Click Log in.

- 3.2. Explore the **Infrastructure → Clusters** menu.

From the **Infrastructure** menu, click **Clusters**. Scroll down to locate the managed clusters. Notice that the `local-cluster ocp4.example.com` cluster, where the hub cluster runs, is automatically managed.

- 3.3. Generate the code to import an existing cluster.

Click **Import Cluster**

In the Import an existing cluster screen, indicate:

- **Name: managed-cluster**

Leave the rest of the values unchanged and click **Save import and generate code**.

The **Save import and generate code** button now displays the **Code generated successfully** message.

Click **Copy command**.

**Note**

The previous steps copies a command to the clipboard to be executed in a terminal after logging in the managed cluster `ocp4-mng.example.com`.

- 3.4. From the terminal, log in to the `ocp4-mng.example.com` cluster as the `admin` user. The API Server URL is `https://api.ocp4-mng.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat https://api.ocp4-
mng.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 3.5. Paste and run the import code in to the terminal and press `Enter`. You can use the keys `Ctrl+Shift+V` or your preferred method for pasting content from the clipboard. The paste command is quite long, and most of it is encoded in base64.

```
[student@workstation ~]$ echo "Ci0tLQph.....G9ydCBhZ2Fpbj4=" | base64 -d
customresourcedefinition.apiextensions.k8s.io/klusterlets.operator.open-cluster-
management.io created
namespace/open-cluster-management-agent created
serviceaccount/klusterlet created
secret/bootstrap-hub-kubeconfig created
clusterrole.rbac.authorization.k8s.io/klusterlet created
clusterrole.rbac.authorization.k8s.io/open-cluster-management:klusterlet-admin-
aggregate-clusterrole created
clusterrolebinding.rbac.authorization.k8s.io/klusterlet created
deployment.apps/klusterlet created
klusterlet.operator.open-cluster-management.io/klusterlet created
[student@workstation ~]$
```

- 3.6. Verify the import of the managed cluster.

Return to the RHACM web console in Firefox. Navigate to **Infrastructure → Clusters**. The `managed-cluster` is now listed in the **Managed clusters** list.

Name	Status	Infrastructure provider	Distribution version	Labels	Nodes
local-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	13 labels	3
managed-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	11 labels	3

The **managed-cluster** appears in the RHACM web console almost immediately after executing the import code. Initially, you will see limited information about the cluster. After a few minutes, all the details of the cluster will become available.

► **4.** Remove the imported cluster.

4.1. From the RHACM web console, detach the imported cluster.

In the **Infrastructure → Clusters** pane, click **Options** button on the right side of the **managed-cluster**. Then, click **Detach cluster**.

Then, type the cluster name **managed-cluster** in the text field and click **Detach**.

The **Status** column changes to **Detaching**. After a few minutes, the **managed-cluster** is completely detached.

► **5.** Uninstall the RHACM operator and all of its components.

5.1. Delete the **multicloudhub** object.

From the RHOCP web console, navigate to **Operators → Installed Operators** and click **MultiCloudHub** in the **Provided APIs** column.

Name	Managed Namespaces	Status	Provided APIs
Advanced Cluster Management for Kubernetes	open-cluster-management	Succeeded Up to date	MultiClusterHub ClusterManager multicloudhub.operator.open-cluster-management.io

Click **Options** to the right side of the **multicloudhub** object. Then, click **Delete MultiClusterHub**.

Name	Kind	Status	Labels
multicloudhub	MultiClusterHub	Phase: Running	No labels

When prompted, click **Delete** to confirm. The **Status** column changes to **Phase: Uninstalling**. After a few minutes, the **multicloudhub** object is completely removed and the page displays the **No operands found** message.



Note

It takes several minutes to remove the **multicloudhub** object. Do not continue with the next step until the **multicloudhub** object is completely removed.

5.2. Uninstall the Advanced Cluster Management for Kubernetes operator.

From the RHOCP web console, navigate to **Operators** → **Installed Operators**. Click **Options** on the right side of the **Advanced Cluster Management for Kubernetes** operator. Then, click **Uninstall Operator**.

The screenshot shows the Red Hat OpenShift Container Platform interface. On the left, there's a sidebar with navigation links like Home, Overview, Projects, Search, API Explorer, Events, Operators, OperatorHub, and Installed Operators. The 'Installed Operators' link is currently selected. The main content area is titled 'Installed Operators' and contains a table with one row. The table columns are Name, Managed Namespaces, Status, and Provided APIs. The single operator listed is 'Advanced Cluster Management for Kubernetes', which is managed by the 'open-cluster-management' namespace, has a status of 'Succeeded' and 'Up to date', and provides APIs for 'MultiClusterHub', 'ClusterManager', 'App Subscription Channel', and 'View 19 more...'. A context menu is open for this operator, with options 'Edit Subscription' and 'Uninstall Operator' visible.

When prompted, click **Uninstall** to confirm. After a few minutes, the operator is completely uninstalled and the page displays the `No Operators found` message.

5.3. Remove the `open-cluster-management` namespace.

From the terminal, run the following command to remove the namespace and any remaining objects in the `open-cluster-management` namespace:

```
[student@workstation ~]$ oc delete project open-cluster-management
project.project.openshift.io "open-cluster-management" deleted
[student@workstation ~]$
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish multicloud-acm
```

This concludes the guided exercise.

► Lab

Installing the RHACM Operator

In this lab, you will install RHACM and import an existing cluster in the lab environment using the command-line interface.

Outcomes

You should be able to perform the following tasks by using the `oc` command-line tool:

- Install the RHACM operator.
- Create the `MultiClusterHub` object.
- Import an existing cluster into RHACM.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the RHACM operator is not already installed.

```
[student@workstation ~]$ lab start multicloud-review
```



Note

The hub cluster must be installed in the cluster `ocp4.example.com`.

The managed-cluster URL is `ocp4-mng.example.com`.

Instructions

1. Install the Advanced Cluster Management for Kubernetes operator on the `ocp4.example.com` cluster.
2. Create the RHACM `MultiClusterHub` object.
3. Prepare RHACM to import a cluster using the name `managed-cluster`.
4. Generate the files to import the `managed-cluster` into RHACM.
5. Import the RHCOP cluster, `ocp4-mng.example.com`, available in the lab environment into RHACM. Use `managed-cluster` as the name of the cluster to import.
6. Log back in to the hub cluster and verify the status of the `managed-cluster`.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade multicloud-review
```

Finish

Do not make any other changes to the lab environment until the next guided exercise. You will continue using this environment in upcoming exercises.

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise.

```
[student@workstation ~]$ lab finish multicloud-review
```

This concludes the lab.

► Solution

Installing the RHACM Operator

In this lab, you will install RHACM and import an existing cluster in the lab environment using the command-line interface.

Outcomes

You should be able to perform the following tasks by using the `oc` command-line tool:

- Install the RHACM operator.
- Create the `MultiClusterHub` object.
- Import an existing cluster into RHACM.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the RHACM operator is not already installed.

```
[student@workstation ~]$ lab start multicloud-review
```



Note

The hub cluster must be installed in the cluster `ocp4.example.com`.

The managed-cluster URL is `ocp4-mng.example.com`.

Instructions

1. Install the Advanced Cluster Management for Kubernetes operator on the `ocp4.example.com` cluster.
 - 1.1. From the `workstation` machine, open a terminal and log in to the `ocp4` cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
...output omitted...
[student@workstation ~]$
```

- 1.2. Create a project named `open-cluster-management`.

```
[student@workstation ~]$ oc new-project open-cluster-management
Now using project "open-cluster-management" on server "https://
api.ocp4.example.com:6443".
...output omitted...
[student@workstation ~]$
```

- 1.3. Create an operator group for the `open-cluster-management` namespace. First, create a file named `operator-group.yaml` with the following contents:

```
[student@workstation ~]$ vi operator-group.yaml
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: acm-operator-group
spec:
  targetNamespaces:
    - open-cluster-management
```

Save the changes and close the editor.

Then, create the `OperatorGroup` object by using the `oc` command.

```
[student@workstation ~]$ oc create -f operator-group.yaml
operatorgroupoperators.coreos.com/acm-operator-group created
[student@workstation ~]$
```

- 1.4. Create a subscription to the Advanced Cluster Management for Kubernetes operator.
First, create a file named `subscription.yaml` with the following contents:

```
[student@workstation ~]$ vi subscription.yaml
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: acm-operator-subscription
spec:
  sourceNamespace: openshift-marketplace
  source: redhat-operators
  channel: release-2.4
  installPlanApproval: Manual
  name: advanced-cluster-management
```

Save the changes and close the editor.

Then, create the `Subscription` object using the `oc` command.

```
[student@workstation ~]$ oc create -f subscription.yaml
subscriptionoperators.coreos.com/acm-operator-subscription created
```

- 1.5. Approve the installation of the operator.
Because you created the subscription with the `Manual` approval strategy, you must approve the installation manually.
First, retrieve the installation plans.

```
[student@workstation ~]$ oc get installplan
NAME          CSV                         APPROVAL   APPROVED
install-4k2q8 advanced-cluster-management.v2.4.1   Manual    false
```

As you can see in the preceding output, the installation is not approved. Approve the installation plan *install-4k2q8* by running the following command.

```
[student@workstation ~]$ oc patch installplan install-4k2q8 --type merge \
--patch '[{"spec":{"approved":true}}]'
installplan.operator.coreos.com/install-4k2q8 patched
```

After the installation plan is approved, the cluster creates a `ClusterServiceVersion` object.

- 1.6. Watch the status of the available `ClusterServiceVersion` objects to verify the status of the installation. Wait until the `PHASE` column shows the message: `Succeeded`.

```
[student@workstation ~]$ watch oc get csv
...output omitted...
NAME          DISPLAY
VERSION ... PHASE
advanced-cluster-management.v2.4.1  Advanced Cluster Management for Kubernetes
2.4.1      ... Succeeded
```

Exit the `watch` command wby pressing `Ctrl+C`.

When the `PHASE` column shows the message: `Succeeded`, the installation of the operator is complete.

2. Create the RHACM `MultiClusterHub` object.

- 2.1. From the terminal, create a file named `mch.yaml` containing the definition of the `MultiClusterHub` object.

```
[student@workstation ~]$ vi mch.yaml
apiVersion: operator.open-cluster-management.io/v1
kind: MultiClusterHub
metadata:
  name: multicluseterhub
  namespace: open-cluster-management
spec: {}
```

Save the changes and close the editor.

- 2.2. Use the `oc` command to create the `MultiClusterHub` object from the `mch.yaml` file.

```
[student@workstation ~]$ oc create -f mch.yaml
multicluseterhub.operator.open-cluster-management.io/multicluseterhub created
```



Note

It takes around 5 minutes to create all the resources of the MultiClusterHub object.

- 2.3. Wait until the MultiClusterHub object creates all its components. Use the `watch` command to monitor the status.

```
[student@workstation ~]$ watch oc get multiclusterhub  
...output omitted...  
NAME STATUS AGE  
multicloudhub Installing 3m21s
```

When the STATUS column displays `Running`, exit the `watch` command with `Ctrl+C`.

The creation of the `MultiClusterHub` object is complete.

3. Prepare RHACM to import a cluster using the name `managed-cluster`.

- 3.1. In the hub cluster, create the `managed-cluster` namespace.

```
[student@workstation ~]$ oc new-project managed-cluster  
...output omitted...
```

- 3.2. Label the namespace with the cluster name to be used by RHACM.

```
[student@workstation ~]$ oc label namespace managed-cluster \  
cluster.open-cluster-management.io/manage dCluster=managed-cluster  
namespace/managed-cluster labeled  
...output omitted...
```

- 3.3. Create a file named `mngcluster.yaml` with the following contents:

```
[student@workstation ~]$ vi mngcluster.yaml  
apiVersion: cluster.open-cluster-management.io/v1  
kind: ManagedCluster  
metadata:  
  name: managed-cluster  
spec:  
  hubAcceptsClient: true
```

Save the changes and close the editor.

- 3.4. Create the `ManagedCluster` object using the previous file.

```
[student@workstation ~]$ oc create -f mngcluster.yaml  
managedcluster.cluster.open-cluster-management.io/managed-cluster created
```

- 3.5. Create a `klusterlet` add-on configuration file named `klusterlet.yaml` with the following contents::

```
[student@workstation ~]$ vi klusterlet.yaml
apiVersion: agent.open-cluster-management.io/v1
kind: KlusterletAddonConfig
metadata:
  name: managed-cluster
  namespace: managed-cluster
spec:
  clusterName: managed-cluster
  clusterNamespace: managed-cluster
  applicationManager:
    enabled: true
  certPolicyController:
    enabled: true
  clusterLabels:
    cloud: auto-detect
    vendor: auto-detect
  iamPolicyController:
    enabled: true
  policyController:
    enabled: true
  searchCollector:
    enabled: true
  version: 2.4.0
```

Save the changes and close the editor.

- 3.6. Use the `klusterlet.yaml` file to create the `klusterlet` configuration in the hub cluster.

```
[student@workstation ~]$ oc create -f klusterlet.yaml
klusterletaddonconfig.agent.open-cluster-management.io/managed-cluster created
```

At this point, the `ManagedCluster-Import-Controller` will generate a secret named `managed-cluster-import`. This secret contains the `import.yaml` file containing the secret to be used in the cluster to import.

4. Generate the files to import the `managed-cluster` into RHACM.

- 4.1. Obtain the necessary files to import the `managed-cluster` cluster.

First, run the following command to obtain the `klusterlet-crd.yaml` file. This file is used to create the `klusterlet` in the `managed-cluster`.

```
[student@workstation ~]$ oc get secret managed-cluster-import -n managed-cluster \
-o jsonpath={.data.c rds\\.yaml} | base64 --decode > klusterlet-crd.yaml
```

Second, obtain the `import.yaml` file by running the following command:

```
[student@workstation ~]$ oc get secret managed-cluster-import -n managed-cluster \
-o jsonpath={.data.i mport\\.yaml} | base64 --decode > import.yaml
```

This file contains the necessary secrets to import the `managed-cluster` into RHACM.

5. Import the RHCOP cluster, `ocp4-mng.example.com`, available in the lab environment into RHACM. Use `managed-cluster` as the name of the cluster to import.

- 5.1. Use the terminal to log in to the ocp4-mng cluster as the `admin` user. The API Server URL is `https://api.ocp4-mng.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4-mng.example.com:6443
...output omitted...
[student@workstation ~]$
```

- 5.2. Use the `klusterlet-crd.yaml` file to create the `klusterlet` custom resource definition.

```
[student@workstation ~]$ oc create -f klusterlet-crd.yaml
customresourcedefinition.apiextensions.k8s.io/klusterlets.operator.open-cluster-
management.io created
```

- 5.3. Now, use the `import.yaml` file to create the rest of the resources necessary to import the cluster into RHACM.

```
[student@workstation ~]$ oc create -f import.yaml
namespace/open-cluster-management-agent created
serviceaccount/klusterlet created
secret/bootstrap-hub-kubeconfig created
clusterrole.rbac.authorization.k8s.io/klusterlet created
clusterrole.rbac.authorization.k8s.io/open-cluster-management:klusterlet-admin-
aggregate-clusterrole created
clusterrolebinding.rbac.authorization.k8s.io/klusterlet created
deployment.apps/klusterlet created
klusterlet.operator.open-cluster-management.io/klusterlet created
```

- 5.4. Verify the status of the pods running in the `open-cluster-management-agent` namespace.

```
[student@workstation ~]$ oc get pod -n open-cluster-management-agent
NAME                               READY   STATUS    RESTARTS   AGE
klusterlet-bfb4cd68f-j2fdt        1/1    Running   0          53s
klusterlet-registration-agent-5f749f5cf9-8xvp6  1/1    Running   0          36s
klusterlet-registration-agent-5f749f5cf9-pnkj6  1/1    Running   0          36s
klusterlet-registration-agent-5f749f5cf9-slk8j  1/1    Running   0          36s
klusterlet-work-agent-7d848f496b-7mnfw       1/1    Running   0          36s
klusterlet-work-agent-7d848f496b-bwkfv        1/1    Running   1          36s
klusterlet-work-agent-7d848f496b-wn2vd        1/1    Running   0          36s
```

- 5.5. Finally, use the `watch` command to validate the status of the agent pods running in the `open-cluster-management-agent-addon` namespace.

```
[student@workstation ~]$ watch oc get pod -n open-cluster-management-agent-addon
NAME                               READY   STATUS    RESTARTS   AGE
klusterlet-addon-appmgr-7f69b84c76-kvhp7      1/1    Running   0          55s
klusterlet-addon-certpolicyctrl-7c97656db8-5s6xk  1/1    Running   0          54s
```

klusterlet-addon-iampolicyctrl-6f8cccf86c-lj2fg 54s	1/1	Running	0
klusterlet-addon-operator-b479bb446-kpc6t 83s	1/1	Running	0
klusterlet-addon-policyctrl-config-policy-769745c7b6-78sw2 54s	1/1	Running	0
klusterlet-addon-policyctrl-framework-6455bc9558-bz8mc 54s	3/3	Running	0
klusterlet-addon-search-7b5bb78f98-h2jlg 53s	1/1	Running	0
klusterlet-addon-workmgr-5b84cf6dc-nxwvt 52s	1/1	Running	0

When all the pods are ready, press **Ctrl+C** to exit the watch command.

6. Log back in to the hub cluster and verify the status of the managed-cluster.

6.1. Log back in to the hub cluster.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
...output omitted...
[student@workstation ~]$
```

6.2. Verify the status of the managed-cluster.

```
[student@workstation ~]$ oc get managedcluster
NAME          HUB ACCEPTED   ... JOINED   AVAILABLE   AGE
local-cluster  true          ... True     True        4h26m
managed-cluster true          ... True     True        42m
```

The JOINED and AVAILABLE status must be True.

Evaluation

As the **student** user on the **workstation** machine, use the **lab** command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade multicloud-review
```

Finish

Do not make any other changes to the lab environment until the next guided exercise. You will continue using this environment in upcoming exercises.

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise.

```
[student@workstation ~]$ lab finish multicloud-review
```

This concludes the lab.

Summary

In this chapter, you learned:

- About the challenges of hybrid cloud environments and multicluster architectures.
- How the Red Hat OpenShift Platform Plus tools help to address the challenges of multicluster architectures.
- How to install RHACM in a hub cluster.
- How to import a managed cluster to RHACM.

Chapter 2

Inspecting Resources from Multiple Clusters Using the RHACM Web Console

Goal

Describe and navigate the RHACM web console. Configure RBAC and search resources across multiple clusters by using the RHACM search engine.

Objectives

- Locate objects across a fleet of managed clusters by using the search engine and enumerate RHACM features through the web console.
- Create different user roles in RHACM and define an authentication model for multicluster management.

Sections

- Navigating the Red Hat Advanced Cluster Management Web Console (and Guided Exercise)
- Configuring Access Control for Multicluster Management (and Guided Exercise)

Lab

Inspecting Resources from Multiple Clusters Using the RHACM Web Console

Navigating the Red Hat Advanced Cluster Management Web Console

Objectives

After completing this section, you should be able to locate objects across a fleet of managed clusters by using the search engine and enumerate RHACM features through the web console.

RHACM Web Console Overview

The RHACM web console represents a single point for access to the management and observation of a fleet of Kubernetes or OpenShift clusters.

The following sections describe the different components of RHACM web console.

Home

The **Home** pane provides information about RHACM and its use cases, along with links to the main product features. The **Home** pane includes the following submenus:

Welcome

Provides information and links to access the main RHACM features.

Overview

Provides a summary and a high-level overview of the details and status of the managed clusters.

Infrastructure

You can use the **Infrastructure** pane to access cluster lifecycle management, bare metal assets management, Ansible automation configuration, and infrastructure environments management. The **Infrastructure** pane includes the following submenus:

Clusters

The following list shows some of the features about clusters that you can use from the **Clusters** pane:

- Creating or upgrading a cluster in many different public cloud providers.
- Importing an existing cluster to the RHACM hub cluster, to manage it.
- Scaling manually the clusters, or enabling autoscaling. The process of resizing a cluster is different if you created the cluster by using RHACM, or if the cluster already existed and you imported it to RHACM.
- Using RHACM features such as cluster sets, cluster pools, and discover clusters.



Note

A list of supported hub cluster providers is available in the RHACM documentation [https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/supported-clouds#supported-hub-cluster-providers].

A list of supported managed cluster providers is available in the RHACM documentation [https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/index#supported-managed-cluster-providers].

Bare metal assets

In RHACM, a bare metal asset is a collection of physical or virtual servers that you can configure to run your OpenShift clusters. RHACM uses the Intelligent Platform Management Interface (IPMI) set of specifications to interacting with the bare metal assets by using their Baseboard Management Controller (BMC) microcontrollers. The **Bare metal assets** pane integrates bare metal assets into RHACM via their BMC address. RHACM uses the integration of bare metal assets to deploy and manage clusters hosted the virtual or physical infrastructure.

Automation

In the **Automation** pane, you can create Ansible job templates and run Ansible jobs automatically during different stages of a cluster lifecycle. To create Ansible templates you need the Ansible Automation Platform Resource Operator installed on the RHACM hub cluster. You will also need an Ansible Tower 3.7.3 or a later version available.

Infrastructure environments

In RHACM, an infrastructure environment is a pool of resources that allows you to create hosts, and create clusters on those hosts. The main component for managing infrastructure environments is the Central Infrastructure Management (CIM) service, that you need to enable. In the **Infrastructure environments** pane, you can create infrastructure environments and access to them to add hosts.

Applications

You can use the **Applications** pane to create, deploy, and manage applications across the fleet of clusters.

You can find more information in the *Introducing the RHACM Application Model* chapter elsewhere in this Course.

Governance

Through the **Governance** pane, you can create and manage policies and policy controllers, and apply those to the fleet of clusters.

You can find more information in the *Deploying the Compliance Operator Across the Fleet of Clusters Using the RHACM Compliance Operator Policy* chapter elsewhere in this Course.

Credentials

In RHACM, a credential stores the access information for a cloud provider. Kubernetes secrets is the way to store credentials. Each credential has 2 keys, the cloud provider access information, and

a DNS name within that cloud provider. The following is the list of types of credentials that RHACM uses:

- Cloud provider credentials: for example, Amazon Web Services, Google Cloud Platform, or Microsoft Azure
- Data center credentials: for example, Red Hat OpenStack Platform or bare metal resources
- Automation and other credentials: for example, for access to Red Hat Ansible Automation Platform
- Centrally managed: type of credential for on premises environments

You can use the **Credentials** pane to create and administer credentials for all different cloud providers and systems.

RHACM Search Engine

The RHACM search engine provides visibility of the Kubernetes objects across the fleet of clusters from a single user interface.

The RHACM search engine is always enabled, and you can access to it through the magnifying glass icon in upper area of RHACM web console.

The function of the search engine is to index and store the Kubernetes objects present in the fleet of clusters, and calculates the relationships with other objects.

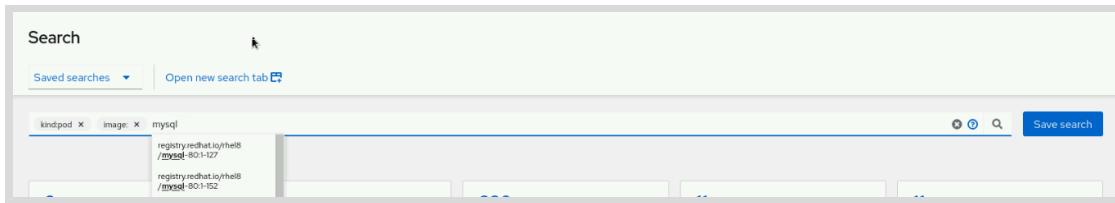
The RHACM search engine architecture is composed of the following components:

- Collector – It is deployed in each of the clusters of the fleet. In the hub cluster the search collector is deployed in the `open-cluster-management` namespace. In the rest of the managed clusters the collector is deployed in the `open-cluster-management-agent-addon` namespace, as part of the `search-collector` add-on. The collector indexes the information of the Kubernetes objects and computes relationships for objects within the managed clusters.
- Aggregator – It is deployed only in the RHACM hub cluster, in the `open-cluster-management` namespace. The function of the aggregator is to collect the data from multiple collectors in different managed clusters. Then the data is wrote to a Redis database present in the `search-redisgraph` stateful set. The aggregator also computes relationships between objects of different clusters, and tracks the activity from the collectors that send data from the managed clusters.
- Search API – Provides an API to access the data in the search index, enforcing role-based access control (RBAC). The search API uses the RBAC of each managed cluster. So, if you are using the RHACM web console you can only search for objects in managed cluster where you already have authorization.
- Search UI – It is deployed as part of the RHACM MultiClusterHub.

Using the Search User Interface

You can type free text in the Search box of the user interface. The search engine uses that text to locate it in the Redis database of indexes, and show results for any object that contains the literal. The UI orders the results in a table with a row for each type of Kubernetes or OpenShift object, represented by the ``Kind: `` filter.

You can refine each search adding more filters to the query. While you type new filters the UI displays the values that are stored for that index.

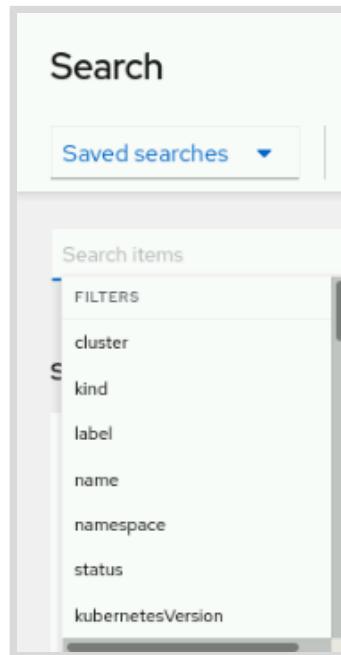


Some of the filters allow arithmetic comparators for fields of numeric nature, as `cpu:`, `replicas:, capacity:,` or `memory:,`

You can use the search UI also for editing the objects. Once you click in the object displayed in a result, you see the YAML definition of the object and the **Edit** button.

What is indexed

The RHACM search engine uses many different filters to classify the objects present in the fleet of clusters. Those filters are the keys of the indexing process. You can see all the filters in the search user interface of the RHACM web console, as displayed in the next image.



You can use any of the filters to perform search, and you can also make free text searches. But you can not make a search refining by all the fields of the object. If the field is not part of a filter, it is not indexed for search.



Note

RHACM search engine uses always the `label:` filter to index every Kubernetes object. It is a good practice to set labels during the different phases of CI/CD to the objects, to make more accurate and quicker searches.



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Web Console* guide at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/web_console/index

For more information about cluster management, refer to the *Red Hat Advanced Cluster Management for Kubernetes Clusters* guide at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/clusters/index

For more information about configuring and tuning the search engine, refer to the *Red Hat Advanced Cluster Management for Kubernetes Web Console* guide at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/web_console/index#search-customization

For more information about credentials, refer to the *Red Hat Advanced Cluster Management for Kubernetes Credentials* guide at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4/html-single/credentials/index

► Guided Exercise

Navigating the Red Hat Advanced Cluster Management Web Console

In this exercise you will list the failed deployments across the fleet of clusters to diagnose and fix a failed MySQL deployment. You will also identify the most used database server. Then, you will ensure that there is no running containers using a specific MySQL image with known vulnerabilities.

Outcomes

You should be able to:

- Use the saved searches to quickly identify the failing deployments across a fleet of clusters in the RHACM web console
- Locate Kubernetes and OpenShift objects across a fleet of clusters in the RHACM web console
- Edit Kubernetes objects across a fleet of clusters in the RHACM web console

Before You Begin

As the student user on the **workstation** machine, use the `lab` command to prepare your system for this exercise. This command creates 15 namespaces in each managed cluster and populate them with 9 different applications using 3 different database servers.



Warning

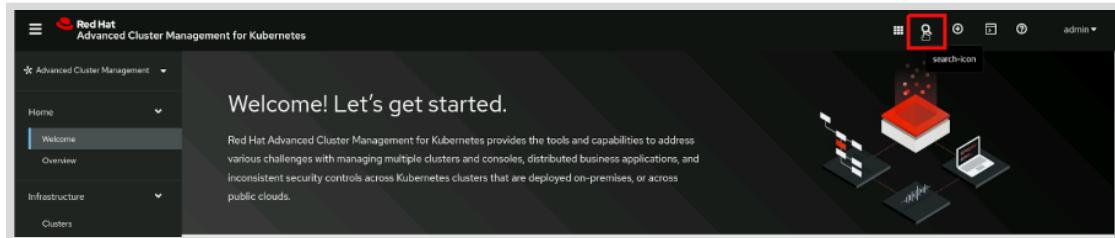
The preparation of this lab will change screen resolution to 1920x1080 for the most comfortable visualization for RHACM console. When the lab is finished resolution is reverted to the existing before the exercise.

```
[student@workstation ~]$ lab start features-console
```

Instructions

- 1. Identify any failed database server deployments across the fleet of clusters using the RHACM web console feature of saved searches.
 - 1.1. From workstation, use Firefox to navigate to the RHACM web console at <https://multicloud-console.apps.ocp4.example.com>.
When prompted, select `htpasswd_provider`
Type the credentials:
 - Username: `admin`
 - Password: `redhat`And then click **Log in**

12. Click the glass icon to navigate to the RHACM search web interface.



Then click the predefined search **Unhealthy pods**. Observe that the syntax of the search is based on the status of the pods.

Name	Namespace	Status	Host IP	Created	Labels
mysql-finance-application-2-5648465868-c7k6l	company-applications-5	Pending		4 minutes ago	app=mysql-finance-application-2 pod-template-hash=5648465868
mysql-finance-application-2-5648465868-thr6j	company-applications-5	Pending		4 minutes ago	app=mysql-finance-application-2 pod-template-hash=5648465868

There are 2 failing pods with names that start with mysql-<ID_OF_THE_POD>. The 2 failing pods are in the namespace called **company-applications-5**. There is a namespace **company-applications-5** in each managed cluster.

13. Click the name of one of the MySQL pods in the Pending status. Scroll down to the status section of the YAML file to identify the cause of the failure.

```

Cluster: managed-cluster Namespace: company-applications-5
185   projected:
186     defaultMode: 420
187     sources:
188       - serviceAccountToken:
189         expirationSeconds: 3607
190         path: token
191       - configMap:
192         items:
193           - key: ca.crt
194             name: kube-root-ca.crt
195             path: ca.crt
196       - downwardAPI:
197         items:
198           - fieldRef:
199             apiVersion: v1
200             fieldPath: metadata.namespace
201             path: namespace
202       - configMap:
203         items:
204           - key: service-ca.crt
205             path: service-ca.crt
206             name: openshift-service-ca.crt
207     status:
208       conditions:
209         - lastProbeTime: null
210         - lastTransitionTime: 2021-11-17T06:04:03Z
211         - message: "No nodes are available: 3 persistentvolumeclaim \"nonexistingPVCname\" not found."
212         - reason: Unschedulable
213         - state: False
214       phase: Pending
215       qosClass: BestEffort
216     status: PodScheduled
217   phase: Pending
218   qosClass: BestEffort

```

The pods are pending because there is a reference to a Persistent Volume Claim (PVC) that does not exist.



Note

The search results can show pods not related to MySQL deployments. Ignore those results and use the ones related to the MySQL deployments.

- 1.4. Navigate back to the predefined search **Unhealthy pods** and click the **related deployment** button. The deployment that contains the pods is **mysql-finance-application-2**. The deployment is present in all the managed clusters.
- 1.5. Click again the glass icon to perform a new search for locating the correct name of the existing PVCs.

Type `namespace:company-applications-5 kind:persistentvolumeclaim` in the search field.

Name	Namespace	Cluster	Status	Persistent volume	Requests	Access mode	Created	Labels
dbclaim	company-applications-5	local-cluster	Bound	pvc-1ebc4964-52d0-4b4e-a10e-828effc6tf64	10Mi		8 minutes ago	app=mysql-persistent-template application=finance-application-2 template=mysql-persistent-template
dbclaim	company-applications-5	managed-cluster	Bound	pvc-ff0f32f6-87bc-43c0-9a62-7f8252ce0b78	10Mi		8 minutes ago	app=mysql-persistent-template application=finance-application-2 template=mysql-persistent-template

The name of the existing PVC is **dbclaim**.

- 1.6. Click again the glass icon to make a new search.

Type `namespace:company-application-5 kind:deployment` in the search field.

- 1.7. Click the name of the first deployment in the **Deployment** results. Edit the yaml file to set **dbclaim** as the PVC name.

```

Cluster: local-cluster Namespace: company-applications-S
Read only mode Edit

137 template:
138   metadata:
139     creationTimestamp: null
140     labels:
141       app: mysql-finance-application-2
142     spec:
143       volumes:
144         - name: db-volume
145           persistentVolumeClaim:
146             claimName: nonexistingclaim
147       containers:
148         - name: mysql
149           image: registry.redhat.io/rhel/mysql-8.0:1-152
150           ports:
151             - name: mysql
152               containerPort: 3306
153               protocol: TCP
154           env:
155             - name: MYSQL_ROOT_PASSWORD
156               value: rootpass
157             - name: MYSQL_USER
158               value: user1
159             - name: MYSQL_PASSWORD
160               value: user1
161             - name: MYSQL_DATABASE
162               value: items
163           resources: {}
164           volumeMounts:
165             - name: db-volume
166               mountPath: /var/lib/mysql/data
167             terminationMessagePath: /dev/termination-log
168             terminationMessagePolicy: File
169             imagePullPolicy: IfNotPresent
170             restartPolicy: Always
171             terminationGracePeriodSeconds: 30

```

Then click the **Save** button to trigger a redeployment automatically using the right name of the PVC.

- 1.8. Repeat the operation to fix the failing deployment in the other cluster.
 - 1.9. Navigate back to the saved searches and review the **Unhealthy pods** saved search button to verify that there are no more MySQL pods failing.
- ▶ 2. Search for any database software deployed across the fleet of managed clusters.
- 2.1. In the search field type **namespace** : and look at the different namespaces across the fleet of clusters offered by the search engine. You can see that the name of the namespaces are always prepended by **company-applications-**.
 - 2.2. Clear the search field and type **kind:deployment**. Then add **kind:deploymentconfig**. Notice that the RHACM console merges the filters using the syntax **kind:deployment, deploymentconfig**. Finally add the free text "mysql".
- Notice that the MySQL instances are deployed as Kubernetes Deployment objects. The label **application** indicates the name of the application. There are 2 different applications using MySQL containers: **globalshop-application**, and **finance-application-2**.

Name	Namespace	Cluster	Desired	Current	Ready	Available	Created	Labels
mysql-finance-application-2	company-applications-5	local-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=finance-application-2
mysql-globalshop-application	company-applications-6	local-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=globalshop-application
mysql-globalshop-application	company-applications-15	managed-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=globalshop-application
mysql-globalshop-application	company-applications-6	managed-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=globalshop-application
mysql-finance-application-2	company-applications-5	managed-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=finance-application-2
mysql-globalshop-application	company-applications-7	managed-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template template=mysql-persistent-template application=globalshop-application
mysql-finance-application-2	company-applications-4	managed-cluster	1	1	1	1	17 hours ago	app=mysql-persistent-template application=finance-application-2

- 2.3. In the search field, remove the free text "mysql" and replace it with "mariadb". Do not remove the kind: deployment, deploymentconfig filter.

Notice that the MariaDB instances are deployed as OpenShift DeploymentConfig objects. Use the label application to locate the name. There are 2 different applications using MariaDB containers: **humanresources-application-1**, and **marketing-application-2**.

- 2.4. In the search field, remove the free text "mysql" and replace it with "postgresql". Do not remove the kind: deployment, deploymentconfig filter.

You will see that PostgreSQL is deployed as OpenShift DeploymentConfig objects.

Use the label application to locate the name. There are 5 different applications using PostgreSQL containers: **finance-application-1**, **finance-application-3** **finance-application-4**, **humanresources-application-2**, and **marketing-application-1**.

The most used database server is PostgreSQL that is used in 5 applications of existing 9.

- 3. Using the RHACM web console, locate any running container using a MySQL image with known vulnerabilities. Red Hat Container Catalog states that updated image in use for MySQL 8 must be `registry.redhat.io/rhel8/mysql-80:1`.



Note

You can see the mapping between MySQL version and Red Hat MySQL container image version in <https://catalog.redhat.com/software/containers/rhel8/mysql-80/5ba0ad4cdd19c70b45cbf48>

There is a version of MySQL image with the tag `registry.redhat.io/rhel8/mysql-80:1-127` that contains important vulnerabilities and must not be used.

The screenshot shows the Red Hat Ecosystem Catalog interface. At the top, there's a navigation bar with links for Home, Software, Container images, MySQL 8.0, Hardware, Software, and Cloud & service providers. Below the navigation, it says "Standalone Image" and "MySQL 8.0". It shows the tag "rhel8/mysql-80" and filters for "Architecture: amd64" and "Tag: 1-127". On the right, there's a "Provided by" section with the Red Hat logo. The main content area has a heading "1-127" with a note about an updated image available. Below this, tabs for Overview, Security (which is selected), Technical Information, Packages, Dockerfile, and Get this image are shown. The Security tab displays a "Health index" with a scale from A to F, where D is highlighted in orange. A note states: "This image is affected by Critical (no older than 90 days) or Important (no older than 12 months) security updates. The Container Health Index analysis is based on RPM packages signed and created by Red Hat, and does not grade other software that may be included in a container image." Below the health index, it lists "53 security vulnerabilities affecting 43 packages" with a breakdown: Critical (0), Important (4), Moderate (46), and Low (2). To the right, it shows "Release category: Generally Available", "Advisory: RHBA-2021:1107", and "Privilege mode: Unprivileged".

Find the applications using the container image `mysql-80` with the tag `1-127`.

- 3.1. Type `registry.redhat.io/rhel8/mysql-80` in the search field,
Notice that the results only contain objects of type pods.
 - 3.2. Click one of the pods to see the reference to the container image in use in its YAML definition. Most of the MySQL running pods are using `registry.redhat.io/rhel8/mysql-80` with tag `1-152`.
 - 3.3. Navigate to the initial search page of RHACM and add "1-127" in the search field.
Do not remove the filter `registry.redhat.io/rhel8/mysql-80`. Inspect the results: you will see two pods running in namespaces with the same name, but in different clusters.
 - 3.4. Click one of the pods to see the image in use in its YAML definition.
 - 3.5. Click the **Logs** tab within the pod page to verify that the running version of MySQL is the old one, 8.0.21, as stated in the Red Hat Container Catalog.
- With this information, you can inform the developers which namespaces contain running pods using a vulnerable MySQL container image.

Finish

On the **workstation** machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish features-console
```

This concludes the guided exercise.

Configuring Access Control for Multicloud Management

Objectives

After completing this section, you should be able to create different user roles in RHACM and define an authentication model for multicloud management.

Role-based Access Control in RHACM

RHACM access control takes advantage of the Red Hat OpenShift Container Platform authentication layer and roles. Administrators can design a multicloud authentication model to define access for different tenants to different groups of clusters, with fine-grained control.

RHACM Cluster Sets

To create different groups of clusters, RHACM defines the concept of a `ClusterSet`. A cluster set tags all the clusters pertaining to it with the `cluster.open-cluster-management.io/clusterset=<managed_clusterset_name>` tag. RHACM uses this tag to identify the clusters belonging to the cluster set, enabling you to take action on all the member clusters at once.

The following table contains the roles created by RHACM to define different access levels to a cluster, all the clusters, or a group of clusters in the same cluster set.

Role	Definition
<code>open-cluster-management:cluster-manager-admin</code>	RHACM super user, with full access. Can create a <code>ManagedCluster</code> resource.
<code>open-cluster-management:admin:<managed_cluster_name></code>	RHACM administrator access to the <code>ManagedCluster</code> resource named <code><managed_cluster_name></code> .
<code>open-cluster-management:view:<managed_cluster_name></code>	RHACM view access to the <code>ManagedCluster</code> resource named <code><managed_cluster_name></code> .
<code>open-cluster-management:managedclusterset:admin:<managed_clusterset_name></code>	RHACM administrator access to the <code>ManagedCluster</code> resource named <code><managed_cluster_name></code> . The user has also administrator access to the resources of type <code>managedcluster.cluster.open-cluster-management.io</code> , <code>clusterclaim.hive.openshift.io</code> , <code>clusterdeployment.hive.openshift.io</code> , and <code>clusterpool.hive.openshift.io</code> with the label <code>cluster.open-cluster-management.io/clusterset=<managed_clusterset_name></code> .

Role	Definition
open-cluster-management:managedclusterset:view:<managed_clusterset_name>	RHACM view access to the ManagedCluster resource named <managed_cluster_name>. The user has also administrator access to the resources of type managedcluster.cluster.open-cluster-management.io, clusterclaim.hive.openshift.io, clusterdeployment.hive.openshift.io, and clusterpool.hive.openshift.io with the label cluster.open-cluster-management.io/clusterset=<managed_clusterset_name>.
open-cluster-management:subscription-admin	Can create Git subscriptions that deploy Kubernetes resources YAML files to multiple namespaces.

RHACM also supports the following default RHOCP roles:

Role	Definition
cluster-admin	RHOCP super user, with full access.
admin, edit, view	RHOCP default roles. A user with a namespace-scoped binding to these roles has access to open-cluster-management resources in a specific namespace. Cluster-wide binding to the same roles gives access to all of the open-cluster-management resources.

RHACM roles are especially useful for assigning user or group permissions to a cluster set. Apart from using the `oc adm` command, the admin and view cluster set roles can be assigned to users in the **Access management** tab within the clusterset details pane in the RHACM web console. You can assign the `admin` and `view` cluster set roles by using the `oc adm` command or by navigating to the **Access management** tab within the **Cluster sets** details pane in the RHACM web console.

Cluster Selectors and Cluster Sets

A cluster cannot pertain to more than one cluster set simultaneously. If you need to define overlapping groups of clusters, you can use cluster labels acting as selectors.

Combining cluster sets and cluster selectors, administrators can use other tools outside of RHACM to act on a subset of clusters belonging to the same or different cluster sets.

Placement Rules

Using placement rules, administrators can define a subset of clusters pertaining to different cluster sets for placing Kubernetes resources.

Placement rules can use `labelSelector`, `claimSelector`, `clusterSet`, and `numberOfClusters` parameters to determine the clusters on which to deploy the resources.

For example, the following `Placement` uses the `claimSelector` parameter to specify that resources should deploy on clusters from the `us-west-1` region.

```
apiVersion: cluster.open-cluster-management.io/v1alpha1
kind: Placement
metadata:
  name: placement2
  namespace: ns1
spec:
  predicates:
    - requiredClusterSelector:
        claimSelector:
          matchExpressions:
            - key: region.open-cluster-management.io
              operator: In
              values:
                - us-west-1
```



Note

For more information, refer to the *Using ManagedClusterSets with Placement* section at https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/clusters/index#placement-managed

Benefits of a Centralized Identity Management Service

One important challenge in multicloud environments is to manage the authentication and authorization mechanisms in a fully controlled manner across different Kubernetes and OpenShift clusters.

As noted previously, RHACM takes advantage of the Red Hat OpenShift Container Platform authentication and authorization layers. The authentication layer of OpenShift uses the OAuth open standard as an authorization framework. By default, OpenShift uses an internal OAuth server. Unfortunately, neither Kubernetes nor OpenShift currently provides a way to federate all the internal OAuth servers of each cluster.

To solve this problem, you can use an external identity manager for central management of the identities of the users that access to your clusters.



Note

You can find the list of supported OAuth providers in the *Supported identity providers* section of the *Red Hat OpenShift Container Platform Authorization and Authentication* guide at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.9/html-single/authentication_and_authorization/index#supported-identity-providers



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Access Control* guide at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/access_control/index

► Guided Exercise

Configuring Access Control for Multicluster Management

In this exercise you will create cluster sets using existing RHOPC clusters. Then you will configure existing users in RHOPC to manage those cluster sets.

Outcomes

You should be able to:

- Create a `ClusterSet`
- Assign a managed cluster to a `ClusterSet`
- Assign predefined RHACM roles to different users
- Remove a `ClusterSet`

Before You Begin

Sample prerequisites

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

The lab environment includes a Red Hat Identity Management server. This command ensures that the necessary users and roles exist in the Red Hat Identity Management server. It also configures an additional LDAP identity provider in the hub cluster.

The following table summarizes the users and groups available in the Red Hat Identity Management server for this exercise.

User	Group
prod-admin	production-administrators
stage-admin	stage-administrators

```
[student@workstation ~]$ lab start features-users
```



Note

RHACM is installed in the hub cluster with URL `ocp4.example.com`.



Note

The route to the RHACM web console is `multicloud-console.apps.ocp4.example.com`.

Instructions

- 1. From the RHACM web console, create a cluster set named `production` and add the cluster named `local-cluster`.

- 1.1. Open Firefox and navigate to the RHACM console in `multicloud-console.apps.ocp4.example.com`.

When prompted, select `htpasswd_provider`

Type the credentials:

- Username: `admin`
- Password: `redhat`

And then click **Log in**

- 1.2. Navigate to **Infrastructure → Clusters** and click **Cluster sets**.

The screenshot shows the Red Hat Advanced Cluster Management for Kubernetes web interface. The top navigation bar includes the Red Hat logo and the title "Advanced Cluster Management for Kubernetes". The left sidebar has a "Home" section with "Welcome" and "Overview" options, and an "Infrastructure" section with "Clusters", "Bare metal assets", "Automation", and "Infrastructure environments". The main content area is titled "Clusters" and shows tabs for "Managed clusters", "Cluster sets" (which is the active tab), "Cluster pools", and "Discovered clusters". A prominent callout box in the center says "Create Cluster on Premise with a cloud like experience" and "The best solution for creating cluster on an On Premise at scale. Easily create ready to go clusters for your applications. Easily create and re-create clusters from hosts that are provided by an infrastructure environment." It features a red hexagonal cluster icon and buttons for "Get started with infrastructure environments" and "Dismiss". Below the callout are search and filter tools, and buttons for "Create cluster", "Import cluster", and "Actions".

- 1.3. Create a cluster set named `production`. Add the `local-cluster` to it.

From within the **Cluster sets** page, scroll down and click **Create cluster set**

The screenshot shows the 'Clusters' page in the RHACM web console. The 'Cluster sets' tab is active. A central message states 'You don't have any cluster sets.' with a 'Create cluster set' button below it. The left sidebar shows the 'Clusters' section is selected under 'Infrastructure'.

Type production in the Cluster set name field and click Create

A modal dialog titled 'Create cluster set' is open. It explains that creating a ManagedClusterSet allows grouping resources and managing access control. The 'Cluster set name' field contains 'production'. The 'Create' button is highlighted in blue.

In the confirmation menu, click Manage resource assignments

The 'Clusters' page shows a success message: 'Cluster set successfully created'. It also displays a table with one row: 'production' under 'Name' and 'Add clusters' under 'Multi-cluster network status'. A 'Manage resource assignments' button is visible at the bottom of the message box.

Select the cluster named local-cluster and click Review

Name	Kind	Current cluster set
local-cluster	ManagedCluster	-
managed-cluster	ManagedCluster	-

In the confirmation screen, click **Save**

Finally, you are redirected to the **Overview** tab of the production cluster set.

- 1.4. Verify the labels corresponding to the production cluster set.

From the **Managed clusters** page of the production cluster set, click **Managed clusters**

Observe the `cluster.open-cluster-management.io/clusterset=production` label.

RHACM uses this labels to identify clusters pertaining to a cluster set.

- 2. From the RHACM web console, create a cluster set named **stage** and add the cluster named **managed-cluster**.

- 2.1. Navigate to **Infrastructure → Clusters** and click **Cluster sets** and repeat all the previous steps to create a **stage** cluster set.

Add the cluster named **managed-cluster** to the **stage** cluster set.

From the **Managed clusters** page of the **stage** cluster set, click **11 labels** to verify that RHACM adds the label `cluster.open-cluster-management.io/clusterset=stage` to the **managed-cluster** cluster.

- 3. From the terminal, assign the pre-existing group **production-administrators** administrator permissions in the **production** cluster set. Also, assign the group **production-administrators** view permissions in the **stage** cluster set.



Note

RHACM automatically creates default `admin` and `view` cluster roles in the RHOCP hub cluster for each new cluster set. For instance, to grant administrator permissions to the `production` cluster set, use the role `open-cluster-management:managedclusterset:admin:<clusterset_name>`.

To grant view permissions to the `production` cluster set, use the role `open-cluster-management:managedclusterset:view:<clusterset_name>`.

- 3.1. Open a terminal and log in to the `ocp4` cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 3.2. Use `oc adm` to assign the cluster role `open-cluster-management:managedclusterset:admin:production` to the existing group `production-administrators`.



Note

The group `production-administrators` already exists in the Red Hat Identity Management server in the lab environment. The `lab start` script triggers a group synchronization between the hub cluster and the Red Hat Identity Management server.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-group \
  open-cluster-management:managedclusterset:admin:production
  production-administrators
clusterrole.rbac.authorization.k8s.io/open-cluster-
management:managedclusterset:admin:production added: "production-administrators"
```

- 3.3. Use `oc adm` to assign the cluster role `open-cluster-management:managedclusterset:view:stage` to the existing group `production-administrators`.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-group \
  open-cluster-management:managedclusterset:view:stage production-administrators
clusterrole.rbac.authorization.k8s.io/open-cluster-
management:managedclusterset:view:stage added: "production-administrators"
```

4. From the terminal, assign the pre-existing group `stage-administrators` administrator permissions in the `stage` cluster set. Do not assign the `stage-administrators` group any permissions to the `production` cluster set.

- 4.1. Use `oc adm` to assign the cluster role `open-cluster-management:managedclusterset:admin:stage` to the existing group `stage-administrators`.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-group \
  open-cluster-management:managedclusterset:admin:stage stage-administrators
clusterrole.rbac.authorization.k8s.io/open-cluster-
management:managedclusterset:admin:stage added: "stage-administrators"
```

The following table summarizes the users, groups and roles at this point.

User	Group	Roles
prod-admin	production-administrators	open-cluster-management:managedclusterset:admin:production, open-cluster-management:managedclusterset:view:stage
stage-admin	stage-administrators	open-cluster-management:managedclusterset:admin:stage

- 5. Log out from RHACM and close Firefox to clean any cached credentials in the browser. Open it again and log in to the RHACM web console with the `prod-admin` user from the Red Hat Identity Management provider. Verify that the `prod-admin` user has administrator permissions on the production cluster set, comprised by the cluster named `local-cluster`. Also, verify that the `prod-admin` user has view only permissions on the stage cluster set.



Note

The `prod-admin` user is member of the `production-administrators` group.

- 5.1. Log out from RHACM.

From the RHACM web console, click the `admin` button at the top right corner, and click `Logout`.

- 5.2. Close all the instances of Firefox. Open a new Firefox window and log in to the RHACM web console with the user `prod-admin` from the Red Hat Identity Management provider.

When prompted, select **Red Hat Identity Management**

Type the credentials:

- Username: prod-admin
- Password: redhat

And then click **Log in**

- 5.3. Navigate to **Infrastructure → Clusters**. Note the differences between the two clusters.

The **Upgrade available** button is active for the **local-cluster**, but inactive for the **managed-cluster**. This happens because the **prod-admin** user has administrator privileges in the **production** cluster set, but only view permissions in the **stage** cluster set.

- 5.4. Click the **Options** button of the **local-cluster**.

Note that as an administrator, all the options are enabled.

Name	Status	Infrastructure provider	Distribution version	Labels	Nodes
local-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	14 labels	3
managed-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	11 labels	

A cluster set administrator can perform actions on clusters such as **Edit labels**, **Upgrade cluster**, **Select channel** for updates, **Search cluster resources**, and **Detach cluster**.

- 5.5. Click the options button of the **managed-cluster**.

Name	Status	Infrastructure provider	Distribution version	Labels	Nodes
local-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	14 labels	3
managed-cluster	Ready	Other	OpenShift 4.8.2 Upgrade available	11 labels	

Note that as a viewer, only the **Search cluster** option is enabled.

- 6. Log out from RHACM and close Firefox to clean any cached credentials in the browser. Open it again and log in to the RHACM web console with the **stage-admin** user from the Red Hat Identity Management provider. Verify that the **stage-admin** user has administrator permissions on the **stage** cluster set, comprised by the cluster named **managed-cluster**. Also, verify that the **stage-admin** don't have any permissions on the **production** cluster set.



Note

The **stage-admin** user is member of the **stage-administrators** group.

- 6.1. Log out from RHACM.

From the RHACM web console, click the **prod-admin** button at the top right corner, and click **Logout**.

- 6.2. Close all the instances of Firefox. Open a new Firefox window and log in to the RHACM web console with the user **stage-admin** from the Red Hat Identity Management provider.

When prompted, select **Red Hat Identity Management**

Type the credentials:

- Username: **stage-admin**
- Password: **redhat**

And then click **Log in**

- 6.3. Navigate to **Infrastructure → Clusters**. Note that the **stage-admin** user can only see clusters belonging to the **stage** cluster set.

The **Upgrade available** button and all the options from the **Options** button are enabled.

This is because the **stage-admin** user has administrator permissions on the **stage** cluster set.

The cluster **local-cluster**, pertaining to the **production** cluster set does not appear in the **Managed clusters** list.

- 7. Log out from the RHACM web console and close all the Firefox instances to remove any cached credentials.

- 7.1. From the RHACM web console, click the **stage-admin** button at the top right corner, and click **Logout**.

Close the Firefox window.

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish features-users
```

This concludes the guided exercise.

▶ Lab

Inspecting Resources from Multiple Clusters Using the RHACM Web Console

In this lab, you assign the cluster-wide role `view` to a group of users so that they can perform searches across the fleet of clusters. Then, you log in with a user from that group and search for a deployment with few replicas. Finally you log in to RHACM with a user with admin permissions in the managed cluster set and fix the problem.

Outcomes

You should be able to:

- Manage RHACM roles and groups to allow searches across a fleet of clusters.
- Locate Kubernetes and OpenShift objects across a fleet of clusters by using the correct group of users.
- Edit Kubernetes objects across a fleet of clusters using the correct group of users.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is installed and running and the managed clusters are present. It also checks that the necessary users and roles exist in the Red Hat Identity Management server. Finally, it creates all the deployments across the fleet of clusters.

The following table summarizes the users and groups available in the Red Hat Identity Management server for this exercise.

User	Group
apac-operator	apac-operators
emea-operator	emea-operators
fleet-searcher	fleet-searchers

The following table summarizes the managed cluster sets and the roles of the groups available in the laboratory environment for this exercise.

Managed Cluster Set	Group	Role of the Group
apac	apac-operators	open-cluster-management:managedclusterset:admin:apac
emea	emea-operators	open-cluster-management:managedclusterset:admin:emea

Managed Cluster Set	Group	Role of the Group
-	fleet-searcher	-

```
[student@workstation ~]$ lab start features-review
```

Instructions

1. Assign the cluster-wide role `view` to the `fleet-searchers` group to allow searches across all the objects in the fleet of clusters.
2. Locate all Kubernetes objects from the applications that have the `app=finance-application` label and contain fewer than 3 replicas.
3. Fix the number of replicas of the two deployments by logging in as a user belonging to the `emea-operators` group.

Evaluation

Lab scripts to be developed.

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises. Do not make any other changes to the lab environment until the next guided exercise. You will continue using it in later guided exercises.

```
[student@workstation ~]$ lab finish features-review
```

This concludes the lab.

► Solution

Inspecting Resources from Multiple Clusters Using the RHACM Web Console

In this lab, you assign the cluster-wide role `view` to a group of users so that they can perform searches across the fleet of clusters. Then, you log in with a user from that group and search for a deployment with few replicas. Finally you log in to RHACM with a user with admin permissions in the managed cluster set and fix the problem.

Outcomes

You should be able to:

- Manage RHACM roles and groups to allow searches across a fleet of clusters.
- Locate Kubernetes and OpenShift objects across a fleet of clusters by using the correct group of users.
- Edit Kubernetes objects across a fleet of clusters using the correct group of users.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is installed and running and the managed clusters are present. It also checks that the necessary users and roles exist in the Red Hat Identity Management server. Finally, it creates all the deployments across the fleet of clusters.

The following table summarizes the users and groups available in the Red Hat Identity Management server for this exercise.

User	Group
apac-operator	apac-operators
emea-operator	emea-operators
fleet-searcher	fleet-searchers

The following table summarizes the managed cluster sets and the roles of the groups available in the laboratory environment for this exercise.

Managed Cluster Set	Group	Role of the Group
apac	apac-operators	open-cluster-management:managedclusterset:admin:apac
emea	emea-operators	open-cluster-management:managedclusterset:admin:emea

Managed Cluster Set	Group	Role of the Group
-	fleet-searcher	-

```
[student@workstation ~]$ lab start features-review
```

Instructions

1. Assign the cluster-wide role `view` to the `fleet-searchers` group to allow searches across all the objects in the fleet of clusters.
 - 1.1. From the `workstation` machine, open a terminal and log in to the `ocp4` cluster as the `admin` user. The password is `redhat`. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
...output omitted...
[student@workstation ~]$
```

- 1.2. Verify the roles assigned to the `apac-operators` group.

```
[student@workstation ~]$ oc get clusterrolebindings.authorization \
--custom-columns=role:.roleRef.name ,groups:.groupNames.* | grep apac-operators

open-cluster-management:managedclusterset:admin:apac      apac-operators
open-cluster-management:admin:local-cluster
system:masters,system:cluster-admins,apac-operators
open-cluster-management:view:local-cluster               system:cluster-
readers,system:cluster-admins,system:masters,apac-operators
[student@workstation ~]$
```

- 1.3. Verify the roles assigned to the `fleet-searchers` group.

```
[student@workstation ~]$ oc get clusterrolebindings.authorization \
--custom-columns=role:.roleRef.name ,groups:.groupNames.* | grep fleet-searchers
[student@workstation ~]$
```

This group is not assigned any role.

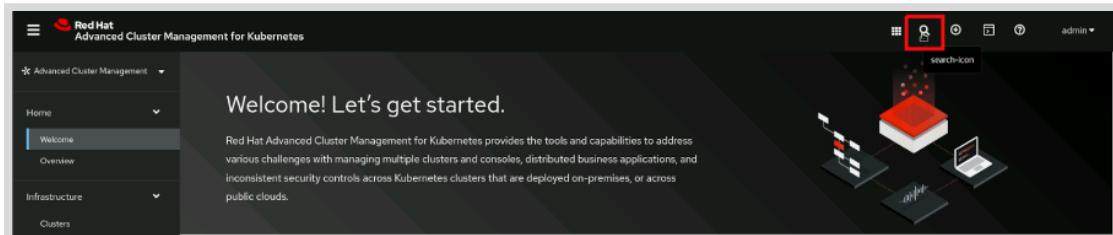
- 1.4. Use the `oc adm` command to assign the `view` cluster role to the `fleet-searchers` group.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-group \
view fleet-searchers
clusterrole.rbac.authorization.k8s.io/view added: "fleet-searchers"
```

- 1.5. Reverify the `fleet-searchers` group roles to identify the additional roles assigned by the RHACM.

```
[student@workstation ~]$ oc get clusterrolebindings.authorization \
-ocustom-columns=role:.roleRef.name ,groups:.groupNames.* | grep fleet-searchers
open-cluster-management:view:local-cluster      system:cluster-
readers,system:cluster-admins,system:masters,apac-operators,fleet-searchers
open-cluster-management:view:managed-cluster    system:cluster-
readers,system:cluster-admins,system:masters,emea-operators,fleet-searchers
view                                         fleet-searchers
[student@workstation ~]$
```

2. Locate all Kubernetes objects from the applications that have the app=finance-application label and contain fewer than 3 replicas.
 - 2.1. From the workstation machine, use Firefox to navigate to the RHACM web console at <https://multicloud-console.apps.ocp4.example.com>.
When prompted, select **Red Hat Identity Management**
Type the following credentials:
 - Username: **fleet-searcher**
 - Password: **redhat**And then click **Log in**.
 - 2.2. Click the magnifying glass icon to navigate to the RHACM search web interface.



Type `label:app=finance-application current:<3` in the search field. The results show 2 different Deployment objects. Both deployments are in the `managed-cluster` cluster, but one is within the `company-applications-5` namespace, and the other one is within the `company-applications-7`.

The RHACM console does not allow the `fleet-searcher` user to modify these two deployments.

- 2.3. Click **Related cluster** in the related objects of the results.

The screenshot shows the RHACM Web Console search interface. At the top, there is a search bar with the query "label:app=finance-application current:3". Below the search bar, there are four cards representing related resources:

- 1 Related cluster
- 2 Related pod
- 2 Related service
- 2 Related replicaset

Then expand the related cluster to see all the information about the managed-cluster.

- 2.4. Locate the cluster set that contains the managed-cluster by finding the value of the label `cluster.open-cluster-management.io/clusterset=`.

The screenshot shows the detailed view of a managed cluster. The cluster has the following configuration:

Name	Value
Available	True
Hub accepted	True
Joined	True
Nodes	3
Kubernetes version	v1.21.1+05lac4f
CPU	12
Memory	48074Mi
Console URL	Launch
Labels	<ul style="list-style-type: none"> cluster.open-cluster-management.io/clusterset=emea clusterID=e5ec7f6-5a2d-4373-ac4c-b7b47f321bf name=managed-cluster cluster.open-cluster-management.io/clusterset=emea

The managed-cluster cluster belongs to the emea cluster set. The emea-operators group can modify the number of replicas.

3. Fix the number of replicas of the two deployments by logging in as a user belonging to the emea-operators group.
 - 3.1. Log out from RHACM and close Firefox to clean any cached credentials in the browser. Open Firefox again and log in to the RHACM web console.
When prompted, select Red Hat Identity Management
Type the credentials:

- Username: emea-operator
- Password: redhat

And then click **Log in**.

- 3.2. Click the magnifying glass icon to navigate to the RHACM search web interface.

Repeat previous search by typing **label:app=finance-application current:<3** in the search field. The same 2 deployments display as when you searched as a member of the fleet-searchers group, but now the members of the emea-operators group can edit the objects.

- 3.3. Click the name of the deployment in the **company-applications-5** namespace.

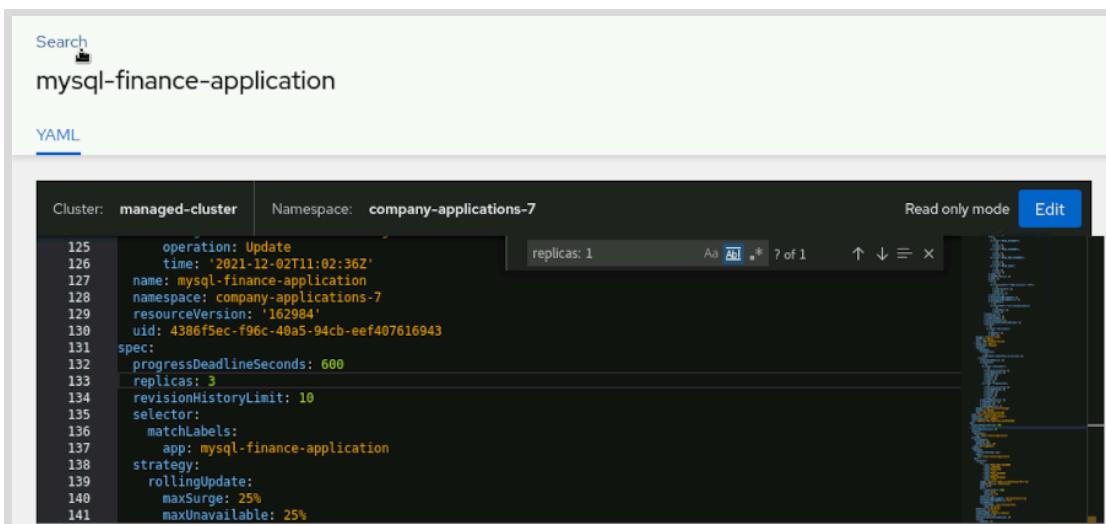
Then click **Edit**.

- 3.4. Find the text **replicas: 1** in the **spec:** section of the deployment. You can use **crtl +f** to find inside the YAML file.

Then change the value to **replicas: 3** and click **Save**.

- 3.5. Repeat the preceding operation in the deployment present in the **company-applications-7** namespace.

- 3.6. Check that the previous search has no more results by clicking the **Search** link.



The screenshot shows the RHACM search interface with a search term of "mysql-finance-application". The results list a single deployment named "mysql-finance-application" in the "company-applications-7" namespace. The "YAML" tab is active, displaying the deployment's YAML configuration. The "replicas: 1" line is highlighted with a blue border, indicating it is being edited. The "Edit" button is visible at the top right of the YAML editor.

```
Cluster: managed-cluster Namespace: company-applications-7
125     operation: Update
126     time: '2021-12-02T11:02:36Z'
127     name: mysql-finance-application
128     namespace: company-applications-7
129     resourceVersion: '162984'
130     uid: 4386f5ec-f96c-40a5-94cb-eef407616943
131 spec:
132   progressDeadlineSeconds: 600
133   replicas: 3
134   revisionHistoryLimit: 10
135   selector:
136     matchLabels:
137       app: mysql-finance-application
138   strategy:
139     rollingUpdate:
140       maxSurge: 25%
141       maxUnavailable: 25%
```

Evaluation

Lab scripts to be developed.

Finish

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises. Do not make any other changes to the lab environment until the next guided exercise. You will continue using it in later guided exercises.

```
[student@workstation ~]$ lab finish features-review
```

This concludes the lab.

Summary

In this chapter, you learned:

- How to access the functionalities of RHACM from the web console.
- How to locate Kubernetes objects across de fleet of clusters by using the RHACM search engine.
- How the role-based access control model of RHACM works, and its default roles.
- How to define an authentication model to manage the fleet of clusters classifying the clusters in RHACM clustersets.

Chapter 3

Deploying and Managing Policies for Multiple Clusters with RHACM

Goal

Deploy and manage policies in multicluster environments by using RHACM governance.

Objectives

- Deploy policies on multiple clusters by using the command line and the RHACM Governance Dashboard.
- Deploy the compliance operator policy and view compliance reports from multiple clusters by using the command line and the RHACM Governance Dashboard.
- Deploy the gatekeeper policy and gatekeeper constraints to multiple clusters by using the command line and the RHACM Governance Dashboard.

Sections

- Deploying and Managing Policies with RHACM (and Guided Exercise)
- Deploying and Configuring the Compliance Operator for Multiple clusters using RHACM (and Guided Exercise)
- Integrating RHACM with Other Policy Engines (and Guided Exercise)

Lab

Deploying and Managing Policies with RHACM

Deploying and Managing Policies with RHACM

Objectives

After completing this section, you should be able to deploy policies on multiple clusters by using the command line and the RHACM Governance Dashboard.

RHACM Governance Overview

Introduction

The enterprises must ensure that the security of their workloads and the configuration hosted on clouds meet corporate standards. RHACM Governance provides a policy-based approach for companies to monitor these corporate standards automatically.

Governance Architecture

Provides a summary and components of governance architecture

Governance Dashboard

Provides a summary of governance dashboard

RHACM Policy Controllers Overview

The policy controllers monitor if your cluster is compliant with a policy. The policy controller also reports the policy status, which you can see on the RHACM governance dashboard.

List of Policy Controller

- Kubernetes configuration policy controller
- Certificate policy controller
- IAM policy controller
- Integrate third-party policy controllers
- Custom policy controller

RHACM Policy Overview

Provides a summary of the RHACM Policy

Describe Policy YAML File

A policy YAML file has some required and some optional parameters. The optional fields and parameters depend on the policy controller.

```
apiVersion: policy.open-cluster-management.io/v1 ①
kind: Policy ②
metadata:
```

```

name: policy-namespace ③
annotations: ④
  policy.open-cluster-management.io/standards: NIST SP 800-53 ⑤
  policy.open-cluster-management.io/categories: CM Configuration Management ⑥
  policy.open-cluster-management.io/controls: CM-2 Baseline Configuration ⑦
spec:
  remediationAction: inform ⑧
  disabled: false ⑨
  policy-templates: ⑩
    - objectDefinition:
        apiVersion: policy.open-cluster-management.io/v1
        kind: ConfigurationPolicy
        metadata:
          name: policy-namespace-example
        spec:
          remediationAction: inform # the policy-template spec.remediationAction
is overridden by the preceding parameter value for spec.remediationAction.
          severity: low
        namespaceSelector:
          exclude: ["kube-*"]
          include: ["default"]
      object-templates:
        - complianceType: musthave
          objectDefinition:
            kind: Namespace # must have namespace 'prod'
            apiVersion: v1
            metadata:
              name: prod
    ---
apiVersion: policy.open-cluster-management.io/v1
kind: PlacementBinding
metadata:
  name: binding-policy-namespace
placementRef:
  name: placement-policy-namespace
  kind: PlacementRule
  apiGroup: apps.open-cluster-management.io
subjects:
  - name: policy-namespace
    kind: Policy
    apiGroup: policy.open-cluster-management.io
  ---
apiVersion: apps.open-cluster-management.io/v1
kind: PlacementRule
metadata:
  name: placement-policy-namespace
spec:
  clusterConditions:
    - status: "True"
      type: ManagedClusterConditionAvailable
  clusterSelector:
    matchExpressions:
      - {key: environment, operator: In, values: ["dev"]}

```

- ① *apiVersion*: Required. The value for this parameter is `policy.open-cluster-management.io/v1`
- ② *kind*: Required. The value for this parameter is `Policy`
- ③ *metadata.name*: Required. The parameter sets name of the policy
- ④ *metadata.annotations*: Optional. This parameter defines standards for policy to validate.
- ⑤ *annotations.policy.open-cluster-management.io/standards*: This parameter defines the name of the security standard.
- ⑥ *annotations.policy.open-cluster-management.io/categories*: This parameter defines the name of the security control category.
- ⑦ *annotations.policy.open-cluster-management.io/controls*: This parameter defines the name of the security control.
- ⑧ *spec.remediationAction*: Optional. The values for this parameter are `enforce` and `inform`.
- ⑨ *spec.disabled*: Required. The values for this parameter are `true` and `false`.
- ⑩ *spec.policy-templates*: This parameter is use for policy creation.

Stable and Community Policies

RHACM can use stable policies, supported by Red Hat, and also can use policies from the community. You can see the most of the stable policies in RHACM Dashboard.

Deploying policy on multicloud with RHACM

- Deploy policy from governance dashboard
- Deploy policy from command line



References

For more information, refer to the _Red Hat Advanced Cluster Management policy yaml structure at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/governance/index#policy-yaml-structure

► Guided Exercise

Deploying and Managing Policies with RHACM

In this exercise, you will create a policy in all the clusters to verify the expiration of the certificates in the `openshift-console` and `openshift-ingress` namespaces. Then, you will analyze the status of the policy, apply a fix, and verify the final status after the remediation.

Outcomes

You should be able to:

- Deploy a certificate policy from the RHACM governance dashboard
- Verify the policy status from the RHACM governance dashboard
- Verify the policy status after applying the remediation

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

```
[student@workstation ~]$ lab start policy-governance
```

Instructions

- 1. Log in to the Hub OpenShift cluster and create a `policy-governance` project.

- 1.1. Open the terminal application on the workstation machine. Log in to the Hub OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Create a new project `policy-governance`.

```
[student@workstation ~]$ oc new-project policy-governance
Now using project "policy-governance" on server "https://
api.ocp4.example.com:6443".
...output omitted...
```

- 2. Log in to the RHACM web console and create the certificate policy.

- 2.1. From the workstation machine, open Firefox and access <https://multicloud-console.apps.ocp4.example.com>.
- 2.2. Click `htpasswd_provider` and log in as the `admin` user with the `redhat` password.

- 3. Create the certificate policy with the following parameters for openshift-console and openshift-ingress namespace:

Field Name	Value
Name	policy-certificatepolicy
Namespace	policy-governance
Specifications	CertificatePolicy - Certificate management expiration
Remediation	Inform

- 3.1. Click **Governance** tab on the left pane to navigate to the governance dashboard. Then, click **Create policy** button. The **Create policy** page appears.

The screenshot shows the Red Hat Advanced Cluster Management for Kubernetes interface. The top navigation bar includes the Red Hat logo, the title 'Advanced Cluster Management for Kubernetes', and user information ('admin'). The left sidebar has a dark theme with several sections: 'Home' (Welcome, Overview), 'Infrastructure' (Clusters, Bare metal assets, Automation), 'Applications' (Governance, Credentials), and 'Visual Web Terminal'. The 'Governance' section is currently selected and highlighted with a red box. The main content area is titled 'Governance' and features a large blue sphere icon with smaller spheres and stars. Below the icon, it says 'No policies found' and provides instructions: 'You don't have any policies, yet. Use the [create policy](#) button to create your resource. View the documentation for more information'. At the bottom of this section is another 'Create policy' button, which is also highlighted with a red box. In the top right corner of the main content area, there is a refresh icon and a status message: 'Last update 42 days ago'.

- 3.2. Fill in the fields as follows, leaving rest of the fields unchanged. Do **not** click **Create** button yet.

All fields marked with an asterisk (*) are mandatory.

Name *

Namespace * ②

Specifications * ②

CertificatePolicy - Ce... X
CertificatePolicy X

Cluster selector ②

Begin typing to search for cluster label to select. If not selected, all cluste...

Standards ②

NIST-CSF X
NIST-CSF X

Categories ②

PR.DS Data Security X
PR.DS Data Security X

Controls ②

PR.DS-2 Data-in-transit X
PR.DS-2 Data-in-transit X

3.3. On the right side of the Create policy page, edit the YAML code as follows:

```

spec:
  remediationAction: inform
  disabled: false
  policy-templates:
    - objectDefinition:
        apiVersion: policy.open-cluster-management.io/v1
        kind: CertificatePolicy
        metadata:
          name: policy-certificatepolicy-cert-expiration
        spec:
          namespaceSelector:
            include:
              - default

```

```

- openshift-console
- openshift-ingress
exclude:
- kube-*
remediationAction: inform
severity: low
minimumDuration: 300h

```

With that customization the policy applies to the default, openshift-console, and openshift-ingress namespaces.

Click **Create** button.

- 3.4. In the Governance page, scroll down and click the **policy-certificatepolicy** policy name from the list of policies. Click **Clusters** tab to review the policy details.

Cluster	Compliance	Template	Message	Last Report	History
managed-cluster	Not compliant	policy-certificatepolicy	NonCompliant; 1 certificates expire in less than 300h0m0s: openshift-ingress: wildcard-tls	4 minutes ago	View history
local-cluster	Compliant	policy-certificatepolicy	Compliant	4 minutes ago	View history

The status of the policy is **Not compliant** for the cluster named **managed-cluster**.



Note

Not compliant policy status indicates that one certificate has already expired or has less than 300h time left in the **openshift-ingress** namespace.

4. Renew the ingress controller wildcard certificate of the **managed-cluster**.

- 4.1. Open the **terminal** application on the **workstation** machine and change to the **~/D0480/labs/policy-governance/** directory.

```
[student@workstation ~]$ cd ~/D0480/labs/policy-governance/
```

- 4.2. Log in to the **managed-cluster** as the **admin** user. The APIServer URL is **https://api.ocp4-mng.example.com:6443**.

```
[student@workstation policy-governance]$ oc login -u admin -p redhat \
https://api.ocp4-mng.example.com:6443
Login successful.
...output omitted...
```

- 4.3. Run the `renew_wildcard.sh` script. The script uses an Ansible Playbook to create a new wildcard certificate with an expiration date set to 3650 days from now.

```
[student@workstation policy-governance]$ ./renew_wildcard.sh
...output omitted...
TASK [Create a combined certificate] ****
ok: [localhost]
TASK [Create a hard-link to the combined certificate] ****
ok: [localhost]
PLAY RECAP ****
localhost      : ok=19      changed=7      unreachable=0      failed=0
skipped=3      rescued=0      ignored=0

configmap/wildcard-bundle data updated
secret/wildcard-tls data updated
```

- 4.4. Move to the home directory.

```
[student@workstation policy-governance]$ cd ~
[student@workstation ~]$
```

► **5.** Review the status of the `policy-certificatepolicy` policy.

- 5.1. Click the `Governance → policy-certificatepolicy` policy name from the list of policies.

The screenshot shows the 'policy-certificatepolicy' details page in the RHACM UI. The 'Policy details' section lists the following information:

Name	policy-certificatepolicy	Categories	PR.DS Data Security
Namespace	policy-governance	Controls	PR.DS-2 Data-in-transit
Status	Enabled	Standards	NIST-CSF
Remediation	Inform	Created	30 minutes ago
Cluster violations	0/2	Automation	Configure

The 'Placement' section shows the cluster selector 'matchExpressions:[]' and the number of clusters '2'. The 'Compliance' section shows the status as 'Compliant: managed-cluster, local-cluster'.

The Policy status is now Compliant for both clusters.



Note

It takes around a minute before the router pods restart and use the updated certificate. Also, it takes a few seconds to change the policy status from Not compliant to Complaint.

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish policy-governance
```

This concludes the guided exercise.

Deploying and Configuring the Compliance Operator for Multiple clusters using RHACM

Objectives

After completing this section, you should be able to deploy the compliance operator policy and view compliance reports from multiple clusters by using the command line and the RHACM Governance Dashboard.

Definition of Compliance

Compliance management is a continuous process to ensure that the IT systems comply with different kinds of policies and requirements. The following list shows some of this kinds of requirements:

- Industry standards
- Security standards
- Corporate policies
- Regulatory policies

Describing the RHOCP Compliance Operator

- You can use the Compliance Operator to verify the compliance state of a cluster.
- The Compliance Operator also verifies the compliance of the Kubernetes objects, as well as the nodes running the cluster.
- The Compliance Operator scans and enforce the security policies you provides using OpenSCAP, a NIST-certified tool.

Explaining the Compliance Operator Components

- Namespace
- Operatorgroup
- Subscription

Deploying the Compliance Operator Across the Fleet of Clusters Using the RHACM Compliance Operator Policy

You can use RHACM console to deploy the policy that deploys RHOCP Compliance Operator in all the managed clusters.

Performing Scans with a Scan Policy

You can use RHACM console to perform scans using scan policies across the fleet of managed clusters.

Verifying Conformance of a Scan Policy

You can use RHACM console to perform scans using scan policies across the fleet of managed clusters.

- Scan and scan binding. Example Essential 8 (E8).

Inspecting Reports of a Scan Policy

You can use RHACM console to verify conformance of a scan policy



References

What is compliance management?

<https://www.redhat.com/en/topics/management/what-is-compliance-management>

For more information, refer to the *Compliance operator policy* chapter in the *Red Hat Advanced Cluster Management Governance Guide* at

https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/governance/index#compliance-operator-policy

► Guided Exercise

Deploying and Configuring the Compliance Operator for Multiple clusters using RHACM

In this exercise, you will deploy the compliance operator on all the clusters in the Asia-Pacific (APAC) location to verify conformance to the Essential 8 (E8) standard. You will also review the compliance report on the RHACM web console.

Outcomes

You should be able to:

- Install the compliance operator from the RHACM governance dashboard
- Deploy the E8-scan policy
- Check the compliant scan result from E8 scan

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

```
[student@workstation ~]$ lab start policy-compliance
```

Instructions

► 1. Log in to the Hub OpenShift cluster and create a `policy-compliance` project.

- 1.1. Open the terminal application on the `workstation` machine. Log in to the Hub OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Create a new project named `policy-compliance`.

```
[student@workstation ~]$ oc new-project policy-compliance
Now using project "policy-compliance" on server "https://
api.ocp4.example.com:6443".
...output omitted...
```

► 2. Log in to the RHACM web console and install the compliance operator.

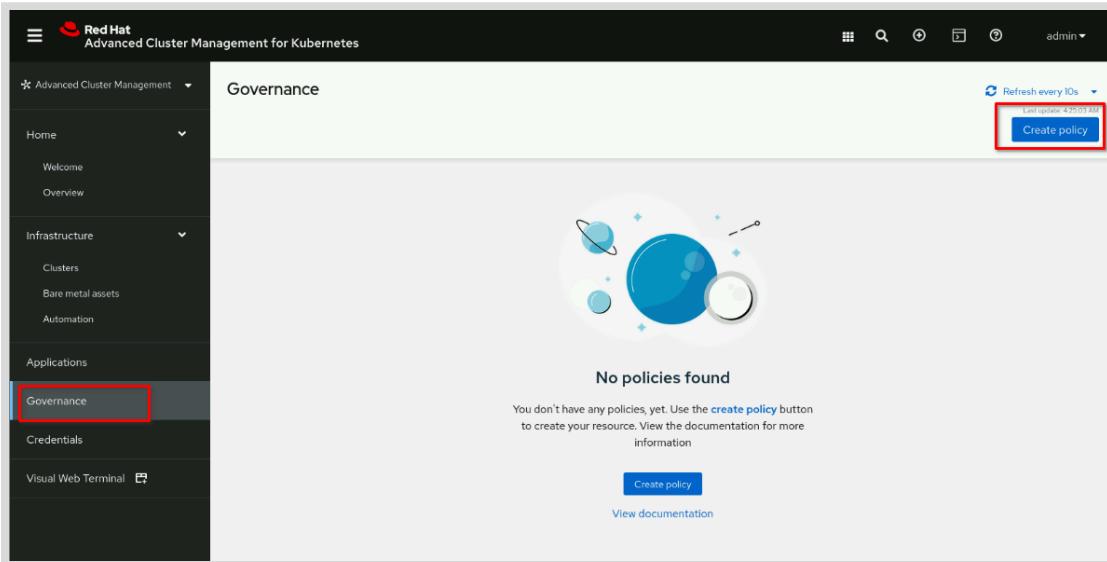
- 2.1. From the `workstation` machine, open Firefox and access <https://multicloud-console.apps.ocp4.example.com>.

2.2. Click `htpasswd_provider` and log in as the `admin` user with the `redhat` password.

► 3. Create the compliance operator policy with the following parameters:

Field Name	Value
Name	policy-complianceoperator
Namespace	policy-compliance
Specifications	ComplianceOperator - Install the Compliance operator
Cluster selector	location: "APAC"
Remediation	Inform

3.1. On the left pane, click **Governance** tab to display the governance dashboard, and then click **Create policy** to create the policy. The **Create policy** page is displayed.



The screenshot shows the Red Hat Advanced Cluster Management for Kubernetes interface. The top navigation bar includes the Red Hat logo, the title 'Advanced Cluster Management for Kubernetes', and a user dropdown for 'admin'. The left sidebar has a dark theme with several sections: 'Home' (selected), 'Welcome', 'Overview', 'Infrastructure' (with 'Clusters', 'Bare metal assets', and 'Automation' sub-options), 'Applications' (with 'Governance' highlighted by a red box), 'Credentials', and 'Visual Web Terminal'. The main content area is titled 'Governance' and features a central graphic of three overlapping circles in blue and white. Below the graphic, the text 'No policies found' is displayed, followed by a message: 'You don't have any policies, yet. Use the [create policy](#) button to create your resource. View the documentation for more information'. At the bottom of the content area are two buttons: 'Create policy' (highlighted with a red box) and 'View documentation'.

3.2. Complete the page with the following details, leaving the other fields unchanged. Then, click **Create**.

- 3.3. Click the **policy-complianceoperator** policy name from the list of policies to check the policy details from the governance dashboard.

Name	policy-complianceoperator
Namespace	policy-compliance
Status	Enabled
Remediation	Inform
Cluster violations	1 / 1
Categories	PRIP Information Protection Processes and Procedures
Controls	PRIP-1 Baseline Configuration
Standards	NIST-CSF
Created	3 minutes ago
Automation	Configure

Placement

Cluster selector	Clusters	Compliance
matchExpressions=[{"key": "location", "operator": "In", "values": ["APAC"]}]	1	Not compliant: managed-cluster

The Policy status shows **Not compliant** for the cluster named **managed-cluster**. The **Not compliant** policy status indicates that the compliance operator is not installed on the **managed-cluster**.

- 3.4. On the left pane, click **Governance** tab. The dashboard shows the **policy-complianceoperator** policy. Click vertical ellipses button to the right of the policy field, and then click **Enforce**. Click **Enforce** again in the confirmation window to change policy mode from **Inform** to **Enforce**.

Enforce mode installs compliance operator to the managed-cluster, and the `policy-complianceoperator` changes to a Compliant status.

- 4. As the admin user, verify the compliance operator installation by using the RHACM web console.

- 4.1. Click the search icon on the top right of the RHACM web console.

- 4.2. Type the following parameters in the search bar to search the compliance-operator subscription.

Field Name	Value
kind	subscription
name	compliance-operator

The search result shows the compliance-operator subscription present in the managed-cluster.

Name	compliance-operator
Namespace	openshift-compliance
Cluster	managed-cluster
Package	
Status	a day ago
Local placement	
Time window	
Created	
Labels	operators.coreos.com/compliance-operator.openshift-compliance

**Note**

APAC location has one cluster, named **managed-cluster**.

- ▶ 5. Clone the **policy collection** repository. The policy-collection repository has policy examples for Open cluster management.
 - 5.1. Open a **terminal** application on the **workstation** machine. Run the following command to clone the policy-collection repository.

```
[student@workstation ~]$ git clone \
https://github.com/open-cluster-management/policy-collection.git
Cloning into 'policy-collection'...
...output omitted...
```

- 5.2. Enter your Git repository clone directory.

```
[student@workstation ~]$ cd policy-collection
```

- ▶ 6. Deploy the Essential 8 (E8) scan policy on the APAC location.
 - 6.1. Edit the **policy-compliance-operator-e8-scan.yaml** YAML file to change the following parameters:

Field Name	Value
remediationAction	enforce
key	location
values	APAC

**Note**

The policy-collection repository has both stable and community policy examples. This course uses the E8 scan policy that is available under the **stable → CM-Configuration-Management** directory.

For more information, please visit <https://github.com/open-cluster-management/policy-collection/blob/main/README.md>.

```
[student@workstation policy-collection]$ cd stable/CM-Configuration-Management
[student@workstation CM-Configuration-Management]$ vim \
policy-compliance-operator-e8-scan.yaml
apiVersion: policy.open-cluster-management.io/v1
kind: Policy
metadata:
  name: policy-e8-scan
  annotations:
    policy.open-cluster-management.io/standards: NIST SP 800-53
    policy.open-cluster-management.io/categories: CM Configuration Management
    policy.open-cluster-management.io/controls: CM-6 Configuration Settings
spec:
  remediationAction: enforce
  disabled: false
  ...output omitted...
  ...output omitted...
spec:
  clusterConditions:
  - status: "True"
    type: ManagedClusterConditionAvailable
  clusterSelector:
    matchExpressions:
    - {key: location, operator: In, values: ["APAC"]}
```

- 6.2. Log in to Hub OpenShift cluster as the **admin** user, and then switch to the **policy-compliance** project.

```
[student@workstation CM-Configuration-Management]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
[student@workstation CM-Configuration-Management]$ oc project policy-compliance
...output omitted...
```

- 6.3. Deploy the E8-scan policy.

```
[student@workstation CM-Configuration-Management]$ oc create -f \
policy-compliance-operator-e8-scan.yaml
policy.policy.open-cluster-management.io/policy-e8-scan created
placementbinding.policy.open-cluster-management.io/binding-policy-e8-scan created
placementrule.apps.open-cluster-management.io/placement-policy-e8-scan created
```

6.4. Move to the home directory.

```
[student@workstation CM-Configuration-Management]$ cd ~
[student@workstation ~]$
```

► 7. Check the compliance operator e8 scan results.

- 7.1. On the left pane, click **Governance** tab. The dashboard shows the **policy-e8-scan** policy. Click **policy-e8-scan** and then select the **Clusters** tab to check the status.
- 7.2. The clusters tab shows three resources: **compliance-e8-scan**, **compliance-suite-e8**, and **compliance-suite-e8-result**. **Compliance-suite-e8-result** will show **Not compliant**.

Cluster	Compliance	Template	Message	Last Report	History
managed-cluster	Not compliant	compliance-suite-e8-results	NonCompliant; violation - compliancecheckresults found: [ocp4-e8-api-server-encryption-provider-cipher, ocp4-e8-ocp-a-d-bpf-disabled, rhcos4-e8-worker-sysctl-kernel-yama-ptrace-scope, rhcos4-e8-worker-sysctl-net-core-bpf-jit-harden] in namespace openshift-compliance	5 minutes ago	View history
managed-cluster	Compliant	compliance-suite-e8	Compliant; notification - compliancesuites [e8] in namespace openshift-compliance found as specified, therefore this Object template is compliant	4 minutes ago	View history
managed-cluster	Compliant	compliance-e8-scan	Compliant; notification - scansettingbindings [e8] in namespace openshift-compliance found as specified, therefore this Object template is compliant	6 minutes ago	View history



Note

It takes 2-3 minutes for **Compliance-suite-e8-result** to show **Not compliant** status.

The **compliance-suit-e8-result** object is non-compliant because the policy applies the **mustnothave** condition to the **compliancecheckresult** object.

- 7.3. Click **View details** and review the details. On the details page, you can see the name of non-compliant objects.

The screenshot shows the RHACM web console interface. On the left, under 'Template details', there is a table with columns: Name, Cluster, Kind, API groups, Compliant, and Details. The 'Name' field is highlighted with a red box and contains the value 'compliance-suite-e8-results'. The 'Details' field also has a red box around its content, which includes JSON-like conditions and a list of compliance check results. On the right, under 'Template yaml', the full YAML configuration is displayed, showing the apiVersion, kind, metadata, and spec sections.

- 7.4. To list the compliance check results, click the search icon on the top right side of the RHACM web console and type kind:configurationpolicy, and then click search.
- 7.5. Click **compliance-suite-e8-result** and check the status. It will show a list of NonCompliant objects.

The screenshot shows the search results for 'compliance-suite-e8-results' in YAML format. It displays a table with 'Cluster' and 'Namespace' columns, both set to 'managed-cluster'. The main area shows the YAML code for the compliance check results, with a red box highlighting the section where it lists various audit rules and their outcomes. The 'Edit' button is visible at the top right of the code editor.

Finish

On the **workstation** machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish policy-compliance
```

This concludes the guided exercise.

Integrating RHACM with Other Policy Engines

Objectives

After completing this section, you should be able to deploy the gatekeeper policy and gatekeeper constraints to multiple clusters by using the command line and the RHACM Governance Dashboard.

Open Policy Agent vs Open Policy Agent Gatekeeper

Open Policy Agent (OPA)

Provides a summary of OPA

OPA Gatekeeper

Provides a summary of OPA Gatekeeper

OPA Gatekeeper Integration with RHACM

Introduction

Gatekeeper Integration with RHACM

Gatekeeper operator installation via RHACM Governance

Install the Gatekeeper operator via RHACM in-built policy.

Deploy Constraint with gatekeeper policy

Components of gatekeeper policy

- ConstraintTemplate and Constraint
- Audit template
- Admission template



References

For more information, refer to the _Red Hat Advanced Cluster Management policy yaml structure at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/governance/index#policy-yaml-structure

For more information, refer to the Managing gatekeeper operator policies section at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/governance/index#managing-gatekeeper-operator-policies

► Guided Exercise

Integrating RHACM with Other Policy Engines

Outcomes

You should be able to:

- Deploy a gatekeeper policy from the RHACM governance dashboard.
- Deploy a gatekeeper policy to exclude given namespace for all constraints.
- Deploy a gatekeeper policy to enforce containers not to use images with the latest tag.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

```
[student@workstation ~]$ lab start policy-gatekeeper
```

Instructions

► 1. Log in to the Hub OpenShift cluster and create a `policy-gatekeeper` project.

- 1.1. Open the terminal application on the `workstation` machine. Log in to the Hub OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Create a new project named `policy-gatekeeper`.

```
[student@workstation ~]$ oc new-project policy-gatekeeper
Now using project "policy-gatekeeper" on server "https://
api.ocp4.example.com:6443".
...output omitted...
```

► 2. Log in to the RHACM web console and install the gatekeeper operator.

- 2.1. From the `workstation` machine, open Firefox and access `https://multicloud-console.apps.ocp4.example.com`.
- 2.2. Click `htpasswd_provider` and log in as the `admin` user with the `redhat` password.

► 3. Create the gatekeeper operator policy with the following parameters:

Field Name	Value
Name	policy-gatekeeperoperator
Namespace	policy-gatekeeper
Specifications	GatekeeperOperator - Install the Gatekeeper operator
Remediation	Enforce

- 3.1. On the left pane, click the **Governance** tab to display the governance dashboard, and then click **Create policy** to create the policy. The **Create policy** page is displayed.
- 3.2. Complete the page with the following details, leaving the other fields unchanged. Then, click **Create**.

The screenshot shows the 'Create policy' interface and the generated Kubernetes manifest. The manifest is as follows:

```

1 apiVersion: policy.open-cluster-management.io/v1
2 kind: Policy
3 metadata:
4   name: policy-gatekeeperoperator
5   namespace: policy-gatekeeper
6 annotations:
7   policy.open-cluster-management.io/standards: NIST-CSF
8   policy.open-cluster-management.io/categories: PRIP
9   policy.open-cluster-management.io/controls: PRIP
10 spec:
11   remediationAction: enforce
12   disabled: false
13   policy-templates:
14     - objectDefinition:
15       apiVersion: policy.open-cluster-management.io/v1
16       kind: ConfigurationPolicy
17       metadata:
18         name: gatekeeper-operator-product-sub
19       spec:
20         remediationAction: inform
21         severity: high
22         object-templates:
23           - complianceType: musthave
24             objectDefinition:
25               apiVersion: operators.coreos.com/v1
26               kind: Subscription
27               metadata:
28                 name: gatekeeper-operator-product
29               namespace: openshift-operators
30             spec:
31               channel: stable
32               installPlanApproval: Automatic
33               name: gatekeeper-operator-product
34               source: redhat-operators
35               sourceNamespace: openshift-marketplace
36     - objectDefinition:
37       apiVersion: policy.open-cluster-management.io/v1
38       kind: ConfigurationPolicy
39       metadata:
40         name: gatekeeper
41       spec:
42         remediationAction: enforce
43         severity: high
44         object-templates:
45           - complianceType: musthave
46             objectDefinition:
47               apiVersion: operators.coreos.com/v1
48               kind: Subscription
49               metadata:
50                 name: gatekeeper
51               namespace: openshift-operators
52             spec:
53               channel: stable
54               installPlanApproval: Automatic
55               name: gatekeeper
56               source: redhat-operators
57               sourceNamespace: openshift-marketplace

```

The 'Create' button is highlighted in the top right of the UI, and the 'Remediation' section in the UI is also highlighted with a red box.

- 3.3. On the **Governance** page, scroll down and click the **policy-gatekeeperoperator** policy name from the list of policies. Click the **Status** tab to review the policy details.

Cluster	Status	Template	Message	Last Report	History
managed-cluster	Compliant	gatekeeper	Compliant; notification - gatekeepers [gatekeeper] found as specified, therefore this Object template is compliant View details	a few seconds ago	View history
managed-cluster	Compliant	gatekeeper-operator-product-sub	Compliant; notification - subscriptions [gatekeeper-operator-product] in namespace openshift-operators found as specified, therefore this Object template is compliant View details	a minute ago	View history
local-cluster	Compliant	gatekeeper	Compliant; notification - gatekeepers [gatekeeper] found as specified, therefore this Object template is compliant View details	a few seconds ago	View history
local-cluster	Compliant	gatekeeper-operator-product-sub	Compliant; notification - subscriptions [gatekeeper-operator-product] in namespace openshift-operators found as specified, therefore this Object template is compliant View details	a minute ago	View history

The the policy status is **Compliant** for all the clusters.



Note

It takes 2–3 minutes to install the gatekeeper operator. The policy status changes from **Not complaint** to **Complaint** after gatekeeper installation.

- ▶ 4. Deploy a gatekeeper policy in the `policy-gatekeeper` namespace to exclude app-stage namespace for all constraints in the stage environment.

- 4.1. Open the `terminal` application on the `workstation` machine and change to the `~/D0480/labs/policy-gatekeeper/` directory.

```
[student@workstation ~]$ cd ~/D0480/labs/policy-gatekeeper/
```

- 4.2. Edit the `policy-gatekeeper-config-exclude-namespaces.yaml` YAML file to change the following parameters:

Field Name	Value
<code>excludedNamespaces</code>	<code>app-stage</code>
<code>key</code>	<code>environment</code>
<code>values</code>	<code>stage</code>

**Note**

The policy-collection repository has both stable and community policy examples. This policy example is from the **community → CM-Configuration-Management** directory.

For more information, visit <https://github.com/open-cluster-management/policy-collection/blob/main/README.md>.

```
[student@workstation policy-gatekeeper]$ vim \
policy-gatekeeper-config-exclude-namespaces.yaml
...output omitted...
apiVersion: config.gatekeeper.sh/v1alpha1
    kind: Config
    metadata:
        name: config
        namespace: openshift-gatekeeper-system
    spec:
        match:
            - excludedNamespaces:
                - hive
                - kube-node-lease
                - kube-public
            ...output omitted...

            ...output omitted...
            - openshift-user-workload-monitoring
            - openshift-vsphere-infra
            - app-stage
        processes:
            - '*'
---
apiVersion: policy.open-cluster-management.io/v1
kind: PlacementBinding
...output omitted...

...output omitted...
clusterSelector:
    matchExpressions:
        - {key: environment, operator: In, values: ["stage"]}
```

- 4.3. Deploy the gatekeeper policy to exclude **app-stage** namespace for all constraints in the stage environment.

```
[student@workstation policy-gatekeeper]$ oc create -f \
policy-gatekeeper-config-exclude-namespaces.yaml
policy.policy.open-cluster-management.io/policy-gatekeeper-config-exclude-
namespaces created
placementbinding.policy.open-cluster-management.io/binding-policy-gatekeeper-
config-exclude-namespaces created
placementrule.apps.open-cluster-management.io/placement-policy-gatekeeper-config-
exclude-namespaces created
```

- 4.4. On the Governance page, scroll down and click the **policy-gatekeeper-config-exclude-namespaces** policy name from the list of policies.

Name	policy-gatekeeper-config-exclude-namespaces	Categories	CM Configuration Management
Namespace	policy-gatekeeper	Controls	CM-2 Baseline Configuration
Remediation	enforce	Standards	NIST SP 800-53
Disabled	false	Created	4 minutes ago
Cluster violations	0/1	Automation	Configure

Placement	
Cluster selector	Status
matchExpressions=[{"key": "environment", "operator": "In", "values": ["stage"]}]	Clusters: 1 Status: Compliant: local-cluster

The the policy status is **Compliant** for stage environment.



Note

The stage environment has one cluster, named **local-cluster**. The production environment has one cluster, named **managed-cluster**.

5. Deploy a gatekeeper policy in the **policy-gatekeeper** namespace to enforce containers not to use images with the latest tag for both stage and production environments.
 - 5.1. Edit the **policy-gatekeeper-container-image-latest.yaml** YAML file to remove **clusterSelector** values.

```
[student@workstation policy-gatekeeper]$ vim \
policy-gatekeeper-container-image-latest.yaml
...output omitted...
spec:
  clusterConditions:
    - status: "True"
      type: ManagedClusterConditionAvailable
  clusterSelector:
    matchExpressions: []
```



Note

The policy-collection repository has both stable and community policy examples. This policy example is from the **community → CM-Configuration-Management** directory.

For more information, visit <https://github.com/open-cluster-management/policy-collection/blob/main/README.md>.

- 5.2. Deploy the gatekeeper policy to enforce containers not to use images with the latest tag for both stage and production environment.

```
[student@workstation policy-gatekeeper]$ oc create -f \
policy-gatekeeper-container-image-latest.yaml
policy.policy.open-cluster-management.io/policy-gatekeeper-containerimagelatest
created
placementbinding.policy.open-cluster-management.io/binding-policy-gatekeeper-
containerimagelatest created
placementrule.apps.open-cluster-management.io/placement-policy-gatekeeper-
containerimagelatest created
```

- 5.3. On the Governance page, scroll down and click the **policy-gatekeeper-containerimagelatest** policy name from the list of policies.

Cluster selector	Clusters	Status
matchExpressions:[]	2	Compliant: managed-cluster, local-cluster

The the policy status is Compliant for both stage and production environments.



Note

It takes 2-3 minutes to for policy to check all object templates. The policy status changes from Not compliant to Compliant after 2-3 minutes.

- ▶ 6. Deploy the stage-hello application to the stage environment by using the ~/D0480/labs/policy-gatekeeper/stage-hello.yaml file.

- 6.1. Log in to the local-cluster as the admin user.

```
[student@workstation policy-gatekeeper]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 6.2. A stage-hello application file is provided to create the namespace, deployment, service, and route. Review the stage-hello.yaml file. This application uses an image with the latest tag

```
[student@workstation policy-gatekeeper]$ cat stage-hello.yaml
apiVersion: project.openshift.io/v1
kind: Project
metadata:
  name: app-stage
spec: {}
---
apiVersion: apps/v1
kind: Deployment
```

```
metadata:
  name: stage-hello
  labels:
    app: stage-hello
    name: stage-hello
  namespace: app-stage
spec:
  replicas: 1
  selector:
    matchLabels:
      app: stage-hello
      name: stage-hello
template:
  metadata:
    labels:
      app: stage-hello
      name: stage-hello
  spec:
    containers:
      - name: stage-hello
        image: quay.io/redhattraining/do480-hello-app:latest
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: stage-hello
    name: stage-hello
  name: stage-hello
  namespace: app-stage
spec:
  ports:
    - port: 8080
  selector:
    name: stage-hello
---
kind: Route
apiVersion: route.openshift.io/v1
metadata:
  name: stage-hello
  labels:
    app: stage-hello
    name: stage-hello
  namespace: app-stage
spec:
  host: stage-hello.apps.ocp4.example.com
  subdomain: ''
  to:
    kind: Service
    name: stage-hello
    weight: 100
  port:
    targetPort: 8080
  wildcardPolicy: None
```

6.3. Deploy the stage-hello application.

```
[student@workstation policy-gatekeeper]$ oc create -f stage-hello.yaml
project.project.openshift.io/app-stage created
deployment.apps/stage-hello created
service/stage-hello created
route.route.openshift.io/stage-hello created
```

6.4. Test the stage-hello application deployment.

```
[student@workstation policy-gatekeeper]$ curl stage-hello.apps.ocp4.example.com
Hello Application!
Image version : latest
```

► 7. Review the policy-gatekeeper-containerimagelatest policy status.

- 7.1. On the Governance page, scroll down and click the `policy-gatekeeper-containerimagelatest` policy name from the list of policies. Click the **Status** tab to review the policy details.

The the policy status is **Compliant** for all the clusters.



Note

The `policy-gatekeeper-containerimagelatest` policy status is **Compliant** despite using the `latest` tag because the `policy-gatekeeper-config-exclude-namespaces` policy excludes the `app-stage` namespace for all constraints in the stage environment.

► 8. Deploy the prod-hello application to the production environment by using the `~/D0480/labs/policy-gatekeeper/prod-hello.yaml` file.

8.1. Log in to the managed-cluster as the `admin` user.

```
[student@workstation policy-gatekeeper]$ oc login -u admin -p redhat \
https://api.ocp4-mng.example.com:6443
Login successful.

...output omitted...
```

- 8.2. A `prod-hello` application file is provided to create the namespace, deployment, service, and route. Review the `prod-hello.yaml` file. This application uses an image with the `latest` tag

```
[student@workstation policy-gatekeeper]$ cat prod-hello.yaml
apiVersion: project.openshift.io/v1
kind: Project
metadata:
  name: app-prod
spec: {}
---
apiVersion: apps/v1
kind: Deployment
```

```
metadata:
  name: prod-hello
  labels:
    app: prod-hello
    name: prod-hello
  namespace: app-prod
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prod-hello
      name: prod-hello
template:
  metadata:
    labels:
      app: prod-hello
      name: prod-hello
  spec:
    containers:
      - name: prod-hello
        image: quay.io/redhattraining/do480-hello-app:latest
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: prod-hello
    name: prod-hello
    name: prod-hello
    namespace: app-prod
spec:
  ports:
    - port: 8080
  selector:
    name: prod-hello
---
kind: Route
apiVersion: route.openshift.io/v1
metadata:
  name: prod-hello
  labels:
    app: prod-hello
    name: prod-hello
    namespace: app-prod
spec:
  host: prod-hello.apps.ocp4-mng.example.com
  subdomain: ''
  to:
    kind: Service
    name: prod-hello
    weight: 100
  port:
    targetPort: 8080
  wildcardPolicy: None
```

8.3. Deploy the prod-hello application.

```
[student@workstation policy-gatekeeper]$ oc create -f prod-hello.yaml
project.project.openshift.io/app-prod created
service/prod-hello created
route.route.openshift.io/prod-hello created
Error from server ([containerimagelatest] Deployment/prod-hello: container 'prod-hello' is using the latest tag for its image (quay.io/redhattraining/do480-hello-app:latest), which is an anti-pattern.
[containerimagelatest] Deployment/prod-hello: container 'prod-hello' is using the latest tag for its image (quay.io/redhattraining/do480-hello-app:latest), which is an anti-pattern.): error when creating "prod-hello.yaml": admission webhook "validation.gatekeeper.sh" denied the request: [containerimagelatest] Deployment/prod-hello: container 'prod-hello' is using the latest tag for its image (quay.io/redhattraining/do480-hello-app:latest), which is an anti-pattern.
[containerimagelatest] Deployment/prod-hello: container 'prod-hello' is using the latest tag for its image (quay.io/redhattraining/do480-hello-app:latest), which is an anti-pattern.
```

We are unable to deploy the application because the deployment is using the latest image tag.

► 9. Review the policy-gatekeeper-containerimagelatest policy status.

- On the Governance page, scroll down and click the **policy-gatekeeper-containerimagelatest** policy name from the list of policies. Click the **Status** tab to review the policy details.

Cluster	Status	Template	Message	Last Report	History	
managed-cluster	Not compliant	policy-gatekeeper-audit-latest	NonCompliant; violation - containerimagelatest not found: [containerimagelatest] found but not as specified	View details	2 minutes ago	View history
managed-cluster	Compliant	policy-gatekeeper-admission-latest	Compliant; notification - events in namespace openshift-gatekeeper-system missing as expected, therefore this Object template is compliant	View details	19 hours ago	View history
managed-cluster	Compliant	policy-gatekeeper-containerimage-latest	Compliant; notification - constrainttemplates [containerimagelatest] found as specified, therefore this Object template is compliant; notification - containerimagelatest [containerimagelatest] found as specified, therefore this Object template is compliant	View details	19 hours ago	View history
local-cluster	Compliant	policy-gatekeeper-admission-latest	Compliant; notification - events in namespace openshift-gatekeeper-system missing as expected, therefore this Object template is compliant	View details	19 hours ago	View history
local-cluster	Compliant	policy-gatekeeper-audit-latest	Compliant; notification - containerimagelatest [containerimagelatest] found as specified, therefore this Object template is compliant	View details	19 hours ago	View history

The status of the **policy-gatekeeper-admission-latest** template is **Not compliant** for **managed-cluster**.

- Click **View details** and review the details. On the details page, scroll down and click **View yaml** for event objects.

The screenshot shows the OpenShift web console with the URL `Policies > policy-gatekeeper-containerimagelatest > Status > policy-gatekeeper-admission-latest`. The page title is "policy-gatekeeper-admission-latest". On the right, there is a large code editor window displaying the YAML configuration for the admission webhook. The status field contains an error message:

```

41 |   fieldsType: FieldsV1
42 |   fieldsV1:
43 |     f:status:
44 |       .: {}
45 |       f:complianceDetails: {}
46 |       f:compliant: {}
47 |       f:relatedObjects: {}
48 |     manager: config-policy-controller
49 |     operation: Update
50 |     time: '2021-12-13T06:25:58Z'
51 |     name: policy-gatekeeper-admission-latest
52 |     namespace: managed-cluster
53 |     ownerReferences:
54 |       - apiVersion: policy.open-cluster-management.io/v1
55 |         blockOwnerDeletion: true
56 |         controller: true
57 |         kind: Policy
58 |         name: policy-gatekeeper.policy-gatekeeper-containerimage
59 |         uid: 93888827-d46f-40ef-b549-e11c165b380
60 |     resourceVersion: '3098865'
61 |     uid: b41f70ef-fc86-4c17-alac-a62a592c4a0d
62 |   spec:
63 |     object-templates:
64 |       - complianceType: mustNotHave
65 |         objectDefinition:
66 |           apiVersion: v1

```

Below the code editor, there is a "Related resources" section with a search bar and a table. The table has columns: Name, Namespace, Kind, API groups, Compliant, Reason, and a "View yaml" button. The "View yaml" button for the row "prod-hello.16c03c7851255041" is highlighted with a red box.

- 9.3. The YAML file displays an error message. The error message indicates that Deployment/prod-hello is using the latest tag.

The screenshot shows the OpenShift web console with the URL `Search > prod-hello.16c03c7851255041`. The "YAML" tab is selected. A red box highlights the error message in the YAML editor:

```

>Error querying for resource: prod-hello.16c03c7851255041
Admission webhook "validation.gatekeeper.sh" denied request, Resource Namespace: app-prod, Constraint: containerimagelatest, Message: Deployment/prod-hello: container 'prod-hello' is using the latest tag for its image (quay.io/redhattraining/do480-hello-app:latest), which is an anti-pattern.

```

- 10. Replace the latest image tag with the v1.0 version. Deploy and test the hello-prod application.

- 10.1. Log in to the managed-cluster as the admin user, and then remove the app-prod project.

```
[student@workstation policy-gatekeeper]$ oc login -u admin -p redhat \
https://api.ocp4-mng.example.com:6443
Login successful.

...output omitted...
[student@workstation policy-gatekeeper]$ oc delete project app-prod
project.project.openshift.io "app-prod" deleted
```

- 10.2. Edit the prod-hello.yaml file to use v1.0 image tag.

```
[student@workstation policy-gatekeeper]$ vim prod-hello.yaml
...output omitted...
spec:
  containers:
    - name: prod-hello
      image: quay.io/redhattraining/do480-hello-app:v1.0
...output omitted...
```

10.3. Deploy the `prod-hello` application.

```
[student@workstation policy-gatekeeper]$ oc create -f prod-hello.yaml
project.project.openshift.io/app-prod created
deployment.apps/prod-hello created
service/prod-hello created
route.route.openshift.io/prod-hello created
```

10.4. Test the `prod-hello` application deployment.

```
[student@workstation policy-gatekeeper]$ curl prod-hello.apps.ocp4-mng.example.com
Hello Application!
Image version : v1.0
```

10.5. Move to the home directory.

```
[student@workstation policy-gatekeeper]$ cd ~
[student@workstation ~]$
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish policy-gatekeeper
```

This concludes the guided exercise.

► Lab

Deploying and Managing Policies with RHACM

Outcomes

You should be able to:

- Deploy a namespace policy from the RHACM governance dashboard.
- Deploy an IAM policy from the RHACM governance dashboard.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

```
[student@workstation ~]$ lab start policy-review
```

Instructions

1. Log in to the Hub OpenShift cluster and create a `policy-review` project.
2. Log in to the RHACM web console and deploy a namespace policy in the `policy-review` namespace to ensure that the `test` namespace does not exist in the production environment.
3. Deploy an IAM policy in the `policy-review` namespace to limit the number of cluster administrators to 2 for all clusters.

Evaluation

As the student user on the workstation machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade policy-review
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish policy-review
```

This concludes the lab.

► Solution

Deploying and Managing Policies with RHACM

Outcomes

You should be able to:

- Deploy a namespace policy from the RHACM governance dashboard.
- Deploy an IAM policy from the RHACM governance dashboard.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

```
[student@workstation ~]$ lab start policy-review
```

Instructions

1. Log in to the Hub OpenShift cluster and create a `policy-review` project.
 - 1.1. Open the `terminal` application on the `workstation` machine. Log in to the Hub OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```
 - 1.2. Create a new project named `policy-review`.

```
[student@workstation ~]$ oc new-project policy-review
Now using project "policy-review" on server "https://api.ocp4.example.com:6443".
...output omitted...
```
2. Log in to the RHACM web console and deploy a namespace policy in the `policy-review` namespace to ensure that the `test` namespace does not exist in the production environment.
 - 2.1. From the `workstation` machine, open Firefox and access `https://multicloud-console.apps.ocp4.example.com`.
 - 2.2. Click `htpasswd_provider` and log in as the `admin` user with the `redhat` password.
 - 2.3. On the left pane, click the `Governance` tab to display the `Governance` dashboard, and then click `Create policy` to create the policy. The `Create policy` page is displayed.
 - 2.4. Complete the page with the following details, leaving the other fields unchanged. Do **not** click `Create` button yet.

Field Name	Value
Name	policy-namespace
Namespace	policy-review
Specifications	Namespace - Must have namespace prod
Cluster selector	environment: "production"
Remediation	Enforce

The screenshot shows the 'Create policy' interface. On the left, there are several dropdown menus and input fields for configuration. On the right, a large code editor displays the generated YAML file for the policy.

```

1  apiVersion: policy.open-cluster-management.io/v1
2  kind: Policy
3  metadata:
4    name: policy-namespace
5    namespace: policy-review
6  annotations:
7    policy.open-cluster-management.io/standards: NIST-CSF
8    policy.open-cluster-management.io/categories: PR.IP Information Protection Processes and Procedures
9    policy.open-cluster-management.io/controls: PR.IP-1 Baseline Configuration
10 spec:
11   remediationAction: enforce
12   disabled: false
13   policy-templates:
14     - objectDefinition:
15       apiVersion: policy.open-cluster-management.io/v1
16       kind: ConfigurationPolicy
17       metadata:
18         name: policy-namespace-prod-ns
19       spec:
20         remediationAction: inform
21         severity: low
22         namespaceSelector:
23           exclude:
24             - kube-*
25           include:
26             - default
27         objectTemplates:
28           complianceType: musthave
29           objectDefinition:
30             kind: Namespace
31             apiVersion: v1
32             metadata:
33               name: prod
34 ...
35  apiVersion: policy.open-cluster-management.io/v1

```

2.5. Use the YAML editor on the right of the **Create policy** page to make the following edits to the YAML file:

```

spec:
  remediationAction: enforce
  disabled: false
  policy-templates:
    - objectDefinition:
        apiVersion: policy.open-cluster-management.io/v1
        kind: ConfigurationPolicy
        metadata:
          name: policy-namespace-prod-ns
      spec:
        remediationAction: inform
        severity: low
        namespaceSelector:
          exclude:
            - kube-*
          include:
            - default
      object-templates:
        - complianceType: mustnothave
          objectDefinition:

```

```
kind: Namespace
apiVersion: v1
metadata:
  name: test
```

Click **Create**.

- 2.6. On the **Governance** page, scroll down and click the **policy-namespace** policy name from the list of policies. Click the **Clusters** tab to review the policy details.

Cluster	Compliance	Template	Message	Last Report	History
managed-cluster	Compliant	policy-namespace-prod-ns	Compliant; notification - namespaces [test] missing as expected, therefore this Object template is compliant	6 minutes ago	View history

The the policy status is **Compliant** for the **managed-cluster**.

3. Deploy an IAM policy in the **policy-review** namespace to limit the number of cluster administrators to 2 for all clusters.
 - 3.1. On the left pane, click the **Governance** tab to display the **Governance** dashboard, and then click **Create policy** to create the policy. The **Create policy** page is displayed.
 - 3.2. Complete the page with the following details, leaving the other fields unchanged. Do **not** click **Create** button yet.

Field Name	Value
Name	policy-iampolicy
Namespace	policy-review
Specifications	IamPolicy - Limit clusteradmin roles

```

1 apiVersion: policy.open-cluster-management.io/v1
2 kind: Policy
3 metadata:
4   name: policy-iampolicy
5   namespace: policy-review
6   annotations:
7     policy.open-cluster-management.io/standards: NIST-CSF
8     policy.open-cluster-management.io/categories: PR.AC Identity M
9     policy.open-cluster-management.io/controls: PR.AC-4 Access Con
10    spec:
11      remediationAction: inform
12      disabled: false
13      policy-templates:
14        - objectDefinition:
15          apiVersion: policy.open-cluster-management.io/v1
16          kind: IamPolicy
17          metadata:
18            name: policy-iampolicy-limit-clusteradmin
19          spec:
20            severity: medium
21            namespaceSelector:
22              include:
23                - '*'
24              exclude:
25                - kube-*
26                - openshift-
27            remediationAction: inform
28            maxClusterRoleBindingUsers: 5
29 ...
30          apiVersion: policy.open-cluster-management.io/v1
31          kind: PlacementBinding
32          metadata:
33            name: binding-policy-iampolicy
34            namespace: policy-review
35          placementRef:

```

- 3.3. Use the YAML editor on the right of the **Create policy** page to make the following edits to the YAML file:

```

spec:
  remediationAction: inform
  disabled: false
  policy-templates:
    - objectDefinition:
        apiVersion: policy.open-cluster-management.io/v1
        kind: IamPolicy
        metadata:
          name: policy-iampolicy-limit-clusteradmin
        spec:
          severity: medium
          namespaceSelector:
            include:
              - '*'
            exclude:
              - kube-*
              - openshift-
        remediationAction: inform
        maxClusterRoleBindingUsers: 2

```

Click **Create** button.

- 3.4. On the **Governance** page, scroll down and click the **policy-iampolicy** policy name from the list of policies. Click the **Clusters** tab to review the policy details.

The screenshot shows the 'Clusters' tab of the 'policy-iampolicy' policy in the RHACM UI. It lists two clusters: 'local-cluster' and 'managed-cluster'. The 'Compliance' column indicates the status: 'Not compliant' for local-cluster and 'Compliant' for managed-cluster. The 'Message' column provides details: 'NonCompliant; The number of users with the cluster-admin role is at least 2 above the specified limit' for local-cluster, and 'Compliant; The number of users with the cluster-admin role is at least 0 above the specified limit' for managed-cluster. A red box highlights the 'Not compliant' status for local-cluster.

The status of the **policy-iampolicy-limit-clusteradmin** template is **Not compliant** for **local-cluster** and **Compliant** for **managed-cluster**.

Evaluation

As the student user on the **workstation** machine, use the **lab** command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade policy-review
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish policy-review
```

This concludes the lab.

Summary



Note

This section is under development.

Chapter 4

Configuring the Observability Stack in RHACM

Goal

Gain insight into the fleet of managed clusters by using RHACM observability components.

Objectives

- Describe the architecture of the observability stack and the benefits of observability in a multicluster environment.
- Customize the RHACM observability stack.

Sections

- Introducing the RHACM Observability Stack (and Guided Exercise)
- Customizing the RHACM Observability Stack (and Guided Exercise)

Lab

Installing and Customizing the RHACM Observability Stack

Introducing the RHACM Observability Stack

Objectives

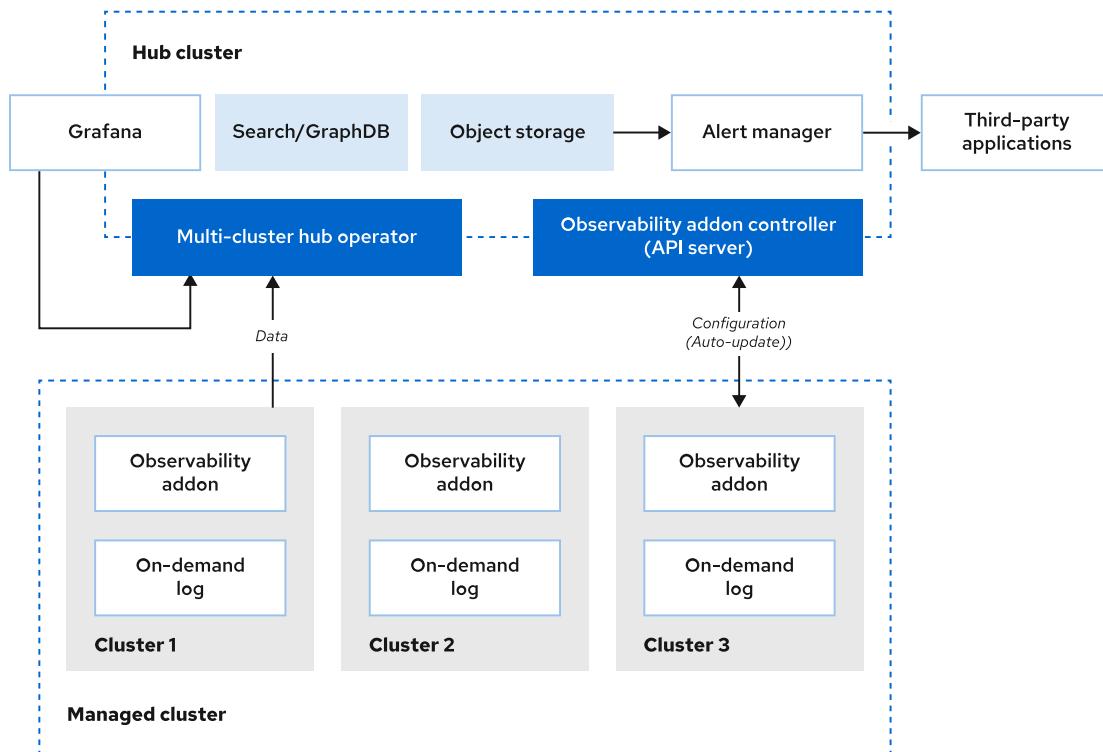
After completing this section, you should be able to describe the architecture of the observability stack and the benefits of observability in a multicluster environment.

Multicluster Observability in RHACM

The observability service in RHACM provides administrators a view across the fleet of clusters from a single window.

In production environments, monitoring, processing, and storing metrics from the fleet of clusters is crucial to anticipating issues, improving the performance of the clusters, and conducting post-mortem analysis.

The following diagram provides an architecture overview of the observability components in RHACM.



When you enable the observability service, the `multicluster-observability-operator` deploys Grafana and Alertmanager instances in the hub cluster. If not specified otherwise, it also deploys the observability addon and on-demand log components in each managed cluster.

The Grafana and Alertmanager instances deployed in the hub cluster receive, process, and display alerts and metrics from all the managed clusters with the observability addon activated.

You can add custom dashboards to Grafana. Also, as explained elsewhere in this chapter, you can create custom recording rules and alerting rules. The Alertmanager instance manages alert forwarding to third-party applications.

Differences Between the RHOCP Monitoring Stack and the RHACM Observability Service

The Grafana and Alertmanager instances provided by RHACM compliment the RHOCP cluster-level metrics.

The RHOCP monitoring stack provides out-of-the-box monitoring for core platform components. Starting with RHOCP version 4.6, you can also monitor your own projects.

The RHACM observability service is designed to provide cluster-level metrics, receiving those from the fleet of clusters. You can use the Grafana instance of the RHACM observability service to monitor cluster-level metrics from across the fleet of clusters.

Enabling the RHACM Observability Service

Due to the requirement for persistent storage, the RHACM observability service is not enabled by default. The observability service supports the following S3 compatible object stores:

- Amazon S3
- Red Hat Ceph
- Google Cloud Storage
- Azure Storage
- Red Hat OpenShift Container Storage
- Red Hat OpenShift on IBM (ROKS)

You must create a `MultiClusterObservability` custom resource in the `open-cluster-management-observability` namespace to enable the observability service. The `MultiClusterObservability` object requires a secret containing the credentials of the S3 bucket storage.



Note

For detailed instructions about enabling the observability service, refer to the *Enable observability service* section at https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/observability/index#enable-observability



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Observability Guide* at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/observability/index

► Guided Exercise

Introducing the RHACM Observability Stack

In this exercise you will deploy the observability stack in all the managed clusters in RHACM and verify the installation in the hub cluster and the managed clusters. Then, you will explore the Grafana dashboard provided by the RHACM observability stack. Finally, you will disable the metrics collection in a cluster used for the stage environment.

Outcomes

You should be able to:

- Install the observability stack in RHACM
- Verify the installation in all the managed clusters
- Access the RHACM Grafana dashboard
- Disable the observability agent in the managed-cluster

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is deployed in the hub cluster, and the `managed-cluster` exists in RHACM.

```
[student@workstation ~]$ lab start observability-install
```

Instructions

► 1. Enable the observability service from the terminal.

- 1.1. Open a terminal and log in to the `ocp4` cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 1.2. Create the namespace `open-cluster-management-observability`.

```
[student@workstation ~]$ oc create namespace open-cluster-management-observability
namespace/open-cluster-management-observability created
```

- 1.3. Extract the pull-secret from the `openshift-config` namespace. Store it in a variable.

```
[student@workstation ~]$ DOCKER_CONFIG_JSON=`oc extract secret/pull-secret \
-n openshift-config --to=-` \
# .dockerconfigjson
[student@workstation ~]$
```

- 1.4. Create the pull-secret in the namespace `open-cluster-management-observability`.

```
[student@workstation ~]$ oc create secret generic \
multicloudhub-operator-pull-secret \
-n open-cluster-management-observability \
--from-literal=.dockerconfigjson="$DOCKER_CONFIG_JSON" \
--type=kubernetes.io/dockerconfigjson
secret/multicloudhub-operator-pull-secret created
[student@workstation ~]$
```

- 1.5. Use your text editor to create a YAML file with the definition of a new object bucket claim (OBC).

```
[student@workstation ~]$ vi obc.yaml
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: thanos-bc
  namespace: open-cluster-management-observability
spec:
  storageClassName: openshift-storage.noobaa.io
  generateBucketName: observability-bucket
```

Save the file and close the editor.

Now, create the `ObjectBucketClaim` using the `oc` command.

```
[student@workstation ~]$ oc create -f obc.yaml
objectbucketclaim.objectbucket.io/thanos-bc created
```

- 1.6. Extract the storage configuration details and the secret.

```
[student@workstation ~]$ oc extract configmap/thanos-bc \
-n open-cluster-management-observability --to=- \
# BUCKET_HOST
s3.openshift-storage.svc
# BUCKET_NAME
observability-bucket-eabe4aca-0d93-49b2-a3dc-ce8e902ec48b
# BUCKET_PORT
443
# BUCKET_REGION

# BUCKET_SUBREGION
```

```
[student@workstation ~]$ oc extract secret/thanos-bc \
-n open-cluster-management-observability --to=-
# AWS_SECRET_ACCESS_KEY
2wNow7IHpJsSxUR9c+OU1CmCIJVkChSe0ZoFrT7V
# AWS_ACCESS_KEY_ID
T6mbdR5Hg6Gf7U2PpXzE
```

- 1.7. Create a secret containingining the `thanos.yaml` configuration file with the information from the previous step.

```
[student@workstation ~]$ vi secret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: thanos-object-storage
  namespace: open-cluster-management-observability
type: Opaque
stringData:
  thanos.yaml: |
    type: s3
    config:
      bucket: observability-bucket-eabe4aca-0d93-49b2-a3dc-ce8e902ec48b
      endpoint: s3.openshift-storage.svc
      insecure: true
      access_key: T6mbdR5Hg6Gf7U2PpXzE
      secret_key: 2wNow7IHpJsSxUR9c+OU1CmCIJVkChSe0ZoFrT7V
```

Use the `oc` command to create the secret in RHOCP.

```
[student@workstation ~]$ oc create -f secret.yaml
secret/thanos-object-storage created
```

- 1.8. Create the multicluster observability object to enable the observability service.

```
[student@workstation ~]$ vi mcobs.yaml
apiVersion: observability.open-cluster-management.io/v1beta2
kind: MultiClusterObservability
metadata:
  name: observability
spec:
  observabilityAddonSpec: {}
  storageConfig:
    metricObjectStorage:
      name: thanos-object-storage
      key: thanos.yaml
```

Use the `oc` command to create the multicluster observability instance in RHOCP.

```
[student@workstation ~]$ oc create -f mcobs.yaml
multiclusterobservability.observability.open-cluster-management.io/observability
created
```

At this point, the observability components deployment starts. You can check the progress of the installation using the following command:

```
[student@workstation ~]$ oc describe mco observability
Name:      observability

...output omitted...

Status:
Conditions:
  Last Transition Time: 2021-11-09T13:34:59Z
  Message:            Installation is in progress
  Reason:             Installing
  Status:              True
  Type:                Installing

...output omitted...
```

- ▶ 2. Verify the deployment of the observability components in all the managed clusters.
 - 2.1. Verify the completion of the installation in the ocp4 cluster, named local-cluster in RHACM.

After a few minutes installing, the status of the observability components transitions to Ready.

```
[student@workstation ~]$ oc describe mco observability
Name:      observability

...output omitted...

Status:
Conditions:
  Last Transition Time: 2021-11-09T13:34:59Z
  Message:            Installation is in progress
  Reason:             Installing
  Status:              True
  Type:                Installing
  Last Transition Time: 2021-11-09T13:36:32Z
  Message:            Observability components are deployed and running
  Reason:             Ready
  Status:              True
  Type:                Ready
Events:    <none>
```

- 2.2. Verify the status of the components of the observability addon in the cluster named managed-cluster in RHACM.
- First, log in to the ocp4-mng cluster.

```
[student@workstation ~]$ oc login -u admin -p redhat https://api.ocp4-
mng.example.com:6443
...output omitted...
```

Then, retrieve the objects of the namespace `open-cluster-management-addon-observability`.

```
[student@workstation ~]$ oc get all -n open-cluster-management-addon-observability
NAME                                     READY   STATUS    RESTARTS
AGE
pod/endpoint-observability-operator-7dd7595ff8-hn8qc   1/1     Running   0
  13m
pod/metrics-collector-deployment-7c5cf96599-cfh5s      1/1     Running   0
  12m

NAME                               READY   UP-TO-DATE   AVAILABLE
AGE
deployment.apps/endpoint-observability-operator   1/1     1           1
  13m
deployment.apps/metrics-collector-deployment     1/1     1           1
  12m

NAME          DESIRED   CURRENT
READY   AGE
replicaset.apps/endpoint-observability-operator-7dd7595ff8   1        1        1
  13m
replicaset.apps/metrics-collector-deployment-546d85d5c6       0        0        0
  12m
replicaset.apps/metrics-collector-deployment-74d5469859       0        0        0
  12m
replicaset.apps/metrics-collector-deployment-7c5cf96599       1        1        1
  12m
replicaset.apps/metrics-collector-deployment-868984fb6d       0        0        0
  12m
```



Note

If the status of the pods is not **Active**, look for error messages in the logs of the pods, or the events in the namespace.

```
[student@workstation ~]$ oc logs pod/metrics-collector-deployment-7c5cf96599-cfh5s -n open-cluster-management-addon-observability
```

```
[student@workstation ~]$ oc get events -n open-cluster-management-addon-observability
```

- ▶ 3. Navigate to the Grafana dashboard available in the **Home → Overview** menu of the RHACM console. Analyze the capacity and utilization data provided about CPU and memory of the clusters.



Note

After enabling the observability components, RHACM shows direct access to the Grafana dashboard from within the RHACM console.

- 3.1. Open Firefox and navigate to the RHACM console in `multicloud-console.apps.ocp4.example.com`.

When prompted, select `htpasswd_provider`

Type the credentials:

- Username: `admin`
- Password: `redhat`

And then click **Log in**

From within the RHACM web console, click **Home → Overview**.

Finally, click **Grafana**

The screenshot shows the Red Hat Advanced Cluster Management for Kubernetes web interface. The left sidebar has a dark theme with categories like 'Home', 'Infrastructure', and 'Applications'. Under 'Home', 'Overview' is selected. The main content area is titled 'Overview' and features a large cloud icon. Below it, the word 'Other' is displayed, followed by a number '2' and the word 'Cluster'. In the top right corner of the main area, there is a 'Grafana' button, a 'Add provider connection' button, and a 'Refresh every 1m' button. The URL in the browser bar is `https://multicloud-console.apps.ocp4.example.com/grafana/d/2b679d600f3b9e7676a7c5ac3643d448/acm-clusters-overview`.



Note

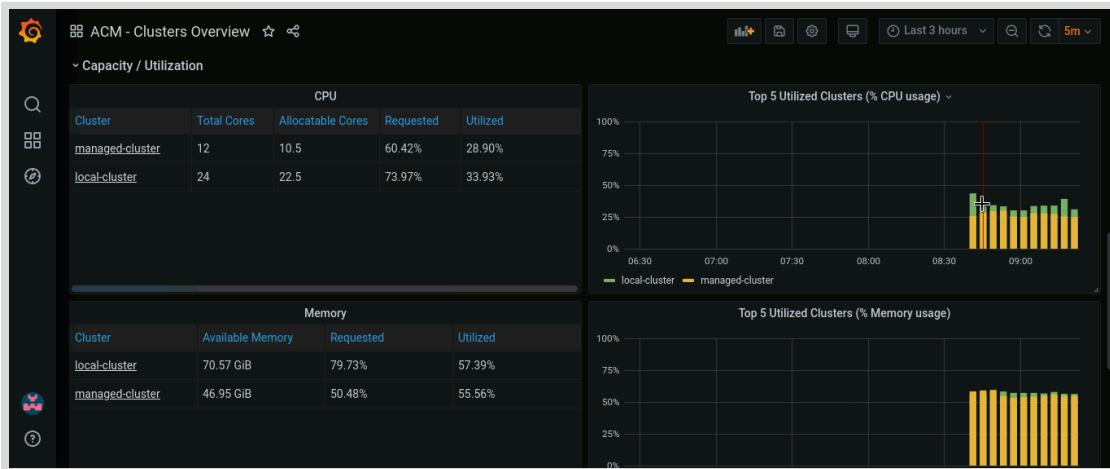
Reload the Grafana dashboard if the panels show any error message.

Finally, you can see the Grafana dashboard with the default configuration.

The screenshot shows a Grafana dashboard titled 'ACM - Clusters Overview'. The dashboard is divided into several sections. On the left, there are navigation icons for search, refresh, and help. The main content area starts with a section titled 'Control Plane Health' which contains two tables: one for 'API Server' and one for 'etcd'. Both tables compare two clusters: 'local-cluster' and 'managed-cluster'. The 'API Server' table shows metrics like Max latency (99th perc), API servers up, and API Errors [1h]. The 'etcd' table shows metrics like Has a Leader, Leader Election Chang, and DB Size. Below this is a section titled 'Optimization' with two more tables: one for 'CPU' and one for 'Memory'. These tables show metrics like Overestimation, Requested, Utilized, and other performance indicators for each cluster. The overall layout is clean with dark-themed cards and light-colored tables.

The Grafana dashboard of RHACM contains information of all the clusters in the fleet. The observability addon, deployed in all the managed clusters, sends information to the central Grafana instance deployed in the hub cluster.

- 3.2. From within the Grafana dashboard, scroll down to find the Capacity / Utilization table.



Grafana gathers information about CPU and memory utilization across all the clusters in the fleet. This way, a system administrator has a single pane of glass to monitor all the clusters from the same dashboard.

- ▶ 4. Disable the metrics collection from the `managed-cluster`. This is useful to avoid overloading the observability service with metrics from clusters used for stage environments.
- 4.1. Open a terminal and log in to the `ocp4` cluster as the `admin` user. The APIServer URL is <https://api.ocp4.example.com:6443>.

```
[student@workstation ~]$ oc login -u admin -p redhat \
  https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 4.2. Add the label `observability=disabled` to the `managedcluster` object named `managed-cluster`.

```
[student@workstation ~]$ oc label managedcluster managed-cluster \
  observability=disabled -n open-cluster-management
managedcluster.cluster.open-cluster-management.io/managed-cluster labeled
```

After applying this label, RHACM removes the objects in the `open-cluster-management-addon-observability` namespace.

- 4.3. Verify that the changes in the `managed-cluster` are applied. Log in to the `ocp4-mng` cluster as the `admin` user. The APIServer URL is <https://api.ocp4-mng.example.com:6443>.

```
[student@workstation ~]$ oc login -u admin -p redhat \
  https://api.ocp4-mng.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 4.4. Retrieve the objects in the `open-cluster-management-addon-observability` namespace.

```
[student@workstation ~]$ oc get all \
-n open-cluster-management-observability-addon
No resources found in open-cluster-management-observability-addon namespace.
```

The expected result is that there are not resources found in the `open-cluster-management-observability-addon` namespace.



Note

The Grafana dashboard can show cached information about the `managed-cluster` after disabling the observability addon. After some time, all the information about the `managed-cluster` disappears.

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish observability-install
```

This concludes the guided exercise.

Customizing the RHACM Observability Stack

Objectives

After completing this section, you should be able to customize the RHACM observability stack.

Observability Service Customizations in RHACM

To benefit from the RHACM observability stack, you can add customizations to the Grafana and Alertmanager instances.

Prometheus recording rules

Recording rules allow you to precalculate expensive expressions. The result is saved as a new set of time series that you can use to create custom Grafana dashboards.



Note

For more information, visit the Prometheus recording rules documentation at https://prometheus.io/docs/prometheus/latest/configuration/recording_rules/

Prometheus alerting rules

Alerting rules allow you to receive alerts based on determined conditions.



Note

For more information, visit the Prometheus alerting rules documentation at https://prometheus.io/docs/prometheus/latest/configuration/alerting_rules/

Prometheus custom metrics

You can create additional custom metrics to be collected from all the managed clusters.



Note

For more information, review the *Adding custom metrics* section at https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/observability/index#adding-custom-metrics

Forwarding Metrics to External Services

Typically, organizations have a centralized alerting engine to receive alerts from different systems.

To facilitate this centralized alerting, you can configure extra alert receivers as explained in the Prometheus Alertmanager documentation at <https://prometheus.io/docs/alerting/latest/configuration/>.



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Observability Guide* at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.3/html-single/observability/index

► Guided Exercise

Customizing the RHACM Observability Stack

In this exercise you will configure an existing RHACM observability service to handle a higher load of monitoring data from the managed clusters. Then, you will create a custom alerting rule that fires when the CPU requests of a cluster are above 70% of the total allocatable CPU. Finally, you will disable the RHACM observability service.

Outcomes

You should be able to:

- Increase the replicas of the observability metrics receiver pods
- Verify that the number of pods increases
- Create a prometheus custom alerting rule
- Verify that the alerts are firing
- Remove the observability components from RHACM

Before You Begin

Before attempting this guided exercise, ensure you have completed the guided exercise **Introducing the RHACM Observability Stack**.

As the student user on the **workstation** machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is deployed in the hub cluster, the `managed-cluster` exists in RHACM, and the observability service is enabled.

```
[student@workstation ~]$ lab start observability-customize
```

Instructions

- 1. From the terminal, increase to six the replicas of the observability metrics receiver pods.
- 1.1. Open a terminal and log in to the `ocp4` cluster as the `admin` user. The APIServer URL is <https://api.ocp4.example.com:6443>.

```
[student@workstation ~]$ oc login -u admin -p redhat \
  https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 1.2. Use the `oc` command to edit the `multiclusterobservability` object in the namespace `openshift-multicloud-observability`. In the `spec` section, add the highlighted snippet to set the number of replicas of the receiver pods to six.

```
[student@workstation ~]$ oc edit multicloudobservability \
-n openshift-multicloud-observability
...output omitted...
spec:
  advanced:
    receive:
      replicas: 6
enableDownsampling: true
```

Save the file and close the editor.



Note

After exiting the editor, the following message displays:

```
multicloudobservability.observability.open-cluster-management.io/
observability edited
```

- 1.3. Use the `watch` command to verify that there are six pods of the `observability-thanos-receive-default` statefulset.

```
[student@workstation ~]$ watch oc get pods \
-n open-cluster-management-observability
...output omitted...
NAME                                     READY   STATUS
RESTARTS   AGE
...output omitted...
observability-thanos-receive-default-0     1/1    Running   0
      55m
observability-thanos-receive-default-1     1/1    Running   0
      54m
observability-thanos-receive-default-2     1/1    Running   0
      54m
observability-thanos-receive-default-3     1/1    Running   0
      2m56s
observability-thanos-receive-default-4     1/1    Running   0
      2m51s
observability-thanos-receive-default-5     1/1    Running   0
      2m46s
...output omitted...
```

Notice that three of the pods are created recently, and the other three were already running.



Note

Increasing the replicas of the receiver pods is necessary in production environments when the number of managed clusters increases.

- ▶ 2. Create a Prometheus alerting rule so that cluster administrators receive alerts when the CPU requests of any cluster are above 70% of the available compute capacity.
 - 2.1. From the terminal, create a file named `custom_rule.yaml` containing the instructions in PromQL.

```
[student@workstation ~]$ vi custom-rules.yaml
kind: ConfigMap
apiVersion: v1
metadata:
  name: thanos-ruler-custom-rules
data:
  custom_rules.yaml: |
    groups:
      - name: cluster-health
        rules:
          - alert: ClusterCPUReq-70
            annotations:
              summary: Notify when CPU requests on a cluster are greater than the
              defined utilization limit
              description: "The cluster {{ $labels.cluster }} has an elevated
              percentage of CPU requests: {{ $labels.clusterID }}."
            expr: |
              sum(namespace_cpu:kube_pod_container_resource_requests:sum) by
              (clusterID, cluster, prometheus) /
              sum(kube_node_status_allocatable{resource="cpu"}) by (clusterID, cluster,
              prometheus) > 0.7
            for: 5s
            labels:
              cluster: "{{ $labels.cluster }}"
              prometheus: "{{ $labels.prometheus }}"
            severity: critical
```

- 2.2. Create the `thanos-ruler-custom-rules` configmap in the `open-cluster-management-observability` namespace.

```
[student@workstation ~]$ oc apply -f custom-rules.yaml \
  -n open-cluster-management-observability
configmap/thanos-ruler-custom-rules created
```

The creation of the configmap triggers a restart of the `observability-thanos-rule-X` pods in the `open-cluster-management-observability` namespace.

- 2.3. Verify that the new configuration applies. You can monitor the restart of the pods in the `open-cluster-management-observability` namespace with the `watch` command.

```
[student@workstation ~]$ watch oc get pods \
-n open-cluster-management-observability
...output omitted...
NAME                                READY   STATUS
RESTARTS   AGE
...output omitted...
observability-thanos-rule-0          2/2    Running   0
  1m
observability-thanos-rule-1          2/2    Running   0
  1m
observability-thanos-rule-2          2/2    Running   0
  2m
...output omitted...
```

Exit the `watch` command with `Ctrl+C`.

After the `observability-thanos-rule-X` pods restart, the alerting rule is visible from the RHACM Grafana dashboard.

► 3. Verify that the custom alert is firing for the cluster named `local-cluster`.

- 3.1. Use Firefox to navigate to the RHACM web console at <https://multicloud-console.apps.ocp4.example.com>.

When prompted, select `htpasswd_provider`

Type the credentials:

- Username: `admin`
- Password: `redhat`

And then click **Log in**

- 3.2. Navigate to Home → Overview and click **Grafana**.

The screenshot shows the Red Hat Advanced Cluster Management for Kubernetes web interface. The left sidebar has a navigation menu with 'Home' selected, which further has 'Overview' selected. The main content area displays a summary card for 'Other' resources, showing the number '2' and the category 'Cluster'. In the top right corner, there is a 'Grafana' link, a 'Add provider connection' button, and a 'Refresh every 1m' button. The URL in the browser bar is <https://multicloud-console.apps.ocp4.example.com/grafana/d/2b679d600f3b9e7676a7c5ac3643d448/acm-clusters-overview>.

A new tab showing the Grafana dashboard opens in the browser.

- 3.3. On the left menu of the Grafana dashboard, click **Explore**

The screenshot shows the ACM - Clusters Overview dashboard. In the Control Plane Health section, there are two tables: 'API Server' and 'etcd'. The 'API Server' table has two rows: 'Explore' and 'local-cluster'. The 'etcd' table has two rows: 'local-cluster' and 'managed-cluster'. In the Optimization section, there are two tables: 'CPU' and 'Memory'. The 'CPU' table has two rows: 'local-cluster' and 'managed-cluster'. The 'Memory' table has two rows: 'local-cluster' and 'managed-cluster'.

- 3.4. In the Metrics text field, type ALERTS{alertname="ClusterCPUReq-70", alertstate="firing"} and click Run Query.

The screenshot shows the Grafana interface with a single query in the Metrics panel: `ALERTS{alertname="ClusterCPUReq-70", alertstate="firing"}`. The query type is set to 'Both'. The results are displayed in a Graph panel, which shows a single data series with a value of 1.000 across the time range from 12:00 to 12:55.

Scroll down to the Table field and verify that the alert with name ClusterCPUReq-70 is firing in the cluster named local-cluster.

The screenshot shows the Grafana interface with the results of the previous query in a Table panel. The table has columns: Time, _name_, alertname, alertstate, cluster, clusterID, prometheus, and severity. There is one row of data: 2021-11-16 12:57:04, ALERTS, ClusterCPUReq-70, firing, local-cluster, 61f2307c-7b3b-49fe-abaf-075799dec734, {{ \$labels.prometheus }}, critical.

- 4. Disable the RHACM observability service.

- 4.1. Open a terminal and log in to the ocp4 cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
  https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 4.2. Use the `oc` to remove the `MultiClusterObservability` object named `observability`.

```
[student@workstation ~]$ oc delete multicluserobervability \
  observability
multicluserobervability.observability.open-cluster-management.io "observability"
deleted
```

- 4.3. Remove the `open-cluster-management-observability` namespace.

```
[student@workstation ~]$ oc delete namespace \
  open-cluster-management-observability
namespace "open-cluster-management-observability" deleted
```



Note

The deletion of the `open-cluster-management-observability` namespace can take a while.

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish observability-customize
```

This concludes the guided exercise.

► Lab

Installing and Customizing the RHACM Observability Stack

In this lab, you will enable the RHACM observability service in all the managed clusters. Then, you will create a custom alerting rule that fires a warning when the memory usage in a cluster exceeds 50%. You will verify that the alert is firing for every cluster that meets the condition. Finally, you will disable the observability service.

Outcomes

You should be able to:

- Install the observability service in RHACM
- Create a Prometheus custom alerting rule
- Verify that the alert is firing in the Grafana dashboard
- Remove the observability components from RHACM

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is deployed in the hub cluster, the `managed-cluster` exists in RHACM, and the observability service is not enabled.

```
[student@workstation ~]$ lab start observability-review
```

Instructions

1. Enable the observability service from the terminal.
2. Create a Prometheus alerting rule so that cluster administrators receive alerts when the memory usage of a cluster exceeds 50%.
3. From Grafana, verify that the custom alert is firing for the clusters `local-cluster` and `managed-cluster`.
4. Disable the RHACM observability service.

Evaluation

As the student user on the workstation machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade observability-review
```

Finish

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish observability-review
```

This concludes the lab.

► Solution

Installing and Customizing the RHACM Observability Stack

In this lab, you will enable the RHACM observability service in all the managed clusters. Then, you will create a custom alerting rule that fires a warning when the memory usage in a cluster exceeds 50%. You will verify that the alert is firing for every cluster that meets the condition. Finally, you will disable the observability service.

Outcomes

You should be able to:

- Install the observability service in RHACM
- Create a Prometheus custom alerting rule
- Verify that the alert is firing in the Grafana dashboard
- Remove the observability components from RHACM

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your system for this exercise.

This command ensures that RHACM is deployed in the hub cluster, the `managed-cluster` exists in RHACM, and the observability service is not enabled.

```
[student@workstation ~]$ lab start observability-review
```

Instructions

1. Enable the observability service from the terminal.

- 1.1. Open a terminal and log in to the `ocp4` cluster as the `admin` user. The API Server URL is `https://api.ocp4.example.com:6443`.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 1.2. Create the `open-cluster-management-observability` namespace.

```
[student@workstation ~]$ oc create namespace open-cluster-management-observability
namespace/open-cluster-management-observability created
```

- 1.3. Extract the pull secret from the openshift-config namespace. Store it in a variable.

```
[student@workstation ~]$ DOCKER_CONFIG_JSON=`oc extract secret/pull-secret \
-n openshift-config --to=-` \
# .dockerconfigjson
[student@workstation ~]$
```

- 1.4. Create the pull secret in the namespace open-cluster-management-observability.

```
[student@workstation ~]$ oc create secret generic \
multicloud-operator-pull-secret \
-n open-cluster-management-observability \
--from-literal=.dockerconfigjson="$DOCKER_CONFIG_JSON" \
--type=kubernetes.io/dockerconfigjson
secret/multicloud-operator-pull-secret created
[student@workstation ~]$
```

- 1.5. Use your text editor to create a YAML file with a new object bucket claim (OBC) definition.

```
[student@workstation ~]$ vi obc.yaml
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: thanos-bc
  namespace: open-cluster-management-observability
spec:
  storageClassName: openshift-storage.noobaa.io
  generateBucketName: observability-bucket
```

Save the file and close the editor.

Now, create the ObjectBucketClaim object by using the oc command.

```
[student@workstation ~]$ oc create -f obc.yaml
objectbucketclaim.objectbucket.io/thanos-bc created
```

- 1.6. Extract the storage configuration details and the secret.

```
[student@workstation ~]$ oc extract configmap/thanos-bc \
-n open-cluster-management-observability --to=- \
# BUCKET_HOST
s3.openshift-storage.svc
# BUCKET_NAME
observability-bucket-eabe4aca-0d93-49b2-a3dc-ce8e902ec48b
# BUCKET_PORT
443
# BUCKET_REGION

# BUCKET_SUBREGION
```

```
[student@workstation ~]$ oc extract secret/thanos-bc \
-n open-cluster-management-observability --to=-
# AWS_SECRET_ACCESS_KEY
2wNow7IHpJsSxUR9c+OU1CmCIJVkChSe0ZoFrT7V
# AWS_ACCESS_KEY_ID
T6mbdR5Hg6Gf7U2PpXzE
```

- 1.7. Create a secret containing the `thanos.yaml` configuration file with the information from the previous step.

```
[student@workstation ~]$ vi secret.yaml
apiVersion: v1
kind: Secret
metadata:
  name: thanos-object-storage
  namespace: open-cluster-management-observability
type: Opaque
stringData:
  thanos.yaml: |
    type: s3
    config:
      bucket: observability-bucket-eabe4aca-0d93-49b2-a3dc-ce8e902ec48b
      endpoint: s3.openshift-storage.svc
      insecure: true
      access_key: T6mbdR5Hg6Gf7U2PpXzE
      secret_key: 2wNow7IHpJsSxUR9c+OU1CmCIJVkChSe0ZoFrT7V
```

Use the `oc` command to create the secret in RHOCP.

```
[student@workstation ~]$ oc create -f secret.yaml
secret/thanos-object-storage created
```

- 1.8. Create the multicluster observability object to enable the observability service.

```
[student@workstation ~]$ vi mcobs.yaml
apiVersion: observability.open-cluster-management.io/v1beta2
kind: MultiClusterObservability
metadata:
  name: observability
spec:
  observabilityAddonSpec: {}
  storageConfig:
    metricObjectStorage:
      name: thanos-object-storage
      key: thanos.yaml
```

Use the `oc` command to create the multicluster observability instance in RHOCP.

```
[student@workstation ~]$ oc create -f mcobs.yaml
multiclusterobservability.observability.open-cluster-management.io/observability
created
```

The observability components deployment begins. Verify the status of the installation by using the following command:

```
[student@workstation ~]$ oc describe mco observability
Name:      observability

...output omitted...

Status:
Conditions:
  Last Transition Time: 2021-11-09T13:34:59Z
  Message:            Installation is in progress
  Reason:             Installing
  Status:              True
  Type:                Installing

...output omitted...
```

When the observability components are ready, the output of the previous command displays the following message:

```
Name:      observability
...output omitted...
Status:
Conditions:
Last Transition Time: 2021-11-25T08:29:07Z
Message:          Observability components are deployed and running
Reason:           Ready
Status:            True
Type:              Ready
...output omitted...
```

2. Create a Prometheus alerting rule so that cluster administrators receive alerts when the memory usage of a cluster exceeds 50%.
 - 2.1. From the terminal, create a file named `custom_rule.yaml` containing the ConfigMap object with the instructions in Prometheus Querying Language (PromQL).

```
[student@workstation ~]$ vi custom-rules.yaml
kind: ConfigMap
apiVersion: v1
metadata:
  name: thanos-ruler-custom-rules
data:
  custom_rules.yaml: |
    groups:
      - name: cluster-health
        rules:
          - alert: MemoryUsage-50
            annotations:
              summary: Notify when the memory usage in idle conditions is greater
              than 50%.
            description: "The cluster {{ $labels.cluster }} has more than 50% of
              the memory in use."
```

```
expr: |
  cluster:memory_usage:ratio > 0.5
for: 5s
labels:
  cluster: "{{ $labels.cluster }}"
  severity: warning
```

- 2.2. Create the `thanos-ruler-custom-rules` ConfigMap object in the `open-cluster-management-observability` namespace.

```
[student@workstation ~]$ oc apply -f custom-rules.yaml \
-n open-cluster-management-observability
configmap/thanos-ruler-custom-rules created
```

Creating the ConfigMap object triggers a restart of the `observability-thanos-rule-X` pods in the `open-cluster-management-observability` namespace.

- 2.3. Verify that the new configuration applies. You can monitor the restart of the pods in the `open-cluster-management-observability` namespace with the `watch` command.

```
[student@workstation ~]$ watch oc get pods \
-n open-cluster-management-observability
...output omitted...
NAME                               READY   STATUS
RESTARTS   AGE
...output omitted...
observability-thanos-rule-0        2/2     Running   0
  1m
observability-thanos-rule-1        2/2     Running   0
  1m
observability-thanos-rule-2        2/2     Running   0
  2m
...output omitted...
```

Exit the `watch` command by pressing `Ctrl+C`.

After the `observability-thanos-rule-X` pods restart, the alerting rule is visible from the RHACM Grafana dashboard.

3. From Grafana, verify that the custom alert is firing for the clusters `local-cluster` and `managed-cluster`.

- 3.1. Use Firefox to navigate to the RHACM web console at `https://multicloud-console.apps.ocp4.example.com`.

When prompted, select `htpasswd_provider`.

Type the credentials:

- Username: `admin`
- Password: `redhat`

Then, click `Log in`.

- 3.2. Navigate to `Home` → `Overview` and click `Grafana`.

A new tab showing the Grafana dashboard opens in the browser.

On the Grafana dashboard left menu, click **Explore**.

In the Metrics text field, type `ALERTS{alertname="MemoryUsage-50", alertstate="firing"}` and click **Run Query**.

Scroll down to the **Table** field and verify that the alert with name `MemoryUsage-50` is `firing` in both clusters.



4. Disable the RHACM observability service.

- 4.1. Open a terminal and log in to the ocp4 cluster as the `admin` user. The API Server URL is <https://api.ocp4.example.com:6443>.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
[student@workstation ~]$
```

- 4.2. Use the `oc` command to remove the `MultiClusterObservability` object named `observability`.

```
[student@workstation ~]$ oc delete multiclusertoobservability \
observability
multiclusertoobservability.observability.open-cluster-management.io "observability"
deleted
```

- 4.3. Remove the `open-cluster-management-observability` namespace.

```
[student@workstation ~]$ oc delete namespace \
open-cluster-management-observability
namespace "open-cluster-management-observability" deleted
```



Note

Deleting the `open-cluster-management-observability` namespace can take a while.

Evaluation

As the **student** user on the **workstation** machine, use the **lab** command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade observability-review
```

Finish

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish observability-review
```

This concludes the lab.

Summary



Note

This section is under development.

Chapter 5

Managing Applications Across Multiple Clusters with ACM

Goal

Deploy and manage applications in a multicluster environment with ACM GitOps

Objectives

- Describe and define GitOps concepts and the resources of the RHACM application model.
- Deliver applications into multiple clusters by using RHACM GitOps.
- Describe Kustomize concepts, features, and deploy new customizations with the Kustomize command-line tool.

Sections

- Introducing RHACM GitOps and the Application Model (and Quiz)
- Managing Multicluster Application Resources with RHACM (and Guided Exercise)
- Customizing Resources with Kustomize for RHACM (and Guided Exercise)

Lab

Managing Applications Across Multiple Clusters with RHACM

Introducing RHACM GitOps and the Application Model

Objectives

After completing this section, you should be able to describe and define GitOps concepts and the resources of the RHACM application model.

Introducing GitOps Concepts

GitOps is a framework that consists of procuring the very best practices from Devops and applying them to Infrastructure as Code (IaC). DevOps organizations can deploy code to production clusters effectively and efficiently twenty four hours a day. GitOps implements similar practices such as version control, code review, and CI/CD to automate the process of managing and provisioning infrastructures. GitOps strives to be highly organized maintaining resource files in a repository, such as Git. The resource files located in the Git repository are referred to as the *source of truth*. In other words, the Git repository contains the resources files that generate the same infrastructure environment every time it's deployed, similar to source code used to build binaries. In today's rapidly growing infrastructures, it is essential to have the capability to manage cloud resources that are required for continuous deployments and integration.

GitOps Principles

Implementing and maintaining GitOps principles are essential to your multicloud environment and consists of the following:

- **Declarative desired state**

Systems are required to use declarative data to export the desired state of the system. The data must be in a format that can be written and read by both machines and humans alike.

- **Immutable desired state versions**

The Git version control system acts as a *single source of truth* and all declarative states are stored in Git. In Git, the system infrastructure changes are available chronologically thus assisting users with troubleshooting, auditing, and rollbacks.

GitOps Components

- **Infrastructure as code**

Infrastructure as code is the practice of storing all configurations for managing and provisioning infrastructure in a repository. This ensures that the same environment is provisioned every time. Similar to any software source code, all files under version control can be modified and then deployed. IaC is used for automating infrastructure provisioning and thereby removes the requirement for developers to manually provision and manage resources such as storage and operating systems.

- **Merge request**

Merge Requests are used effectively as a catalyst for all the infrastructure updates. Merge request enable the project developers to collaborate and review IaC prior to deployment to production systems.

- **Continuous integration**

Continuous Integration (CI) is the discipline of **integrating changes in the main branch as often as possible**. Developers use short-lived branches or small change sets and integrate them frequently into the main branch, ideally several times a day. This speeds up the integration, makes code review easier, and reduces potential problems.

Continuous Integration normally entails validation of integrated changes by using automated integration tests. After the changes have been integrated and validated, the team can then decide when to deploy a new release. By itself, Continuous Integration does not involve automating the release process.

Continuous Integration Tools

At a high level, you need the following tools to implement Continuous Integration:

- **A version control system**, which stores your source code in a repository. The repository uses a *main* or *trunk* branch to keep track of the main development line. Feature development occurs in separate branches, which after review and validation, are integrated into the main branch. The most popular version control system is Git.
- **A CI automation service**. This is usually a server or a daemon that watches the repository for changes. If the repository changes, the CI automation service checks out the new code and verifies that everything is correct. RHACM accomplishes this by deploying the subscription operator to monitor Git for repository changes.

GitOps Workflow

GitOps workflows use Git as the version control system which contains the configurations for managing and provisioning infrastructures. GitOps and IaC use a declarative approach as opposed to an imperative approach to configuration management. A declarative model in programming focuses on writing code to describe what an application should do, while an imperative model focus on how it should do it.

A typical GitOps workflow consists on the following:

1. The GitOps team stores all IaC resource files Git.
2. A branch is created from the *main* repository which will contain the updated resource file.
3. After the branch is updated, the administrator pushes the change to Git and creates a pull request.
4. A GitOps colleague reviews the code to ensure proper configuration.
5. The updated code is then approved and merged.
6. The updated branch is then deployed via the subscription operator in RHACM.

Introducing the RHACM Application Model

The foundation of the application model is the ability to subscribe to channel repositories that contain the resources for managing deployments. The application resources definitions are created, managed, and updated using YAML file spec sections. The following inumerates resources available for the application model:

Subscriptions

The subscription (`subscription.apps.open-cluster-management.io`) enables clusters to subscribe to a source repository, also known as a `channel`. The subscription type can be a `Git` repository, `Helm` release registry, or `Object storage` repository. Both the hub and managed cluster can use subscriptions. Subscriptions are associated with the channel and identifies new or updated resource templates.

The subscription operator pulls from the source repository and deploys to the targeted managed clusters without checking the hub cluster first. The subscription operator can then monitor for new or updated resources. For example, the following displays a typical subscription YAML file.

```
apiVersion: apps.open-cluster-management.io/v1 ①
kind: Subscription ②
metadata:
  name: web-app-cluster1 ③
  namespace: web-app ④
  labels:
    deployment: hello
  annotations:
    apps.open-cluster-management.io/github-branch: main ⑤
    apps.open-cluster-management.io/github-path: mysql ⑥
spec:
  channel: web-app/web-app-channel ⑦
  placement:
    placementRef:
      kind: PlacementRule ⑧
      name: cluster1 ⑨
```

- ① The API version is a required value and set to `apps.open-cluster-management.io/v1`.
- ② The `Subscription` is a required value and indicates the `kind` or type of resource.
- ③ The `name` is a required value used for identifying the subscription.
- ④ The `namespace` is a required value and sets the namespace resource to use for the subscription.
- ⑤ `github-branch` specifies the branch that contains the deployment.
- ⑥ `github-path` specifies the path to the resource configuration files.
- ⑦ An Optional value that sets namespace name Namespace/Name and defines the channel for the subscription.
- ⑧ The `PlacementRule` is a required value and indicates that the resource is a placement rule
- ⑨ The `name` is a required value that indicates the name of the cluster where the rule is placed.

Channels

The application channels `channel.apps.open-cluster-management.io` are source definitions of repositories that the cluster can access. The channels can be a Github repository, Helm charts, objects stores, and deployable resource namespaces on the hub cluster. With the exception of

Github, all channels require individual namespaces. The following examples illustrates a typical channel configuration YAML file:

```
apiVersion: apps.open-cluster-management.io/v1 #API name of the object.
kind: Channel
metadata:
  name: web-app-channel
  namespace: web-app
spec:
  pathname: 'https://github.com/redhattraining/do480-apps' ①
  type: GitHub ②
```

- ①** The pathname is a required value that specifies the URL to the repository that contains the IaC resource files.
- ②** The type is a required value that sets the type of repository selected to contain the IaC resource files.

Placement Rules

Placement rules *placementrule.apps.open-cluster-management.io* are used to ensure that the application subscription is placed on the correct cluster in your infrastructure. Placement rules assist in managing a multicloud deployment and can be shared across multiple subscriptions. For example the following illustrates a typical placement rule YAML file.

```
apiVersion: apps.open-cluster-management.io/v1 #API name of the object
kind: PlacementRule
metadata:
  name: cluster1
  namespace: web-app
spec:
  clusterSelector: ①
    matchLabels: ②
      clusterid: cluster1
```

- ①** The clusterSelector is an optional value and identifies the target clusters.
- ②** The matchLabel is an optional value that indicates the labels that must exist on the target clusters.

Applications

Applications *application.app.k8s.io* in Red Hat Advanced Cluster Management for Kubernetes, are used for grouping Kubernetes resources that build an application. For example the following illustrates a typical application YAML file.

```
apiVersion: app.k8s.io/v1 #API name of the object
kind: Application
metadata:
  name: web-app
  namespace: web-app ①
spec:
```

```
selector:  
  matchLabels: ②  
    app: web-app
```

- ① The namespace resource identifies the namespace to use for the application.
- ② The `matchLabel` resource is optional and specifies the label that must exist on the target clusters.



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Applications guide* at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4

► Quiz

Introducing RHACM GitOps and the Application Model

- 1. **Which one of the following statements describing GitOps best practices is the most correct?**
- a. GitOps is a combination of both an internal registry and automated processes.
 - b. GitOps implements the best practices such as version control, code review, and CI/CD to manually provision and monitor pods.
 - c. GitOps implements the best practices such as version control, code review, and CI/CD to automate the process of managing and provisioning infrastructures.
 - d. GitOps implements the best practices such as version control, code review, and CI/CD to virtualize the process of managing and provisioning infrastructures.
- 2. **Which of the following statements describes best practices for maintaining GitOps resource YAML files?**
- a. Resource YAML files are stored on multiple systems, similar to block chains.
 - b. Resource YAML files are stored in a multicluster Postgres database.
 - c. Resource YAML files are stored on the local storage of infrastructure nodes.
 - d. Resource YAML files are stored in a repository such as Git.
- 3. **Which of the following sentences list the three major GitOps components?**
- a. Infrastructure as code (IaC), Centralized Processes, Continuous Delivery/Continuous Integration (CD/CI).
 - b. Continuous Delivery/Continuous Integration (CD/CI), Pull Requests, Peer Review.
 - c. Infrastructure as code (IaC), Merge Request, Continuous Delivery/Continuous Integration (CD/CI).
 - d. Declarative Configurations, Resource YAML files, Infrastructure as code (IaC).
 - e. Application channels lists repositories that the clusters can access.
 - f. Application channels lists are only stored in Helm charts.
 - g. Application channels ensure the application is deployed on a specific cluster.
 - h. Application channels group resources together to build a deployment.
- 4. **Which one of the following statements about GitOps workflows is correct?**
- a. GitOps workflows use an imperative programming model.
 - b. GitOps workflows use both Git and SVN as version control system.
 - c. GitOps workflows use a declarative approach to configuration management.
 - d. GitOps workflows use a systematic approach to infrastructure deployment.

► **5. Which one of the following statements best describes Infrastructure as code (IaC)?**

- a. Infrastructure as code (IaC) is the practice of storing all configurations for managing and provisioning infrastructure on each OpenShift node.
- b. Infrastructure as code (IaC) ensures that the same infrastructure is not provisioned every time.
- c. Infrastructure as code (IaC) enables and encourages administrators to manually provision and manage resources.
- d. Infrastructure as code (IaC) is used for automating provisioning.

► Solution

Introducing RHACM GitOps and the Application Model

- 1. **Which one of the following statements describing GitOps best practices is the most correct?**
- a. GitOps is a combination of both an internal registry and automated processes.
 - b. GitOps implements the best practices such as version control, code review, and CI/CD to manually provision and monitor pods.
 - c. GitOps implements the best practices such as version control, code review, and CI/CD to automate the process of managing and provisioning infrastructures.
 - d. GitOps implements the best practices such as version control, code review, and CI/CD to virtualize the process of managing and provisioning infrastructures.
- 2. **Which of the following statements describes best practices for maintaining GitOps resource YAML files?**
- a. Resource YAML files are stored on multiple systems, similar to block chains.
 - b. Resource YAML files are stored in a multicluster Postgres database.
 - c. Resource YAML files are stored on the local storage of infrastructure nodes.
 - d. Resource YAML files are stored in a repository such as Git.
- 3. **Which of the following sentences list the three major GitOps components?**
- a. Infrastructure as code (IaC), Centralized Processes, Continuous Delivery/Continuous Integration (CD/CI).
 - b. Continuous Delivery/Continuous Integration (CD/CI), Pull Requests, Peer Review.
 - c. Infrastructure as code (IaC), Merge Request, Continuous Delivery/Continuous Integration (CD/CI).
 - d. Declarative Configurations, Resource YAML files, Infrastructure as code (IaC).
 - e. Application channels lists repositories that the clusters can access.
 - f. Application channels lists are only stored in Helm charts.
 - g. Application channels ensure the application is deployed on a specific cluster.
 - h. Application channels group resources together to build a deployment.
- 4. **Which one of the following statements about GitOps workflows is correct?**
- a. GitOps workflows use an imperative programming model.
 - b. GitOps workflows use both Git and SVN as version control system.
 - c. GitOps workflows use a declarative approach to configuration management.
 - d. GitOps workflows use a systematic approach to infrastructure deployment.

► **5. Which one of the following statements best describes Infrastructure as code (IaC)?**

- a. Infrastructure as code (IaC) is the practice of storing all configurations for managing and provisioning infrastructure on each OpenShift node.
- b. Infrastructure as code (IaC) ensures that the same infrastructure is not provisioned every time.
- c. Infrastructure as code (IaC) enables and encourages administrators to manually provision and manage resources.
- d. Infrastructure as code (IaC) is used for automating provisioning.

Managing Multicluster Application Resources with RHACM

Objectives

After completing this section, you should be able to deliver applications into multiple clusters by using RHACM GitOps.

Introducing RHACM GitOps

RHACM provides GitOps capabilities as part of the application deployment lifecycle. RHACM strives to implement GitOps components and adhere to the GitOps principles. RHACM combines Infrastructure as Code (IaC), merge requests, and continuous integration with the GitOps push pattern or method. The GitOps push pattern uses a CI/CD pipeline to push changes to the multicloud applications. The CI/CD pipeline is triggered by a merge commit in Git. The subscription operator monitors the channel for new or updated resources. After a merge commit, the subscription operator downloads the updated resources directly from the storage location and deploys them to the targeted managed clusters. This GitOps pattern is agent-less so there isn't a concern with running agents in each infrastructure component. This GitOps pattern uses well documented CI/CD tooling that is familiar to a large segment of the administrators and thereby enables a quick start to projects with a large pool for collaboration. The deployments are standardized and use the same deployment methodology.

Organizing a Git Repository for RHACM GitOps

The RHACM application lifecycle can consume multiple types of directory structures. The directory structure provides the logical implementation and deployment of the resource YAML files for one or multiple directories. For example, the mysql application contains two folders and one sub-folder named mysql-app, subscriptions, and mysql-sub respectively. The subscription folder targets the mysql-sub folder that contains a single subscription and is downloaded and applied to the hub cluster. From the mysql-sub folder, subscriptions and policies are applied to the managed clusters based on the placement rule. Placement rules determine which managed-clusters are affected by each subscription.

```
└── mysql-app
    └── subscriptions
        └── mysql-sub
```

The directory structure is created for the following subscription flow: subscription > mysql > mysql-apps.

The subscription in the mysql-sub folder is applied from the CLI terminal to the hub cluster. Subscriptions and policies are then downloaded and applied to the hub cluster from the mysql-sub folder. The subscriptions and policies in the mysql-sub folder then run on the managed clusters based on the placement rule. Placement rules determine which managed-clusters are affected by each subscription. The subscriptions or policies determine what is applied to the clusters that match their placement. Moreover, the subscriptions content from the mysql-sub folder is applied to all managed clusters, common applications, and common configurations that match the placement rule.

Managing Git Repository Applications

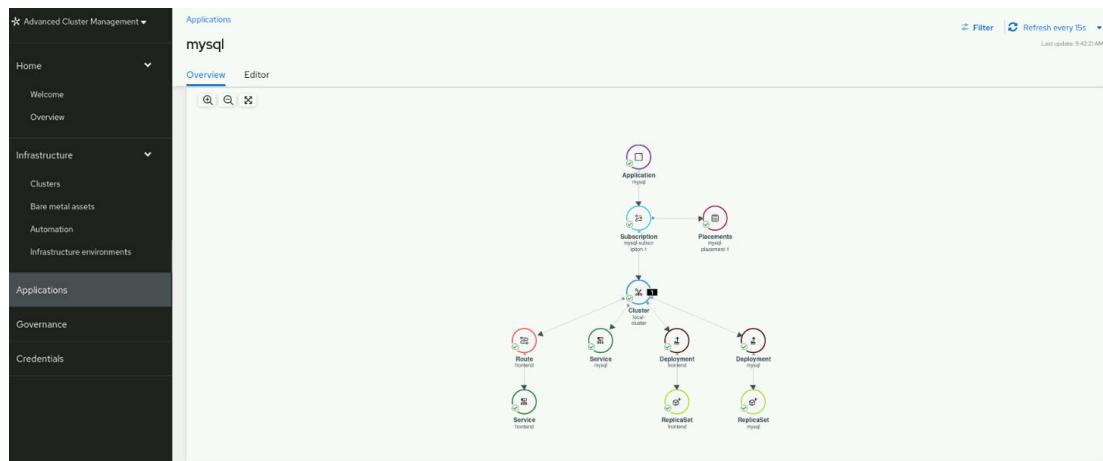
RHACM provides application management functions with concise options for building and deploying applications and updates. Applications are managed across multicluster environments through implementing subscription-based and channel automation. All the application resources are located in the Git repository. Git repositories, Helm repositories, and Object storage repositories are supported, however this course focuses on implementing a Git repository.

Application Console

The application console dashboard manages the application lifecycle. The dashboard includes capabilities to create, manage, and view the application status.

Resource topology

RHACM provides a resource topology page that displays a visual representation of the application and corresponding resources. The visual representation of the state of the resources assists in troubleshooting and confirming application issues.



Advanced configuration

The advanced configuration tab provides the capability to view all application tables and terminology. The Advance configuration tab also includes access to filter for the following resources:

- **Subscription**
- **Placement rules**
- **Channels**

Name	Namespace	Channel	Applications	Clusters	Time window	Created
application-chart-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
assisted-service-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
cluster-lifecycle-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
console-chart-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
discovery-operator-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
grc-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
hive-clusterimagesets-subscription-fast-O	open-cluster-management	acm-hive-openshift-releases-chn-O	0	Local	9 hours ago	
management-ingress-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
policyreport-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	
search-prod-sub	open-cluster-management	charts-v1	0	Local	9 hours ago	

Deploying an application using git repository from RHACM

In preparation for deploying an application using git repository from RHACM, you should define Kubernetes resources for the application in advance, and add them in a Git repository. The RHACM HUB will use these resources to create the desired application architecture in the managed clusters.

To deploy and manage an application with RHACM, you will need to create some custom resources. After creating the resource YAML file, you can verify that the resource definition creates without an error using the `kubectl` command.

Namespace

There must be a namespace resource per deployed application. The following example namespace definition displays the namespace resource created for a MySQL application.

```
apiVersion: v1
kind: Namespace
metadata:
  name: mysql
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f namespace.yaml
namespace/mysql created
```

Subscriptions

Create the subscriptions resource YAML file for the MySQL application. The following subscription resource is an example of the subscription resource created for a MySQL application. The `apiVersion` and full group, `apps.open-cluster-management.io` must be specified in the resource YAML file. By default, the subscription operator subscribes to the `master` branch and you want to use `main`. You can subscribe to a different branch by specifying the branch name annotation in the subscription. You can also specify the path directory, `mysql`, that is used to access the custom resources for deployment.

```
apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
metadata:
  name: mysql-development-subscription
  labels:
    app: mysql
  annotations:
    apps.open-cluster-management.io/github-path: mysql
    apps.open-cluster-management.io/github-branch: main
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f subscription.yaml
subscription/mysql-development-subscription created
```

Channels

Create the channel resource YAML file. Channels define the source repositories that a cluster can subscribe to with a subscription. The following example channel definition displays an example of a Git channel for the `mysql` application. Within the hub cluster, the channel uses the `mysql` namespace. The Channel also specifies the `https://github.com/redhattraining/do480-apps` pathname, which points to the storage location of resources used in the deployment.

```
apiVersion: apps.open-cluster-management.io/v1
kind: Channel
metadata:
  name: mysql-channel
  namespace: mysql
spec:
  pathname: 'https://github.com/redhattraining/do480-apps'
  type: GitHub
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f channel.yaml
channel.apps.open-cluster-management.io/mysql-channel created
```

Placementrule

Create the placementrule resource YAML file. Placementrules describe and specify the target clusters where subscriptions are deployed. The following example placementrule definition specifies the `mysql` namespace and that the application can only occupy the local cluster *hub-cluster*.

```
apiVersion: apps.open-cluster-management.io/v1
kind: PlacementRule
metadata:
  labels:
    app: mysql
  name: mysql-placement-1
  namespace: mysql
spec:
```

```
clusterSelector:  
  matchLabels:  
    'local-cluster': 'true'
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f channel.yaml  
placementrule.apps.open-cluster-management.io/mysql-channel created
```

Services

Create the service resource YAML file. Services are used for pod communication in the network. The following example placementrule definition specifies the service targets pods based on the selector specified label. The service requires a selector to identify it works with any pods with the label `app: todonodejs`

```
apiVersion: v1  
kind: Service  
metadata:  
  labels:  
    app: todonodejs  
    name: mysql  
  name: mysql  
spec:  
  ports:  
    - port: 3306  
  selector:  
    name: mysql
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f service.yaml  
service.apps.open-cluster-management.io/mysql created
```

Routes

Create the route resource YAML file. Routes are used by OpenShift to expose a Service outside the cluster. These routes must have a resource YAML file as part of the deployment or the administrators would have to create the route manually.

```
apiVersion: route.openshift.io/v1  
kind: Route  
metadata:  
  labels:  
    app: todonodejs  
    name: route-frontend  
  name: frontend  
  namespace: mysql  
spec:  
  host: todo.apps.ocp4.example.com  
  path: "/todo"  
  to:  
    kind: Service
```

```
name: frontend
weight: 100
wildcardPolicy: None
```

Use `kubectl apply -f` to verify the resource definition.

```
[student@workstation base]$ kubectl apply -f channel.yaml
route.route.openshift.io/frontend created
```

Deploying an application using git repository from RHACM

After creating the required resource YAML files for the application and pushing them to a Git repository, you are ready to deploy the application with RHACM.

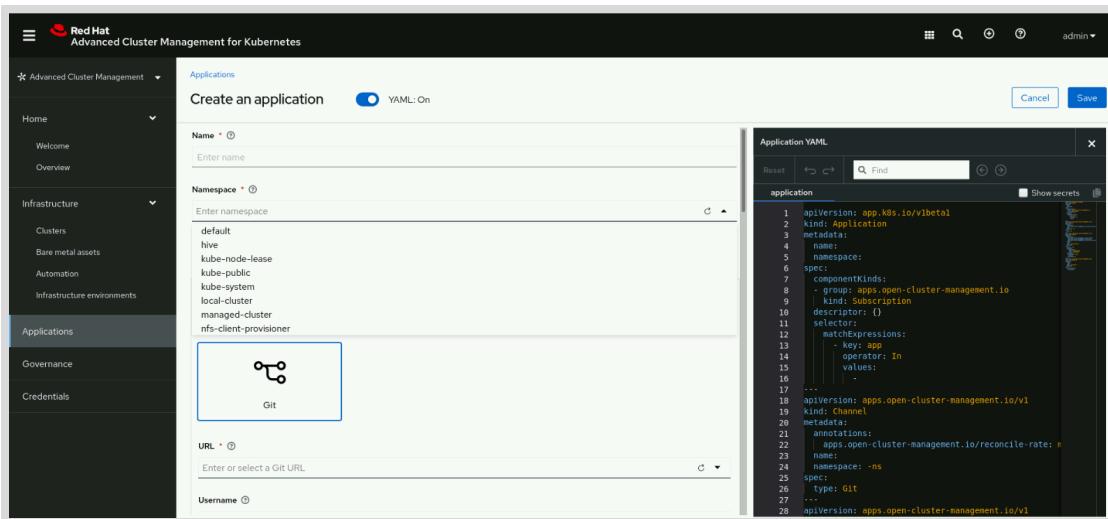
Log in to the console.

Navigate to the left menu and select **Applications** and click **Create application**. RHACM provides a YAML editor that is built directly into the application console. The YAML web UI enables administrators to view, configure, and update the resources in the YAML file directly without leaving the console. The YAML web UI updates and displays in real time as the administrator configures the application.

The screenshot shows the RHACM application creation interface. On the left, there's a form for entering the application name and namespace, along with a dropdown for selecting a repository type (Git, Helm, or Object storage). On the right, there's a modal window titled 'Application YAML' containing the YAML code for the application definition:

```
apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name:
  namespace:
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
      descriptor: {}
      selector:
        matchExpressions:
          - key: app
            operator: In
            values:
              -
```

Enter a valid Kubernetes name and namespace for the application. The drop-down menu provides a list of current namespaces based on your access role. Administrators can select a namespace from drop-down menu or create a new one if they have been assigned the correct access role.



The Location for repository resources menu contains three repository types Git, Helm, and Object storage. Selecting the Git repository displays the following fields:

Git Repository Fields.

Field	Value
URL	Git repository location path.
Branch	The ranch that contains the applicable resource files.
Path	The path to directory that contains the applicable resource files.
Commit hash	To subscribe to a specific Git commit.
Tag	To subscribe to a specific tag.
Reconcile	The default is to reconcile on a commit merge. The other option, replace, replaces the existing resource with the Git resource.
Pre/Post deployment task	Set the Ansible Tower secret for jobs that you want to run before or after the subscription deploys the application resources. The Ansible Tower tasks must be placed within prehook and posthook folders in this repository.

```

apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: 
  namespace: 
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
      selector: {}
      matchExpressions:
        - key: app
          operator: In
          values:
            - ...
  ...
  apiVersion: apps.open-cluster-management.io/v1
  kind: Channel
  metadata:
    annotations:
      apps.open-cluster-management.io/reconcile-rate: #
    name: 
    namespace: -ns
  spec:
    type: Git
  ...
  apiVersion: apps.open-cluster-management.io/v1

```

The **Select clusters to deploy** field allows you to define and update the placementrule allocated for the application. You have the option to deploy the placementrule to the local *hub* cluster, to the online *managed* clusters and the local cluster, or only deploy to clusters that match a specified label. You also have the option to **Select existing placement configuration** if you want to use a previously defined placementrule for an application in an existing namespace.

```

apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: 
  namespace: 
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
      selector: {}
      matchExpressions:
        - key: app
          operator: In
          values:
            - ...
  ...
  apiVersion: apps.open-cluster-management.io/v1
  kind: Channel
  metadata:
    annotations:
      apps.open-cluster-management.io/reconcile-rate: #
    name: 
    namespace: -ns
  spec:
    type: Git
  ...
  apiVersion: apps.open-cluster-management.io/v1

```

From settings, you can specify the application behavior for the **Deployment window** and **Time window**. The **Time window** allows resource subscriptions to begin deployments only during specific days and hours. You can define the subscription **Deployment window** type as **Always active**, **active within specified interval**, or **blocked within specified interval** for the specified time window. For example, you could define an **active** time window between 3:00 PM and 3:30 PM every Monday for web application updates. This **Time window** would allow the application deployments to occur every Monday between 3:00 PM and 3:30 PM. In the event that a deployment begins during a defined time window and is still running when the time window elapses, the deployment will still continue until the deployment is completed. If you were to use **blocked** instead of **active** with the previous example, application deployments could not begin between 3:00 PM and 3:30 PM, but could begin at any other time. The subscription operator continues to monitor regardless of the defined deployment windows. The following is an example **Time window** YAML definition that defines an active, 10:00 PM to 12:00 PM, eastern time zone time window that occurs every Monday.

```

apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
..._output omitted...
spec:
  channel:
    placement:
      placementRef:
        kind: PlacementRule
        name: mysql-placement-1
  timewindow:
    windowtype: active
    location: "America/New_York"
    daysofweek: ["Monday"]
    hours:
      - start: "10:00PM"
        end: "12:00AM"

```

```

application
  apiVersion: app.k8s.io/v1beta1
  kind: Application
  metadata:
    name:
    namespace:
  spec:
    componentKinds:
      - group: apps.open-cluster-management.io
        kind: Subscription
        selector: {}
    selector:
      matchExpressions:
        - key: app
          operator: In
          values:
            - ...
    ...
  ...
  apiVersion: apps.open-cluster-management.io/v1
  kind: Channel
  metadata:
  annotations:
    | apps.open-cluster-management.io/reconcile-rate: 23
    | name: ...
    | namespace: ...
  spec:
    type: Git
  ...

```



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Applications guide* at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4

► Guided Exercise

Managing Multicloud Application Resources with RHACM

In this exercise you will use an RHACM GitOps workflow to deploy a MySQL task-list application to the Development clusters based on an active time window interval.

Outcomes

You should be able to:

- Deliver applications into multiple clusters by using RHACM GitOps.
- Implement application placement rules.
- Define active time interval for application deployments.

Before You Begin

A GitHub account is required to complete this general exercise.

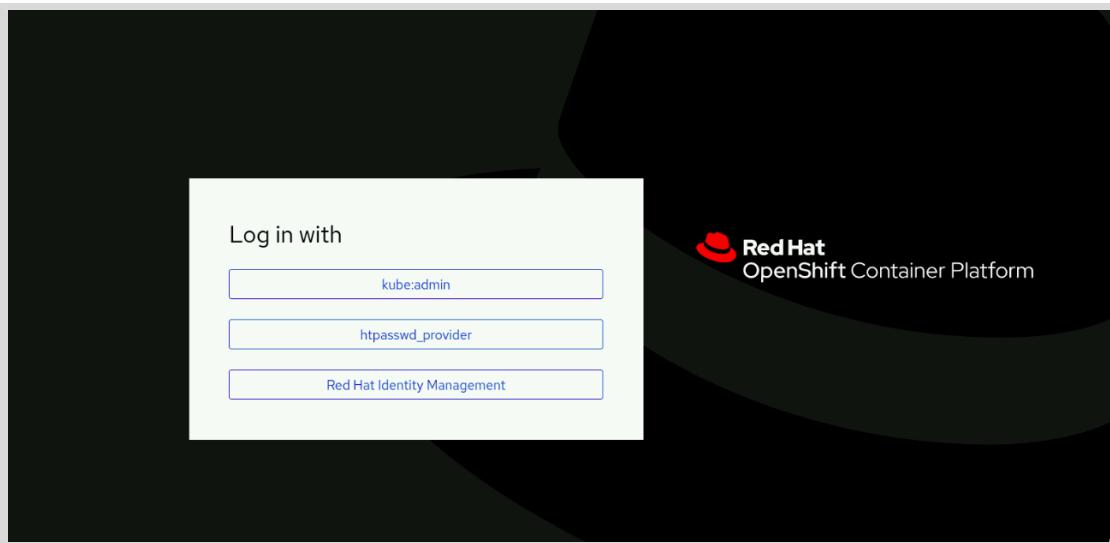
As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the RHACM is installed and that the Development clusters are ready and available.

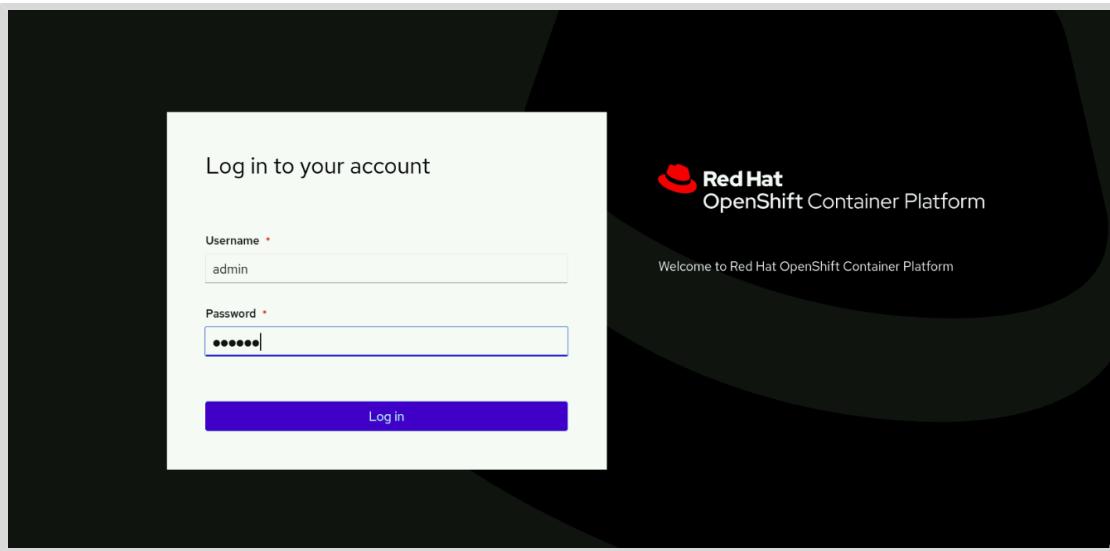
```
[student@workstation ~]$ lab start applications-resources
```

Instructions

- ▶ 1. Fork the `do480-apps` course GitHub repository at <https://github.com/redhat-training/do480-apps/>. This repository contains the YAML files required to build the application.
 - 1.1. From `workstation`, use Firefox to navigate to the `do480-apps` repository at <https://github.com/redhat-training/do480-apps/>.
 - 1.2. In the top-right corner of the GitHub page, click **Fork** to create a fork for your repository.
- ▶ 2. Access the RHACM web console at <https://multicloud.console.apps.ocp4.example.com>
 - 2.1. From the `workstation` machine, open Firefox and access <https://multicloud-console.apps.ocp4.example.com>.



2.2. Click **htpasswd_provider** and log in as the **admin** user with the **redhat** password.



► 3. Use ACM Git Ops to create a new MySQL application based on the following criteria.

MySQL for Development Clusters.

Field	Value
Name	mysql
Namespace	mysql
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	main
Path	mysql

3.1. Navigate to the left menu and select **Applications** and click **Create application**.

3.2. Select **YAML: On** to view the YAML in the console as you create the application.

```

apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name:
  namespace:
spec:
  componentKinds:
  - group: apps.open-cluster-management.io
    kind: Subscription
    selector:
      matchExpressions:
      - key: app
        operator: In
        values:
        - -

```

3.3. In both the **Name** and **namespace** fields, type **mysql**. Applications from the local hub-cluster require a namespace for every application to contain the resources that represent the application in the **managed-clusters**.

```

apiVersion: v1
kind: Namespace
metadata:
  name: mysql
...
apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: mysql
  namespace: mysql
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
      selector:
        matchExpressions:
          - key: app
            operator: In
            values:
              - mysql
...

```

3.4. Select **Git** for the **Repository type** to specify the location of your deployable resources. Input the URL of your forked do480-apps repository for the channel source.

```

apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: mysql
  namespace: mysql
spec:
  componentKinds:
    - group: apps.open-cluster-management.io
      kind: Subscription
      selector:
        matchExpressions:
          - key: app
            operator: In
            values:
              - mysql
...
apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
metadata:
  annotations:
    apps.open-cluster-management.io/git-br: https://github.com/mjarrett0/do480-apps
    apps.open-cluster-management.io/git-pa: mysql
    apps.open-cluster-management.io/reconcile: mysql
  labels:
    app: mysql
  name: mysql-subscription-1
  namespace: mysql

```

3.5. Type **main** for the branch and **mysql** for the path respectively.

```

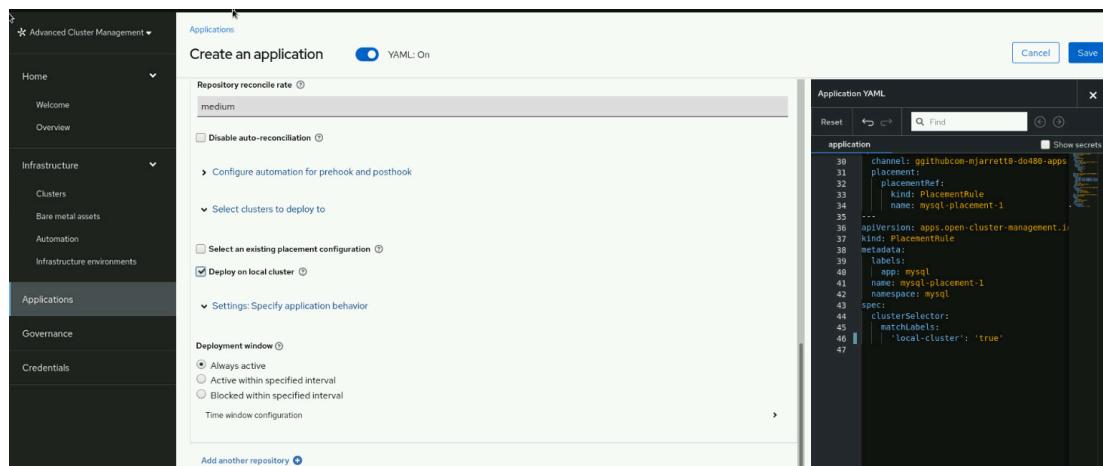
componentKinds:
  - group: apps.open-cluster-management.io
    kind: Subscription
    selector:
      matchExpressions:
        - key: app
          operator: In
          values:
            - mysql
...
apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
metadata:
  annotations:
    apps.open-cluster-management.io/git-br: main
    apps.open-cluster-management.io/git-pa: mysql
    apps.open-cluster-management.io/reconcile: mysql
  labels:
    app: mysql
  name: mysql-subscription-1
  namespace: mysql
spec:
  channel: gitHubUser:mjarrett0:do480-apps
  placement:
    placementRef:
      kind: PlacementRule
      name: mysql-placement-1

```

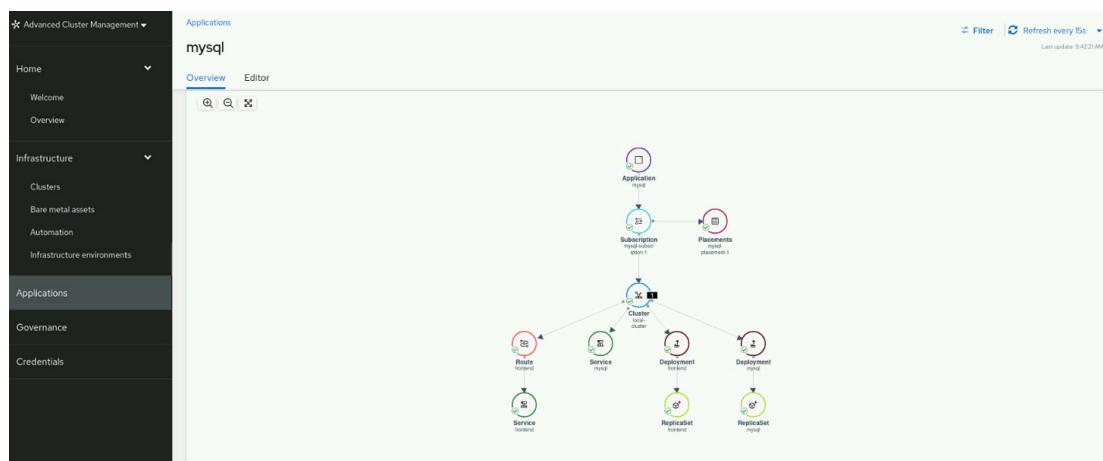
View the subscription and channel configuration in web yaml UI.

```
apiVersion: apps.open-cluster-management.io/v1
kind: Subscription
metadata:
  annotations:
    apps.open-cluster-management.io/git-branch: main
    apps.open-cluster-management.io/git-path: mysql
    apps.open-cluster-management.io/reconcile-option: merge
  labels:
    app: mysql
  name: mysql-subscription-1
  namespace: mysql
spec:
  channel: ggithubcom-mjarrett0-do480-apps-ns/ggithubcom-mjarrett0-do480-apps
```

- 3.6. Scroll down and select **Select clusters to deploy to** to configure a placement rule for the application. Check the box located to the left of **Deploy on local cluster** which creates a placement rule to only deploy the application on the local cluster.

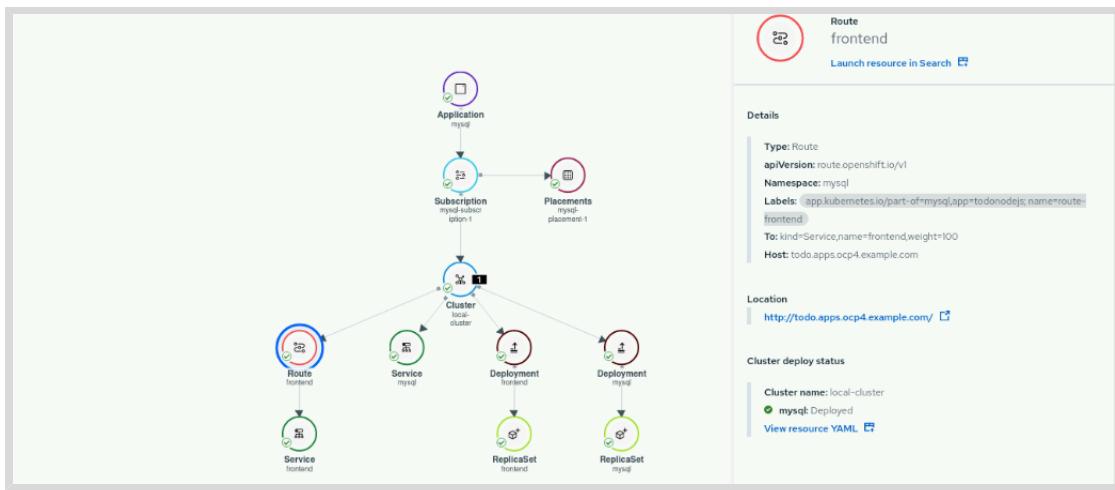


- 3.7. Click the **Save** button, then you are redirected to the applications **Topology** page.



- 4. Verify you can access the MySQL application frontend for the Development clusters.

- 4.1. In the **Topology** page, select the **router** pod, and click the url located under **Location** to view the application frontend on the Development cluster.



4.2. In the subsequent web page, append /todo/ to the url to display the application.

The screenshot shows a web browser displaying the 'To Do List Application'. The page has two main sections: 'To Do List' on the left showing a table of tasks, and 'Add Task' on the right with a form for entering task details. The URL in the address bar is 'todo.apps.ocp4.example.com/todo/'.

► 5. Define an active time window to the deploy updated versions of application resources.

5.1. Click the Applications link at the top of the page.

The screenshot shows the ACM interface with the 'mysql' application selected. The left sidebar shows navigation links for Home, Infrastructure (Clusters, Bare metal assets, Automation, Infrastructure environments), Applications (selected), Governance, and Credentials. The main area displays the application topology for 'mysql', showing the flow from Application mysql through Subscription mysql-autoscale-plan-1 to Cluster local-cluster, which then connects to Route frontend, Service mysql, Deployment frontend, and Deployment mysql.

- 5.2. Select the menu at the end of the mysql application row and click **edit application** to update the application resources.

A screenshot of the ACM application list interface. The mysql application is listed under the Name column. A context menu is open at the end of the mysql row, with 'Edit application' highlighted. Other options in the menu include 'View application', 'Search application', and 'Delete application'. The interface has columns for Name, Namespace, Clusters, Resource, Time window, and Created.

- 5.3. Scroll down and click **Settings: Specify application behavior**. Select the **Active with specified interval** radio button which displays the **Time window configuration**

A screenshot of the 'Settings: Specify application behavior' section. Under 'Deployment window', the 'Active within specified interval' radio button is selected, while 'Always active' and 'Blocked within specified interval' are unselected. The interface includes a 'View application' link.

- 5.4. Specify the day of the week for the time interval. Select the day you are performing this general exercise. For example, today is Thursday and the current time is 05:30. The active time window is defined for between 05:30 and 05:45.

A screenshot of the 'Time window configuration' dialog. Under 'Days of the week', 'Thursday' is selected. Under 'Time zone', 'America/New_York' is chosen. The 'Start Time' is set to '5:30 AM' and the 'End Time' is set to '5:45 AM'.

- 5.5. Click the **Update** button and you are redirected to the applications **Topology** page.



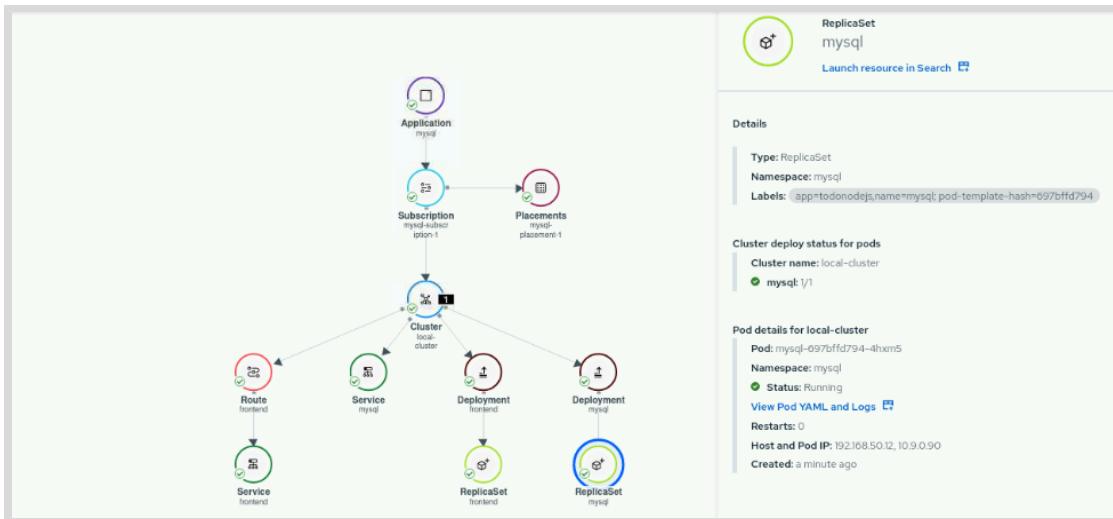
- ▶ 6. Navigate to your forked do480-apps repository in the **main** branch and edit the `do480-apps/mysql/deployment.yaml` file. Update the container image to `mysql-80:152`.
- 6.1. In your github fork, ensure you are in the **main** branch. Click `mysql` and then click the `deployment.yaml` file link to access the file. Next, click the pencil icon to edit the YAML file.
 - 6.2. In the YAML web editor, update the current image tag to `mysql-80:1-156`. Then, scroll to the bottom of the page and click the **Commit changes** button to update the **main** branch.

```

spec:
  containers:
    - image: registry.redhat.io/rhel8/mysql-80:1-152
      name: mysql
      env:

```

- ▶ 7. Return to the RHACM console and view the applications **Topology** page to verify the new pod was created.
- 7.1. In the **Topology** page, select the `mysql` `ReplicaSet` pod. View the **Created** time for the pod as the resource updates and creates a new pod. After a couple minutes, you will see the newly created `mysql` `ReplicaSet` pod.



- 8. Verify resource updates cannot occur outside the active time window. For example, the current time is 5:35, the timezone is US/NY and the time window interval is defined to end at 5:45

- 8.1. After the time window interval expires, in your github fork, ensure you are in the `main` branch. Click `mysql` and then click the `deployment.yaml` file link to access the file. Next, click the pencil icon to edit the YAML file.
- 8.2. In the GitHub YAML editor, update the current image tag to `mysql-80:latest`. Then, scroll to the bottom of the page and click the **Commit changes** button to update the `main` branch.

```
spec:
  containers:
    - image: registry.redhat.io/rhel8/mysql-80:latest
      name: mysql
      env:
```

- 8.3. Return to the RHACM console and view the applications **Topology** page to verify the new pod was created.
- 8.4. In the **Topology** page, select the `mysql` `ReplicaSet` pod. The pod has not been re-created due to the expiration of the active time window interval.

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish applications-resources
```

This concludes the guided exercise.

Customizing Resources with Kustomize for RHACM

Objectives

After completing this section, you should be able to describe Kustomize concepts, features, and deploy new customizations with the Kustomize command-line tool.

Introducing Kustomize

Kustomize is a configuration management solution that allows an administrator to make declarative changes to application configurations and components while preserving the original base YAML files. Kustomize overlays declarative YAML artifacts, or patches that specifically override the general settings without actually modifying the original base files. For example, you could have a multiphase application for multiple departments. The application configuration is predominately identical amongst the different departments, however there are a few settings that are specific to each department and therefore require some customization. Kustomize implements the `kustomization.yaml` file that contains specific settings that override the parameters of the original base configuration. Kustomize enables the GitOps team to consume base YAML file updates while maintaining use-case specific customization overrides.

Kustomize implements the concepts of `bases` and `overlays`. Bases are essentially the base minimum YAML files that contain the settings shared among the application target environments. For example, you could have a base set of files to deploy a database. Overlays, inherit from either one or more bases and are used patch manifests for specific environments or clusters. In reference to the previous example, administrators could have a database with base files for deploying PostGreSql and an overlay that patches the base manifest for an environment to use a specific type of storage.

Kustomize CLI

The Kustomize CLI was embedded in Kubernetes since the March 2019 release of version 1.14. and also maintains a standalone version.

Kustomize benefits

Reusability

Kustomize allows administrators to reuse the base YAML files across multiple environments.

Templateless

Kustomize does not have a templating language and uses the same standard YAML syntax as Kubernetes.

Debug isolation

Kustomize overlays contains the non common or environment specific setting. Due to the isolation of these problematic settings, troubleshooting issues becomes significantly less cumbersome as you only need to focus on the overlays performance and compare that to the performance of the base.

Kustomize files

Kustomize breaks down application YAML files structure into base YAML files and overlay YAML files for environmental specific settings. After the customization parameters have been defined, Kustomize rebuilds and deploys a new manifest with `kubectl kustomize` and `kubectl apply` respectively. Kustomize consists of the following file types:

- **Base YAML files**

The base YAML files contain the application settings that are identical or common to multiple environments.

- **Overlay YAML files**

The patch overlay YAML files contain the application settings that are unique and will override the base file parameters.

- **Kustomization files**

The kustomization YAML files define how the settings in the overlay YAML files should interact with the base YAML files to build the new manifest. The new manifest builds consisting of the base YAML files and their overlays and are called **variants**. Generally, in a multicloud infrastructure, each environment has a corresponding `kustomization.yaml` file that contains specific settings for different environments such as **production** and **development**.

Kustomize Features

The following table lists a subset of the features available from Kustomize:

Kustomize Features

==== | **namePrefix** | prepends a specific prefix to resource names. | **commonLabels** | replaces label values for a specific keyword. | **images** | replaces the image with a customized image name. | **configMapGenerator** | creates K8s ConfigMap resources used as environment variables in the pod definition. | **patchesStrategicMerge** | merges specific configuration segments into to the common base configuration. ===

The infrastructure consists of two specific environments, named **development** and **production**, that run MySQL applications on their respective clusters. Although both environments consist mostly of the same configuration settings, there are a few differences.

First the administrator procures all of the required resource YAML files.

The following is an example of the directory structure:

```

— base
|   — deployment.yaml
|   — kustomization.yaml
|   — service.yaml
— overlays
    — dev
        — kustomization.yaml
    — production
        — kustomization.yaml
        — pvc.yaml

```

The base directory must contain a `kustomization.yaml` file that specifies the resource YAML files that are required to build a new manifest for the application.

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- deployment.yaml
- service.yaml
```

The base directory contains the resource files for the application and must be built with the `kubectl kustomize` command. This command takes the YAML source (via a path or URL) and creates a new YAML that can be piped into the `kubectl create` command.

```
[student@workstation base]$ kubectl kustomize .
```

Patch overlays are created to facilitate the customization of both the production and development environments.

namePrefix

The `namePrefix` directive is used to add a prefixed set of characters to the beginning of resource names. The following `kustomization.yaml` examples apply the prefix `prod` to all the resource names for the production environment and the prefix `dev` to all the resource names in the development environment.

```
#Production
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

#path to base directory
bases:
- ../../base

#Add name prefix "prod"
namePrefix: prod-
```

```
#Development
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

#path to base directory
bases:
- ../../base

#Add name prefix "dev"
namePrefix: dev-
```

commonLabels

The production and development cluster sets require a label for the application that references the cluster set environment. The following `kustomization.yaml` file defines the labels

prod-todolist.js and dev-todolist.js for the keyword app on the production and development cluster sets respectively.

```
#Production  
commonLabels:  
  app: prod-todolist.js
```

```
#Development  
commonLabels:  
  app: dev-todolist.js
```

Note that with commonLabels does not append or prepend characters but directly replaces the values.

images

The production and development cluster sets require an image for the application that references the cluster set environment. The images feature replaces the placeholder mysql-image in the deployment.yaml file. The prod and devel cluster sets images are defined in their kustomization.yaml file using newName and newTag. The following kustomization.yaml files defines the newName value as quay.io/example/todo-single for both the production and development cluster sets. The production cluster sets use the stable version 1.1 that has been tested and secured. However, the development cluster sets implements the latest open source release of the image, v2.0-beta.

```
#Production  
images:  
  - name: mysql-db-image  
    newName: quay.io/example/todo-single  
    newTag: v1.1
```

```
#Development  
images:  
  - name: mysql-db-image  
    newName: quay.io/example/todo-single  
    newTag: v2.0-beta
```

patchesStrategicMerge

The development cluster sets use the default emptyDir setting with ephemeral storage. However, the production environment requires persistent volumes and persistent volume claims for storage. The PVCs are defined in a separate YAML file named pvc.yaml. You can use the patchesStrategicMerge to apply the volumes defined in the pvc.yaml. In the kustomization.yaml file for the production environment, the pvc.yaml is defined under resources.

```
#Production  
patchesStrategicMerge:  
  - pvc.yaml
```



References

For more information, refer to the *Red Hat Advanced Cluster Management for Kubernetes Applications guide* at
https://access.redhat.com/documentation/en-us/red_hat_advanced_cluster_management_for_kubernetes/2.4

► Guided Exercise

Customizing Resources with Kustomize for RHACM

In this exercise you will use Kustomize to maintain a common set of deployment resources and build overlays. The overlays modify the base deployment resources for specific environments. You will build deployment overlays that modify available storage and replicas on the production clusters.

Outcomes

You should be able to:

- Build a base environment for Kustomize.
- Build overlays to modify the base deployment.
- Configure kustomization.yaml to add resources.
- Deploy an application modified and managed with Kustomize

Before You Begin

Fork the do480-apps course GitHub repository at <https://github.com/redhattraining/do480-apps/>. This repository contains the YAML files required to build the application.

As the student user on the workstation machine, use the `lab` command to prepare your system for this exercise.

This command ensures that the RHACM is installed and that the Development clusters are ready and available.

```
[student@workstation ~]$ *lab start applications-kustomize*
```

Instructions

- 1. From workstation, use Kustomize to build the base directory and corresponding `kustomization.yaml` file in the `~/D0480/labs/applications-kustomize` directory. The base directory contents are located at http://github.com/<your_fork>/do480-apps in the `kustomize` branch.
- 1.1. Open the terminal application on the workstation machine. Log in to the Hub OpenShift cluster as the admin user.

```
[student@workstation ~]$ oc login -u admin -p redhat \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Create a new project named mysql.

```
[student@workstation ~]$ oc new-project mysql
Now using project "mysql" on server "https://api.ocp4.example.com:6443".
...output omitted...
```

- 1.3. Navigate to the D0480/labs/applications-kustomize project directory and create the base directory which will contain the set of common yaml resource files.

```
[student@workstation ~]$ cd D0480/labs/applications-kustomize
```

```
[student@workstation applications-kustomize]$ mkdir base
```

- 1.4. Use Git to clone your do480-apps forked repository to access the common set of yaml resource files

```
[student@workstation applications-kustomize]$ git clone \
https://github.com/<your_fork>/do480-apps
Cloning into 'do480-apps'...
remote: Enumerating objects: 144, done.
remote: Counting objects: 100% (144/144), done.
remote: Compressing objects: 100% (103/103), done.
remote: Total 144 (delta 56), reused 104 (delta 36), pack-reused 0
Receiving objects: 100% (144/144), 18.03 KiB | 18.03 MiB/s, done.
Resolving deltas: 100% (56/56), done.
```

- 1.5. Navigate to the do480-apps directory and change to the kustomize branch to access the mysql deployment yaml files.

```
[student@workstation applications-kustomize]$ cd do480-apps
```

```
[student@workstation do480-apps]$ git checkout kustomize
Branch 'kustomize' set up to track remote branch 'kustomize' from 'origin'.
Switched to a new branch 'kustomize'
```

- 1.6. Navigate to the base directory, copy the contents of the kustomize directory to your base directory and list the contents.

```
[student@workstation do480-apps]$ cd ../base/
```

```
[student@workstation base]$ cp ../../do480-apps/kustomize/* .
```

```
[student@workstation base]$ ll
total 32
-rw-rw-r-- 1 student student 165 Dec 24 06:10 applications.yaml
-rw-rw-r-- 1 student student 807 Dec 24 06:10 deployment-frontend.yaml
-rw-rw-r-- 1 student student 692 Dec 24 06:10 deployment.yaml
-rw-rw-r-- 1 student student 65 Dec 24 06:10 namespace.yaml
-rw-rw-r-- 1 student student 227 Dec 24 06:10 placementrule.yaml
```

```
-rw-rw-r--. 1 student student 287 Dec 24 06:10 route.yaml
-rw-rw-r--. 1 student student 166 Dec 24 06:10 service-frontend.yaml
-rw-rw-r--. 1 student student 165 Dec 24 06:10 service.yaml
```

- 1.7. Build the base directory `kustomization.yaml` file which contains the set of resources and associated customizations. Your `kustomization.yaml` file should contain the following:

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- applications.yaml
- deployment-frontend.yaml
- deployment.yaml
- service.yaml
- service-frontend.yaml
- namespace.yaml
- placementrule.yaml
- route.yaml
```

- 1.8. Verify the `kustomization.yaml` file builds the deployment without errors with the `kubectl kustomize` command.

```
[student@workstation base]$ kubectl kustomize .
apiVersion: v1
kind: Namespace
metadata:
  name: mysql
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: todonodejs
    name: frontend
    name: frontend
spec:
  ports:
  - port: 8080
  selector:
    name: frontend
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: todonodejs
    name: mysql
    name: mysql
spec:
  ports:
  - port: 3306
  selector:
```

```
name: mysql
---
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: todonodejs
    name: frontend
  name: frontend
  namespace: mysql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: todonodejs
      name: frontend
  template:
    metadata:
      labels:
        app: todonodejs
        name: frontend
    spec:
      containers:
        - env:
            - name: MYSQL_ENV_MYSQL_DATABASE
              value: items
            - name: MYSQL_ENV_MYSQL_USER
              value: user1
            - name: MYSQL_ENV_MYSQL_PASSWORD
              value: mypa55
            - name: APP_PORT
              value: "8080"
          image: quay.io/redhattraining/todo-single:v1.0
          name: todonodejs
        ports:
          - containerPort: 8080
            name: nodejs-http
        resources:
          limits:
            cpu: "0.5"
    ---
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: todonodejs
    name: mysql
  name: mysql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: todonodejs
      name: mysql
  template:
```

```
metadata:
  labels:
    app: todonodejs
    name: mysql
spec:
  containers:
    - env:
        - name: MYSQL_ROOT_PASSWORD
          value: r00tpa55
        - name: MYSQL_USER
          value: user1
        - name: MYSQL_PASSWORD
          value: mypa55
        - name: MYSQL_DATABASE
          value: items
      image: registry.redhat.io/rhel8/mysql-80:1-156
      name: mysql
      ports:
        - containerPort: 3306
          name: mysql
    ---
apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: mysql
spec:
  selector:
    matchLabels:
      app: mysql
  ---
apiVersion: apps.open-cluster-management.io/v1
kind: PlacementRule
metadata:
  labels:
    app: mysql
  name: mysql-placement-1
  namespace: mysql
spec:
  clusterSelector:
    matchLabels:
      local-cluster: "true"
  ---
apiVersion: route.openshift.io/v1
kind: Route
metadata:
  labels:
    app: todonodejs
    name: route-frontend
  name: frontend
  namespace: mysql
spec:
  host: todo.apps.ocp4.example.com
  path: /todo
  to:
```

```
kind: Service
name: frontend
weight: 100
wildcardPolicy: None
```

Now you have a base set of deployment resources that can be shared amongst both the **Production** and **production** administrators.

- ▶ 2. Examine the `deployment.yaml` in the base directory and note that the storage is not configured. The production clusters require that the mysql applications use persistent volumes and persistent volume claims for storage. Build an overlay for the production clusters that uses persistent storage.
 - 2.1. Create the `overlay` directory and the corresponding `production` and `development` subdirectories. Ensure the `overlays` directory is at the same directory level as the base directory.

```
[student@workstation base]$ tree ../
base
overlays/
└── development
└── production
```

- 2.2. Create the `dbclaim-pv.yaml` file resource, which contains the persistent volume claim, in the `overlays/production/` subdirectory. Your `dbclaim-pv.yaml` file should contain the following:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim
  namespace: mysql
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Mi
  storageClassName: nfs-storage
```

- 2.3. Create the corresponding `kustomization.yaml` file in the `overlay/production/` subdirectory to add the new resource. Your `kustomization.yaml` file should contain the following:

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

bases:
  - ../../base

resources:
- dbclaim-pv.yaml
```

```
[student@workstation base]$ cat ../overlays/production/kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

bases:
- ../../base

resources:
- dbclaim-pv.yaml
```

- 2.4. Use the `kubectl kustomize` command to build the deployment for the production clusters with the new storage resource. Verify the new storage resource has been added to the deployment.

```
[student@workstation base]$ kubectl kustomize ../overlays/production
...output_omitted...
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pv-claim
  namespace: mysql
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 10Mi
  storageClassName: nfs-storage
...output_omitted...
```

- 2.5. The output of the previous command can now be applied to the production clusters. Use the `kubectl kustomize` command piped to `kubectl apply -f` to build and deploy the application to the production clusters with the new storage resource.

```
[student@workstation base]$ kubectl kustomize ../overlays/production | kubectl
  apply -f -
namespace/mysql created
service/frontend created
service/mysql created
persistentvolumeclaim/mysql-pv-claim created
deployment.apps/frontend created
deployment.apps/mysql created
application.app.k8s.io/mysql created
placementrule.apps.open-cluster-management.io/mysql-placement-1 created
route.route.openshift.io/frontend created
```

- 2.6. Verify the deployment and confirm that the new storage resource is deployed.

```
[student@workstation base]$ oc get pods
NAME                  READY   STATUS    RESTARTS   AGE
frontend-5cbdf5bf85-94qwl   1/1     Running   0          12m
mysql-6c5d98dd9b-q94c8     1/1     Running   0          12m
```

```
[student@workstation base]$ oc get pvc
NAME           STATUS  VOLUME                                     CAPACITY
ACCESS MODES   STORAGECLASS   AGE
mysql-pv-claim  Bound   pvc-c95583c8-3fea-45d4-bcba-330a33eea6ff  10Mi
              nfs-storage   5m32s
```

- 2.7. From **workstation**, use Firefox to navigate to the **mysql** application at <http://todo.apps.ocp4.example.com/todo/> to verify the application.
- 2.8. Use the **kubectl kustomize** command to confirm that the base yaml files have not been altered.

```
[student@workstation base]$ kubectl kustomize .
apiVersion: v1
kind: Namespace
metadata:
  name: mysql
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: todonodejs
    name: frontend
    name: frontend
spec:
  ports:
  - port: 8080
  selector:
    name: frontend
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: todonodejs
    name: mysql
    name: mysql
spec:
  ports:
  - port: 3306
  selector:
    name: mysql
---
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
labels:
  app: todonodejs
  name: frontend
  name: frontend
  namespace: mysql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: todonodejs
      name: frontend
  template:
    metadata:
      labels:
        app: todonodejs
        name: frontend
    spec:
      containers:
        - env:
            - name: MYSQL_ENV_MYSQL_DATABASE
              value: items
            - name: MYSQL_ENV_MYSQL_USER
              value: user1
            - name: MYSQL_ENV_MYSQL_PASSWORD
              value: mypa55
            - name: APP_PORT
              value: "8080"
          image: quay.io/redhattraining/todo-single:v1.0
          name: todonodejs
          ports:
            - containerPort: 8080
              name: nodejs-http
          resources:
            limits:
              cpu: "0.5"
      ...
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: todonodejs
    name: mysql
  name: mysql
spec:
  replicas: 1
  selector:
    matchLabels:
      app: todonodejs
      name: mysql
  template:
    metadata:
      labels:
        app: todonodejs
        name: mysql
    spec:
```

```
containers:
  - env:
      - name: MYSQL_ROOT_PASSWORD
        value: r00tpa55
      - name: MYSQL_USER
        value: user1
      - name: MYSQL_PASSWORD
        value: mypa55
      - name: MYSQL_DATABASE
        value: items
    image: registry.redhat.io/rhel8/mysql-80:1-156
    name: mysql
    ports:
      - containerPort: 3306
        name: mysql
    ...
  apiVersion: app.k8s.io/v1beta1
  kind: Application
  metadata:
    name: mysql
  spec:
    selector:
      matchLabels:
        app: mysql
    ...
  apiVersion: apps.open-cluster-management.io/v1
  kind: PlacementRule
  metadata:
    labels:
      app: mysql
    name: mysql-placement-1
    namespace: mysql
  spec:
    clusterSelector:
      matchLabels:
        local-cluster: "true"
    ...
  apiVersion: route.openshift.io/v1
  kind: Route
  metadata:
    labels:
      app: todonodejs
      name: route-frontend
    name: frontend
    namespace: mysql
  spec:
    host: todo.apps.ocp4.example.com
    path: /todo
    to:
      kind: Service
      name: frontend
      weight: 100
    wildcardPolicy: None
```

- 3. Create an overlay for the production clusters to increase the number of mysql database replicas from one to three.

- 3.1. Verify the current number of mysql replicas.

```
[student@workstation base]$ oc get pods
NAME                  READY   STATUS    RESTARTS   AGE
frontend-5cbdf5bf85-94qwl  1/1     Running   0          12m
mysql-6c5d98dd9b-q94c8    1/1     Running   0          12m
```

- 3.2. Modify the kustomization.yaml file to increase the number of mysql replica pods. Your kustomization.yaml file should contain the following:

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

bases:
- ../../base

resources:
- dbclaim-pv.yaml

replicas:
- name: mysql
  count: 3
```

```
[student@workstation base]$ cat ./overlays/production/kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

bases:
- ../../base

resources:
- dbclaim-pv.yaml

replicas:
- name: mysql
  count: 3
```

- 3.3. Use the kubectl kustomize command piped to kubectl apply -f to build and deploy the application and verify the number of mysql replicas.

```
[student@workstation base]$ kubectl kustomize ./overlays/production | kubectl
apply -f -
namespace/mysql unchanged
service/frontend unchanged
service/mysql unchanged
persistentvolumeclaim/mysql-pv-claim unchanged
deployment.apps/frontend configured
deployment.apps/mysql configured
```

```
application.app.k8s.io/mysql unchanged
placementrule.apps.open-cluster-management.io/mysql-placement-1 unchanged
route.route.openshift.io/frontend unchanged
```

```
[student@workstation base]$ oc get pods
NAME           READY   STATUS    RESTARTS   AGE
frontend-5cbdf5bf85-94qwl   1/1     Running   0          12m
mysql-6c5d98dd9b-25mss      1/1     Running   0          6s
mysql-6c5d98dd9b-2nj2j      1/1     Running   0          6s
mysql-6c5d98dd9b-q94c8      1/1     Running   0          12m
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish applications-kustomize
```

This concludes the guided exercise.

► Lab

Managing Applications Across Multiple Clusters with RHACM

In this lab, you will use RHACM GitOps application management to promote a new version of the `todo` application from the **Development** clusters to the **Production** clusters.

Outcomes

You should be able to:

- Configure and deploy a MySQL application with RHACM Git Ops.
- Assign additional labels to clusters managed with RHACM.
- Promote an application from development to production.

Before You Begin

A GitHub account is required to complete this lab.

The following command ensures that the RHACM is installed and that the **Development** and **Production** clusters are ready and available.

```
[student@workstation ~]$ lab start application-review
```

Instructions

1. Fork the `do480-apps` course GitHub repository at <https://github.com/redhattraining/do480-apps/>. This repository contains the YAML files required to build the application.
2. Use the RHACM web console to add the labels `usage=development` and `usage=production` to the `local-cluster` and `managed-cluster` clusters respectively. The web console URL is <https://multicloud.console.apps.ocp4.example.com>
3. Use ACM Git Ops to create a new MySQL application based on the following:

Mysql for development clusters.

Field	Value
Name	mysql
Namespace	mysql
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	main
Path	mysql
Label name	environment
Label value	development

Add an additional repository.

Mysql for Production clusters.

Field	Value
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	production
Path	mysql
Label name	environment
Label value	production

4. Verify you can access the MySQL application frontend for the Development and Production clusters.
5. In Github, use your fork of the do480-apps repository to promote the Development MySQL image mysql-80:1-156 to Production.
6. Return to the RHACM web console and refresh the application Topology page. Verify the Mysql application is working and using the correct image.

Finish

As the student user on the workstation machine, use the lab command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish application-review
```

This concludes the lab.

► Solution

Managing Applications Across Multiple Clusters with RHACM

In this lab, you will use RHACM GitOps application management to promote a new version of the todo application from the Development clusters to the Production clusters.

Outcomes

You should be able to:

- Configure and deploy a MySQL application with RHACM Git Ops.
- Assign additional labels to clusters managed with RHACM.
- Promote an application from development to production.

Before You Begin

A GitHub account is required to complete this lab.

The following command ensures that the RHACM is installed and that the Development and Production clusters are ready and available.

```
[student@workstation ~]$ lab start application-review
```

Instructions

1. Fork the do480-apps course GitHub repository at <https://github.com/redhattraining/do480-apps/>. This repository contains the YAML files required to build the application.
 - 1.1. From workstation, use Firefox to navigate to the do480-apps repository at <https://github.com/redhattraining/do480-apps/>.
 - 1.2. In the top-right corner of the GitHub page, click **Fork** to create a fork for your repository.
2. Use the RHACM web console to add the labels `usage=development` and `usage=production` to the `local-cluster` and `managed-cluster` clusters respectively. The web console URL is <https://multicloud.console.apps.ocp4.example.com>
 - 2.1. From workstation, use Firefox to navigate to the RHACM web console at <https://multicloud.console.apps.ocp4.example.com>.
When prompted, select `htpasswd_provider`
Type the following credentials
 - Username: `admin`
 - Password: `redhat`And then click **Log in**

- 2.2. In the left navigation menu, select **Clusters** to add new labels to the **local-cluster** and **managed-cluster** clusters.
 - 2.3. On the **local-cluster** row, select the drop-down menu at the end of the row, and click **edit label**.
 - 2.4. Type **environment=development** for the new label and click **Save**.
 - 2.5. On the **managed-cluster** row, select the drop-down menu at the end of the row, and click **edit label**.
 - 2.6. Type **environment=production** for the new label and click **Save**.
3. Use ACM Git Ops to create a new MySQL application based on the following:

Mysql for development clusters.

Field	Value
Name	mysql
Namespace	mysql
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	main
Path	mysql
Label name	environment
Label value	development

Add an additional repository.

Mysql for Production clusters.

Field	Value
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	production
Path	mysql
Label name	environment
Label value	production

- 3.1. Navigate to the left menu and select **Applications** and click **Create application**.
- 3.2. Create a MySQL application using the following:

Do **not** click save.

MySQL for Development Clusters.

Field	Value
Name	mysql
Namespace	mysql
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	main
Path	mysql
Label name	environment
Label value	development

- 3.3. Add an additional repository to the MySQL application for the Production cluster based on the following and click Save:

MySQL for Production Clusters.

Field	Value
Repository types	GitHub
URL	github.com/<your fork>/do480-apps/
Branch	production
Path	mysql
Label name	environment
Label value	production

4. Verify you can access the MySQL application frontend for the Development and Production clusters.
- 4.1. In the **Topology** page and navigate to **Subscription → mysql-subscription-1**. Select the **router** pod and click the url located under **Location** to view the application frontend on the Development cluster. In the subsequent web page, append **/todo/** to the url to display the application.
- 4.2. In the **Topology** page and navigate to **Subscription → mysql-subscription-2**. Select the **router** pod and click the url located under **Location** to view the application frontend on the Production cluster. In the subsequent web page, append **/todo/** to the url to display the application.
5. In Github, use your fork of the do480-apps repository to promote the Development MySQL image **mysql-80:1-156** to Production.

- 5.1. In your github fork, switch to the **production** branch. Click **mysql** and then click the **deployment.yaml** link to access the file. Next, click the **pencil** icon to edit the YAML file.
- 5.2. Scroll down and replace the current image tag with **mysql-80:1-156**. Then scroll to the bottom of the page, select **Create a new branch for this commit and start a pull request**, accept the default branch, and click **Proposed changes**.
- 5.3. To complete the workflow, click **Create pull request**, **Merge pull request**, and finally **Confirm merge**. After your pull request is successfully merged and closed, click **Delete branch**.
6. Return to the RHACM web console and refresh the application **Topology** page. Verify the **Mysql** application is working and using the correct image.
 - 6.1. In the **Topology** page, select the **Subscription → mysql-subscription-2** to view the **managed-cluster** topology.
 - 6.2. Select the **mysql ReplicaSet** pod and click **View Pod YAML and Logs** to view the YAML file and verify the image tag.
 - 6.3. In the **YAML** console UI, scroll down and verify the image is **mysql-80:1-156**.
 - 6.4. In the left navigation menu, select the **Applications → mysql** link to view the application.
 - 6.5. In the **Topology** page, select **Subscription → mysql-subscription-2** and select the **router** pod. Click **todo.apps.ocp4-mng.example.com/**, located under **Location**, to view the application. In the subsequent web page, append **/todo/** to the url to display the application.

Finish

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise. This is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish application-review
```

This concludes the lab.

Summary

In this chapter, you learned:

- About the benefits of implementing the GitOps concepts of **Infrastructure as Code (IaC)**, **Merge Commits**, and **Continuous Integration and Continuous Deployment (CI/CD)**.
- How to apply the GitOps principles of **declarative desired state** and **immutable desired state versions** for the application lifecycle in RHACM.
- How to deploy a subscription based application with RHACM.
- How to leverage **Kustomize** to apply and manage custom settings for multicloud environments.

